

Time Series Analysis for Economics and Finance

Zeming Wang (PhD)

September 2025

Table of contents

Preface	7
I The Basics	9
1 Time Series Data	10
Raw data	10
Growth rates	10
Seasonally-adjusted data	11
Logarithms	12
2 Decomposition	13
2.1 Time Series Components	13
2.2 Moving Averages	13
2.3 Classical Decomposition	15
2.4 Seasonal Adjustment	16
3 ACF and PACF	18
3.1 Autocorrelation	18
3.2 Partial Autocorrelation	19
4 Stationarity	21
4.1 Stationary Process	21
4.2 Ergodicity	23
4.3 White Noise	24
II ARIMA Model	25
5 Model vs Spec	26
5.1 Classification	26
5.2 Model vs Spec	26
6 AR Models	28
6.1 AR(1) Process	28
6.2 Lag Operator	33

6.3	AR(p) Process	33
7	MA Models	36
7.1	MA(1) Process	36
7.2	MA(q) Process	38
7.3	MA(∞) Process	39
8	ARMA Models	40
8.1	ARMA(p,q)	40
8.2	ARIMA(p,d,q)	41
9	Wold Theorem	42
9.1	Wold Decomposition	42
9.2	Causality and Invertibility*	44
III	Time Series Regression	46
10	Preliminaries	47
10.1	Chapter Overview	47
10.2	Asymptotic Theorems for i.i.d Random Variables	47
10.3	OLS for i.i.d Random Variables	48
11	OLS for Time Series	51
11.1	Asymptotic Theorems for Dependent Random Variables	51
11.2	OLS for Time Series	53
11.3	HAC Standard Errors	54
11.4	Example	55
12	MLE for ARMA Models	56
13	Forecasting	59
13.1	Intuitive Approach	59
13.2	Best Linear Predictor	60
13.3	Forecasting with ARMA Models	61
13.3.1	Forecasting with AR(p)	61
13.3.2	Forecasting with MA(q)	62
13.3.3	Forecasting with ARMA(p,q)	63
13.4	Applications	64
14	Dynamic Causal Effect	67
15	The Structural Shock Framework	69

16 Estimating Dynamic Multipliers	71
16.1 Distributed Lags	71
16.2 Local Projections	72
16.3 Example of Observable Exogenous Shocks	73
16.4 Example of Constructed Structural Shocks	76
17 Instrument Variables	80
 IV Nonstationary Time Series	 83
18 Spurious Regression	84
19 Trend Stationary	89
20 Unit Root Process	93
21 Brownian Motion	96
21.1 Continuous random walk	96
21.2 Some properties	96
21.3 Itô calculus	99
21.4 Unit root process	102
22 Unit Root Process (contd)	106
22.1 Univariate case	106
22.2 Spurious regression	107
22.3 Cures for spurious regression	108
22.4 Summary	109
23 Unit Root Test	110
23.1 Dickey-Fuller Test	110
23.2 Augmented Dickey-Fuller Test	111
23.3 Phillips-Perron Test	112
24 Cointegration	113
24.1 Cointegration and super-consistency	113
24.2 Inference under cointegration	115
24.3 Conclusion	118
 V Vector Autoregression	 119
25 System of Equations	120

26 Vector Processes	122
26.1 Definitions	122
26.2 VAR and VMA	123
26.3 Stationary Conditions	124
26.4 Autocovariance Matrix	124
27 Estimating VAR	126
27.1 Stacked Form	126
27.2 Vectorized Form	127
28 Granger Causality	129
28.1 Definition	129
28.2 Granger Causality Test	129
28.3 Granger Causality in VAR	130
28.4 Likelihood Ratio Test	131
29 Structural VAR	133
29.1 The Structural Framework	133
29.2 Invertibility	134
29.3 Identification	135
29.3.1 Recursive restriction	135
29.3.2 Non-recursive restriction	136
30 IRF and FEVD	138
30.1 Estimating SVAR	138
30.2 Impulse-Response Functions	138
30.3 IRF Standard Errors	140
30.4 FEVD	140
30.5 Example	141
31 Factor Models	144
31.1 Principle Component Analysis	144
31.2 Factor-augmented VAR	147
32 Unit Roots in VAR	150
32.1 Monte Carlo	151
32.2 Conclusions	152
33 VECM*	154
33.1 Cointegrated Systems	154
33.2 Canonical Correlation	155
33.3 Johansen's Procedure	157
33.4 Hypothesis Testing	160
Test 1: At most h cointegrations	160

Test 2: h cointegrations vs $h+1$	161
VI Bayesian Analysis	162
34 Intro to Bayes	163
34.1 The Sunrise Problem	163
34.2 The Bayesian Approach	165
34.3 Frequentist vs Bayesian	166
35 Linear Model	168
35.1 Linear Regression with Known σ^2	168
35.2 Linear Regression with Unknown σ^2	169
35.3 Credible Interval	170
36 Bayesian VAR	171
36.1 Vectorized Form	171
36.2 Minnesota Prior	172
36.3 The Posterior	173
37 Conjugate Priors	174
38 Gibbs Sampling	175
38.1 Computational Bayes	175
38.2 Gibbs Sampler for Linear Regression Models	175
38.3 Gibbs Sampler for General Models	176
39 Metropolis–Hastings	177
39.1 Dependent Sampling	177
39.2 Random Walk Metropolis	177
39.3 Remarks	178
References	180

Preface

Empirical finance and economic studies often involves the analysis of time series data, such as GDP, inflation, and interest rates, which are distinct from those utilized in cross-sectional studies. The goal of this book is to bridge the gap between introductory time series textbooks and theoretical econometrics. In modern applied research, a rudimentary comprehension of the subject often proves insufficient. Though computational tasks can be executed through simple computer commands, practitioners must go below the surface to understand the intricacies and limitations of the techniques involved. However, an exhaustive exploration of advanced econometric theories would be excessive for practical purposes. For instance, introductory textbooks would caution against running OLS on non-stationary time series, citing the risk of spurious regression. Students often accept this as a rule of thumb without a grasp of its underlying rationale. Yet, delving into intricate topics such as Itô calculus is unnecessary for empirical researchers.

This book seeks to acquaint readers with the time series topics essential for understanding and conducting empirical research, with a focus on macroeconomic applications. In addition to introducing basic concepts and applications, the book endeavors to elevate comprehension to a deeper level by elucidating the “why” alongside the “what” and “how.” However, the objective is not to provide an exhaustive treatment replete with formal proofs; rather, emphasis is placed on providing intuitive explanations. Consequently, readers may encounter instances of informal proofs where a more formal approach is deemed unnecessary for applied works. This book can be read as intermediary materials between undergraduate econometrics and more rigorous treatments of the subject, such as Hamilton’s *Time Series Analysis*.

The materials presented are drawn from or influenced by various sources, which are listed in the References at the end of the book without being cited individually in the context.

Regarding notations, I use lowercase letters for random variables, such as x_t and y_t . Realizations of random variables are expressed as x_1 , x_2 , and so on. The context will make it clear whether I am referring to a random variable or its realizations. Capital letters are reserved for matrices, such as A and B . Vectors and matrices are sometimes written in bold for emphasizing, such as \mathbf{X} and \mathbf{y} ; but mostly, in plain format, X and y , provided that they will not lead to confusion. Greek letters are preferred for parameters, such as α and β . Estimators are indicated with a hat, such as $\hat{\alpha}$ and $\hat{\beta}$.

I use the statistical language R whenever programming is involved. I am aware that there are many time series solutions available in R. To avoid burdening readers with excessive packages, I stick to base R as much as possible with a little help from the *zoo* package.

I would like to emphasize that my knowledge and understanding of the subject are limited, and I acknowledge that there may be mistakes or areas where I could have provided a more accurate explanation. I deeply appreciate any feedback or corrections from readers that could improve the accuracy and clarity of this book.

Part I

The Basics

1 Time Series Data

Raw data

We are not so interested in the raw data without any transformation, as it is hard to read information from it. Take the upper-left panel below as an example. There is an overall upward trend. But we are more interested in: how much does the economy grow this year? Is it better or worse than last year? The answers are not obvious from the raw data. Besides, there are obvious seasonal fluctuations. Usually the first quarter has the lowest value in a whole year, due to holidays that significantly reduce the working days in the first quarter. But this does not necessarily mean the economic condition of the first quarter is worse than others. The seasonality prohibits us from sensibly comparing the values of two consecutive quarters.

Growth rates

The headline GDP growth rates are usually reported by comparing the current quarter with the same quarter from last year: $g = \frac{x_t - x_{t-4}}{x_{t-4}} \times 100$. As mentioned above, due to seasonal patterns, comparing two consecutive quarters does not make sense. The year-on-year growth rate tells us how fast the economy grows. However, it loses the information about absolute levels. That comes with several drawbacks. For instance, it is hard to tell whether the economy recovers to the pre-pandemic level after the shock. Due to the unprecedented impact of the pandemic, the GDP for 2020 is exceptionally low, which renders the growth rate for 2021 artificially high. This is undesirable, because it does not mean the economy in 2021 is actually good. We would like a growth rate that shirks off the excessive influence of past observations.

That's why we sometimes prefer (annualized) quarterly growth rate: $g = \frac{x_t - x_{t-1}}{x_{t-1}} \times 400$. Due to seasonal patterns, two consecutive quarters are not comparable directly. But, since this pattern is the same every year, it is possible to remove the seasonal fluctuations. This is called *seasonally adjustment*. We will not cover the seasonal adjustment method here. But this is something that can be done. After seasonally adjusting the data, we can calculate the growth rate based on two consecutive values (annualized by multiplying 4). The bottom-right panel shows the seasonally-adjusted quarterly growth. Note that it is no longer biased upward in 2021 as the year-on-year growth rate.

Seasonally-adjusted data

This is usually the data format we would prefer in time series analysis. [FRED](#) reports both seasonally-adjusted and non-seasonally-adjusted series. The method for seasonal adjustment is a science in itself. Popular algorithms include X-13-ARIMA developed by the United States Census Bureau, TRAMO/SEATS developed by the Bank of Spain, and so on.

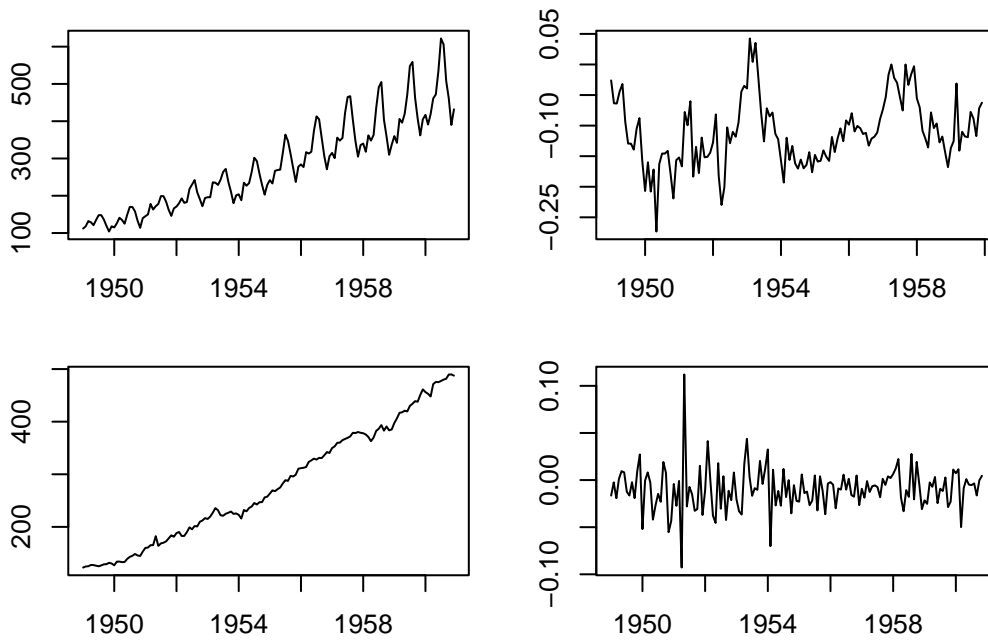
```
par(mfrow=c(2,2), mar=rep(2,4))

# raw data
x <- AirPassengers
plot(x)

# year-over-year growth rate
x_yoy <- x/lag(x,12) - 1
plot(x_yoy)

# seasonally-adjusted series
x_sa <- seasonal::seas(x) |> seasonal::final()
plot(x_sa)

# monthly growth rate
x_mom <- x_sa/lag(x_sa, 1) - 1
plot(x_mom)
```



Logarithms

We like to work with log values. Lots of economic time series exhibit exponential growth, such as GDP. Taking logs convert them to linear. Another amazing fact about logs is that the difference between two log values can be interpreted as percentage growth. To see this, by Taylor's theorem, we have $\ln(\Delta x + 1) \approx \Delta x$ for small values of Δx . Therefore,

$$\ln x_t - \ln x_{t-1} = \ln \left(\frac{x_t}{x_{t-1}} \right) = \ln \left(\frac{x_t - x_{t-1}}{x_{t-1}} + 1 \right) \approx \frac{x_t - x_{t-1}}{x_{t-1}}.$$

So it is very handy to just difference the log levels to get the growth rates. Log-difference can also be interpreted as the continuously compounded rate of change:

$$\frac{x_t}{x_{t-1}} = e^g \implies g = \ln x_t - \ln x_{t-1}.$$

Log-difference also has the property of summability: summing up a sequence of log-differences gives back the log level provided the initial level. It is not as handy if you want to recover the levels from percentage growth.

$$\ln x_t = x_0 + \sum_{j=1}^t (\ln x_j - \ln x_{j-1}).$$

2 Decomposition

2.1 Time Series Components

It is helpful to think about a time series as composed of different components: a trend component, a seasonal component, and a remainder.

$$x_t = T_t + S_t + R_t.$$

The formula assumes the “additive” composition. This assumption is appropriate if the magnitude of the fluctuations does not vary with the absolute levels of the time series. If the magnitude of fluctuations is proportional to the absolute levels, a “multiplicative” decomposition is more appropriate:

$$x_t = T_t \times S_t \times R_t.$$

Note that a multiplicative decomposition of a time series is equivalent to an additive decomposition on its log levels:

$$\ln x_t = \ln T_t + \ln S_t + \ln R_t.$$

Decomposing a time series allows us to extract information that is not obvious from the original time series. It also allows us to manipulate the time series. For example, if the seasonal component can be estimated, we can remove it to obtain seasonally-adjusted series, $x_t^{SA} = x_t - S_t$, or $x_t^{SA} = x_t/S_t$. The question is how to estimate the components given a time series.

2.2 Moving Averages

Moving averages turn out to be handy in estimating trend-cycles by averaging out noisy fluctuations. A moving average of order m (assuming m is an odd number) is defined as

$$\text{MA}(x_t, m) = \frac{1}{m} \sum_{j=-k}^k x_{t+j},$$

where $m = 2k + 1$. For example, a moving average of order 3 is

$$\text{MA}(x_t, 3) = \frac{1}{3}(x_{t-1} + x_t + x_{t+1}).$$

Note that x_t is centered right in the middle and the average is symmetric. This also means, if we apply this formula to real data, the first and last observation will have to be discarded. If the order m is an even number, the formula will no longer be symmetric. To overcome this, we can estimate a moving average over another moving average. For example, we can estimate a moving average of order 4, followed by a moving average of order 2. This is denoted as 2×4 -MA. Mathematically,

$$\begin{aligned} \text{MA}(x_t, 2 \times 4) &= \frac{1}{2}[\text{MA}(x_{t-1}, 4) + \text{MA}(x_t, 4)] \\ &= \frac{1}{2} \left[\frac{1}{4}(x_{t-2} + x_{t-1} + x_t + x_{t+1}) + \frac{1}{4}(x_{t-1} + x_t + x_{t+1} + x_{t+2}) \right] \\ &= \frac{1}{8}x_{t-2} + \frac{1}{4}x_{t-1} + \frac{1}{4}x_t + \frac{1}{4}x_{t+1} + \frac{1}{8}x_{t+2}. \end{aligned}$$

Note that how the 2×4 -MA averages out the seasonality for time series with seasonal period 4, e.g. quarterly series. The formula puts equal weight on every quarter — the first and last terms refer the same quarter and their weights combined to $\frac{1}{4}$.


In general, we can use m -MA to estimate the trend if the seasonal period is an odd number, and use $2 \times m$ -MA if the seasonal period is an even number.

```
library(forecast) # for the `ma` function

x <- AirPassengers

# compute 2x12 MA, since it is a monthly series
x_ma = ma(ma(x, 12), 2)

# plot the original series together with its moving average
plot(x); lines(x_ma, col = "red")
```



2.3 Classical Decomposition


Moving averages give us everything we need to perform classical decomposition. Classical decomposition, invented 1920s, is the simplest method to decompose a time series into trend, seasonality and remainder. It is outdated nowadays and has been replaced by more advanced algorithms. Nonetheless, it serves as a good example for introductory purpose on how time series decomposition could possibly be achieved.

The algorithm for additive decomposition is as follows.

1. Estimate the trend component T_t by applying moving averages. If the seasonal period is an odd number, apply the m -th order MA. If the seasonal period is even, apply the $2 \times m$ MA.
2. Calculate the detrended series $x_t - T_t$.
3. Calculate the seasonal component S_t by averaging all the detrended values of the season. For example, for quarterly series, the value of S_t for Q1 would be the average of all values in Q1. This assumes the seasonal component is constant over time. S_t is then adjusted to ensure all values summed up to zero.
4. Subtracting the seasonal component to get the remainder $R_t = x_t - T_t - S_t$.

```
# classical decomposition
log(AirPassengers) |>
  decompose() |>
  plot()
```

Decomposition of additive time series



The example performs additive decomposition to the logged quarterly GDP series. Note how the constant seasonal component is removed, leaving the smooth and nice-looking up-growing trend. The remainder component tells us the irregular ups and downs of the economy around the trend-cycle. Isn't it amazing that a simple decomposition of the time series tells us a lot about the economy?

2.4 Seasonal Adjustment

By decomposing a time series into trend, seasonality and remainder, it readily gives us a method for seasonal adjustment. Simply subtracting the seasonal component from the original data, or equivalently, summing up the trend and the remainder components, would give us the seasonally-adjusted series.

The following example compares the seasonally-adjusted series using the classical decomposition with the state-of-the-art [X-13ARIMA-SEATS](#) algorithm. We can see the series based on classical decomposition is more volatile, suggesting the classical decomposition is less robust to unusual values.

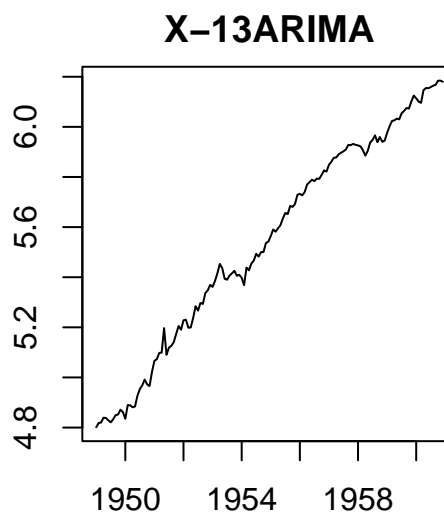


```

par(mfrow=c(1,2), mar=rep(2,4))

# seasonal adjustment based on classical decomposition
x <- log(AirPassengers)
x_comp <- decompose(x)
x_sa_classic = x - x_comp$seasonal
plot(x_sa_classic, main = "Classical SA")

# seasonal adjustment with X-13ARIMA-SEATS
x_seas <- seasonal::seas(x)
x_sa_x13 <- seasonal::final(x_seas)
plot(x_sa_x13, main = "X-13ARIMA")

```



3 ACF and PACF

A time series is notationally represented by $\{\dots, y_{t-1}, y_t, y_{t+1}, y_{t+2}, \dots\}$, which is a sequence of random variables. We think of each variable at a time point t as a random variable, whose realized value is drawn from some distribution.

A distinguishing feature of this sequence is temporal dependence. That is, the distribution of y_t conditional on previous value of the series depends on the outcome of those previous observations. It is of particular interest how observations are correlated across time. A big part of the time series analysis is to exploit this correlation.

3.1 Autocorrelation

The temporal dependence is characterized by the correlation between y_t and its own lags y_{t-k} .

Definition 3.1. The k -th order autocovariance of y_t is defined as

$$\gamma_k = \text{cov}(y_t, y_{t-k}).$$

The k -th order autocorrelation is defined as

$$\rho_k = \frac{\text{cov}(y_t, y_{t-k})}{\text{var}(y_t)} = \frac{\gamma_k}{\gamma_0}.$$

If we plot the autocorrelation as a function of the lag length k , we get the autocorrelation function (ACF). Here is an example of the ACF of China's monthly export growth (log-difference). The lag on the horizontal axis is counted by seasonal period. Because it is monthly data, 1 period is 12 months. We can see the autocorrelation is the strongest for the first two lags. Longer lags are barely significant. There are spikes with 12-month and 24-month lags, indicating the seasonality is not fully removed from the series.


```

par(mfrow=c(1,2))

# compute the daily return of DAX index
x <- EuStockMarkets[, 'DAX']
y <- x/lag(x,1)-1
plot(y)

# compute the ACF among daily returns
acf(y)

```



3.2 Partial Autocorrelation

ACF measures the correlation between y_t and y_{t-k} regardless of their relationships with the intermediate variables $y_{t-1}, y_{t-2}, \dots, y_{t-k+1}$. Even if y_t is only correlated with the first-order lag, it is automatically made correlated with the k -th order lag through intermediate variables. Sometime we are interested in the correlation between y_t and y_{t-k} partialling out the influence of intermediate variables.

Definition 3.2. The partial autocorrelation function (PACF) considers the correlation between the remaining parts in y_t and y_{t-k} after partialling out the intermediate effect of $y_{t-1}, y_{t-2}, \dots, y_{t-k+1}$.

$$\phi_k = \begin{cases} \text{corr}(y_t, y_{t-1}) = \rho_1, & \text{if } k = 1; \\ \text{corr}(r_{y_t|y_{t-1}, \dots, y_{t-k+1}}, r_{y_{t-k}|y_{t-1}, \dots, y_{t-k+1}}), & \text{if } k \geq 2; \end{cases}$$


where $r_{y|x}$ means the remainder in y after partialling out the intermediate effect of x .

In practice, ϕ_k can be estimated by the regression

$$y_t = \mu + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_k y_{t-k} + \epsilon_t.$$

The estimated coefficient $\hat{\phi}_k$ is the partial autocorrelation after controlling the intermediate lags.

```
par(mfrow=c(1,2))  
acf(y)  
pacf(y)
```



4 Stationarity

4.1 Stationary Process

Definition 4.1. A stochastic process is said to be **strictly stationary** if its properties are unaffected by a change of time origin. In other words, the joint distribution at any set of time is not affected by an arbitrary shift along the time axis.

Definition 4.2. A stochastic process is called **covariance stationary** (or **weak stationary**) if its means, variances, and covariances are independent of time. Formally, a process $\{y_t\}$ is covariance stationary if for all t it holds that

- $\mathbb{E}(y_t) = \mu < \infty$;
- $\text{var}(y_t) = \gamma_0 < \infty$;
- $\text{cov}(y_t, y_{t-k}) = \gamma_k$, for $k = 1, 2, 3, \dots$

Stationarity is an important concept in time series analysis. It basically says the statistical properties of a time series are stable over time. Otherwise, if the statistical properties vary with time, statistics estimated from past values, such as autocorrelations, would be much less meaningful. Strict stationarity requires the joint distribution being stable, that is moments of any order would be stable over time. In practice, mostly we only care about the first- and second-order moments, that is means and variances and covariances. Therefore, covariance stationary is sufficient.

Figure 4.1 shows some examples of stationary and non-stationary time series. Only the first one is stationary (it is generated from *i.i.d* normal distribution). The second one is not stationary as its mean is not constant over time. The third one is not stationary as its variance is not constant. The last one is not stationary either, because its covariance is not constant.

Real-life time series are rarely stationary. But they can be transformed to (quasi) stationary by differencing. Figure 4.2 shows some examples of the first-order (log) differences of real-life time series. They more or less exhibit some properties of stationarity, but not perfectly stationary. The series can be further “stationarized” by taking a second-order difference. But these examples are acceptable to be treated as stationary in our models. Even if they are not perfectly stationary, the model can be thought of being used to “extract” their stationary properties.




Figure 4.1: Stationary and non-stationary time series

Percentage change




Figure 4.2: Stationary and non-stationary time series (real data)

Proposition 4.1. *For stationary series, it holds that $\gamma_k = \gamma_{-k}$.*

Proof. By definition,

$$\gamma_k = \mathbb{E}[(y_t - \mu)(y_{t-k} - \mu)],$$

$$\gamma_{-k} = \mathbb{E}[(y_t - \mu)(y_{t+k} - \mu)].$$

Since y_t is stationary, γ_k is invariant with time. Let $t' = t + k$, we have

$$\begin{aligned} \gamma_k &= \mathbb{E}[(y_{t'} - \mu)(y_{t'-k} - \mu)] \\ &= \mathbb{E}[(y_{t+k} - \mu)(y_t - \mu)] \\ &= \gamma_{-k}. \end{aligned}$$

□

4.2 Ergodicity

Temporal dependence is an important feature of time series variables. This dependence is both a blessing and a curse. Autocorrelation enables us to make predictions based on past experiences. However, as we will see in later chapters, it also invalidates theorems that usually require *iid* assumptions. Ideally, we would like the temporal dependence to be not too strong. This is the property of ergodicity.

Definition 4.3. A stationary process $\{y_t\}$ is **ergodic** if

$$\lim_{n \rightarrow \infty} |\mathbb{E}[f(y_t \dots y_{t+k})g(y_{t+n} \dots y_{t+n+k})]| = |\mathbb{E}[f(y_t \dots y_{t+k})]| |\mathbb{E}[g(y_{t+n} \dots y_{t+n+k})]|.$$

Heuristically, ergodicity means if two random variables are positioned far enough in the sequence, they become almost independent. In other words, ergodicity is a restriction on dependency. An ergodic process allows serial correlation, but the serial correlation disappears if the two observations are far apart. Ergodicity is important because as we will see in later chapters, the Law of Large Numbers or the Central Limit Theorem will not hold without it.

Theorem 4.1. *A stationary time series is ergodic if $\sum_{k=0}^{\infty} |\gamma_k| < \infty$.*

Proof. A rigorous proof is not necessary. It is enough to give an intuition why autocorrelation disappears for far apart variables. Note that $\sum_{k=0}^{\infty} |\gamma_k|$ is monotonic and increasing, it converges. Therefore, $\gamma_k \rightarrow 0$ by Cauchy Criterion. □

4.3 White Noise

White noise is a special stationary process that is an important building block of many time series models.

Definition 4.4. A stochastic process w_t is called **white noise** if it has constant mean 0 and variance σ^2 and no serial correlation $\text{cov}(w_t, w_{t-k}) = 0$ for any $k \neq 0$. The white noise process is denoted as

$$w_t \sim \text{WN}(0, \sigma^2).$$

This is the weakest requirement for white noise. It only requires no serial correlation. We may impose further assumptions. If every w_t is independent, it becomes independent white noise $w_t \sim \perp \text{WN}(0, \sigma^2)$. Independence does not imply identical distribution. If every w_t is independently and identically distributed, it is called *i.i.d* white noise, $w_t \stackrel{iid}{\sim} \text{WN}(0, \sigma^2)$. If the distribution is normal, it becomes the most perfect white noise, that is *i.i.d* Gaussian white noise, $w_t \stackrel{iid}{\sim} N(0, \sigma^2)$. The first plot of Figure 4.1 is a demonstration of the *i.i.d* Gaussian white noise. In most cases, the weakest form of white noise is sufficient.

Exercise

Prove that a white noise process is stationary.

Part II

ARIMA Model

5 Model vs Spec

5.1 Classification

Time series models can be broadly sorted into four categories based on whether we are dealing with stationary or non-stationary time series, or whether the model involves only one variable or multiple variables.

Table 5.1: Time series model classification

	Stationary	Nonstationary
Univariate	ARMA	Unit root
Multivariate	VAR	Cointegration

5.2 Model vs Spec

We use the word “model” rather loosely in economics and econometrics. Anything that deals with the quantified relationships between variables can be called a model. A general equilibrium model is a model. A regression is also a model.

To make things less confusing, we would use the word “model” more restrictively in this chapter. We reserve the word **model** to those representing the **data generating processes** (DGPs). That is, when we write down a model in an equation, we literally mean it. If we say y_t follows an AR(1) model:

$$y_t = \phi y_{t-1} + \epsilon_t,$$
$$\epsilon_t \sim N(0, \sigma^2).$$

We literally mean y_t is determined by its previous value and an contemporary innovation drawn from a Gaussian distribution.

A model is distinguished from a **specification**. Suppose $\{y_t\}$ represent the GDP series, we can estimate a regression:

$$y_t = \phi y_{t-1} + e_t$$

This is a specification not a model. Because the DGP of GDP data is unknown, definitely not an AR(1). We can nonetheless fit this spec with the data and get an estimated $\hat{\phi}$. If e_t satisfies some nice properties, for example, uncorrelated with the regressor, then we know this $\hat{\phi}$ is consistent.

When we run regressions with real-life data, we are actually working with specifications. They are not the DGPs of the random variables. But they allow us to recover some useful information from the data when certain assumptions are met. Mostly we are interested in the relationships between variables. A specification describes this relationship, even though it does not describe the full DGP.

This chapter deals with models in the abstract sense. The next chapter will discuss how to fit a model or a spec with real data.

6 AR Models

6.1 AR(1) Process

We start with the simplest time series model — autoregressive model, or AR model. The simplest form of AR model is AR(1), which involves only one lag,

$$y_t = \mu + \phi y_{t-1} + \epsilon_t, \quad (6.1)$$

where $\epsilon_t \sim \text{WN}(0, \sigma^2)$. The model can be extended to include more lags. An AR(p) model is defined as

$$y_t = \mu + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \epsilon_t.$$

We focus on AR(1) first. The model states that the value of y_t is determined by a constant, its previous value, and a random innovation. We call the last term ϵ_t *innovation*, not an error term. It is not an error, it is a random contribution that is unknown until time t . It should also not be confused with the so-called “structural shock”, which is attached with a structural meaning and will be discussed in later chapters.

The model is *probabilistic*, as oppose to *deterministic*, in the sense that some information is unknown or deliberately omitted, so that we do not know the deterministic outcome, but only a probability distribution.

Note

Think about tossing a coin: if every piece of information is incorporated in the model, including the initial speed and position, the air resistance, and so on; then we can figure out the exact outcome, whether the coin will land on its head or tail. But this is unrealistic. Omitting all these information, we can model the process as a Bernoulli distribution. The probability model will not give a deterministic outcome, but only a distribution with each possible value associated with a probability.

i Note

The assumption that a process is only determined by its past values and a white noise innovation seems very restrictive. But it is not. Think about the three assumptions for technical analysis of the stock market (there are still many investors believing this): (1) The market discounts everything, (2) prices move in trends and counter-trends, and (3) price action is repetitive, with certain patterns reoccurring. Effectively, it is saying we can predict the stock market by the past price patterns. If we were to write a model for the stock market based on these assumptions, AR(p) isn't a bad choice at all.

Note that the model can be rewritten as

$$y_t - \frac{\mu}{1-\phi} = \phi \left(y_{t-1} - \frac{\mu}{1-\phi} \right) + \epsilon_t,$$

assuming $\phi \neq 1$. If we define $\tilde{y}_t = y_t - \frac{\mu}{1-\phi}$, we can get rid of the constant term:

$$\tilde{y}_t = \phi \tilde{y}_{t-1} + \epsilon_t. \quad (6.2)$$

It can be easily shown, if y_t is stationary, $\frac{\mu}{1-\phi}$ is the stationary mean. Because this mechanical transformation can always be done to remove the constant. We can simply ignore the constant term without loss of generality.

i Note

Working with demeaned variables greatly simplify the notation. For example, assuming $\mathbb{E}(y_t) = 0$, the variance is simply the second-order moment $\mathbb{E}(y_t^2)$; the covariance can be written as $\mathbb{E}(y_t y_{t-k})$.

For a constant-free AR(1) model, we can rewrite the model as follows:

$$\begin{aligned}
y_t &= \phi y_{t-1} + \epsilon_t \\
&= \phi(\phi y_{t-2} + \epsilon_{t-1}) + \epsilon_t \\
&= \phi^2 y_{t-2} + \phi \epsilon_{t-1} + \epsilon_t \\
&= \phi^2(\phi y_{t-3} + \epsilon_{t-2}) + \phi \epsilon_{t-1} + \epsilon_t \\
&= \phi^3 y_{t-3} + \phi^2 \epsilon_{t-2} + \phi \epsilon_{t-1} + \epsilon_t \\
&\vdots \\
&= \phi^t y_0 + \sum_{j=0}^{t-1} \phi^j \epsilon_{t-j} \\
&= \sum_{j=0}^{\infty} \phi^j \epsilon_{t-j}.
\end{aligned} \tag{6.3}$$

The exercise shows an AR(1) process can be reduced to an MA process, which will be discussed in the next section. It says the value of y_t is determined by its initial value (if it has one) and the accumulated innovations in the past. It is our deeds in history that shapes our world today.

Note

The property that an AR process can be rewritten as an infinite MA process with absolute summable coefficients $\sum_{j=0}^{\infty} |\phi^j| < \infty$ is called *causal*. This must not be confused with the causal effect in econometrics (defines in the *ceteris paribus* sense). To avoid confusion, we avoid use this term as much as possible.

Now we focus our attention on the critical parameter ϕ . If $|\phi| > 1$, the process is explosive. We are not interested in explosive processes. If a real-world time series grows exponentially, we take logarithm to transform it to linear. So in most of our discussions, we rule out the case of explosive behaviour.

If $|\phi| < 1$, $\phi^j \rightarrow 0$ as $j \rightarrow \infty$. This means the influence of innovations far away in the past decays to zero. We will show that the series is stationary and ergodic.

If $|\phi| = 1$, we have $y_t = \sum_{j=0}^{\infty} \text{sgn}(\phi)^j \epsilon_{t-j} = \sum_{j=0}^{\infty} \tilde{\epsilon}_{t-j}$. This means the influence of past innovations will not decay no matter how distant away they are. This is known as a *unit root process*, which will be covered in later chapters. But it is clear that the process is not stationary. Consider the variance of y_t conditioned on an initial value:

$$\text{var}(y_t|y_0) = \text{var}\left(\sum_{j=0}^{t-1} \epsilon_{t-j}\right) = \sum_{j=0}^{t-1} \text{var}(\epsilon_{t-j}) = \sum_{j=0}^{t-1} \sigma^2 = t\sigma^2.$$

The variance is increasing with time. It is not constant. Figure 6.1 simulates the AR(1) with $\phi = 0.5$ and $\phi = 1$ respectively.

```
y = arima.sim(list(ar=0.5), n=1000)
z = arima.sim(list(order=c(0,1,0)), n=1000)
plot(cbind(y,z), plot.type="multiple", nc=2, ann=F,
     mar.multi=rep(2,4), oma.multi = rep(0,4))
```




Figure 6.1: Simulation of AR(1) processes

Proposition 6.1. *An AR(1) process with $|\phi| < 1$ is covariance stationary.*

Proof. Let's compute the mean, variance and covariance for the AR(1) process.

$$\mathbb{E}(y_t) = \mathbb{E} \left[\sum_{j=0}^{\infty} \phi^j \epsilon_{t-j} \right] = \sum_{j=0}^{\infty} \phi^j \mathbb{E}[\epsilon_{t-j}] = 0.$$

$$\begin{aligned} \text{var}(y_t) &= \text{var} \left[\sum_{j=0}^{\infty} \phi^j \epsilon_{t-j} \right] = \sum_{j=0}^{\infty} \phi^{2j} \text{var}[\epsilon_{t-j}] \\ &= \sigma^2 \sum_{j=0}^{\infty} \phi^{2j} = \frac{\sigma^2}{1 - \phi^2}. \end{aligned}$$

For the covariances,

$$\begin{aligned}
\gamma_1 &= \mathbb{E}(y_t y_{t-1}) = \mathbb{E}((\phi y_{t-1} + \epsilon_t) y_{t-1}) \\
&= \mathbb{E}(\phi y_{t-1}^2 + \epsilon_t y_{t-1}) \\
&= \phi \mathbb{E}(y_{t-1}^2) + 0 \\
&= \frac{\phi \sigma^2}{1 - \phi^2}; \\
\gamma_2 &= \mathbb{E}(y_t y_{t-2}) = \mathbb{E}((\phi y_{t-1} + \epsilon_t) y_{t-2}) \\
&= \mathbb{E}(\phi y_{t-1} y_{t-2} + \epsilon_t y_{t-2}) \\
&= \phi \mathbb{E}(y_{t-1} y_{t-2}) \\
&= \phi \gamma_1 = \frac{\phi^2 \sigma^2}{1 - \phi^2}; \\
&\vdots \\
\gamma_j &= \frac{\phi^j \sigma^2}{1 - \phi^2}.
\end{aligned}$$

All of them are independent of time t . By Definition 4.2, the process is covariance stationary. \square

So the ACF decays gradually as $\phi^j \rightarrow 0$. What about the PACF? Estimating the PACF is equivalent to regressing y_t on its lags. Since there is only one lag, the PACF should have non-zero value only for the first lag, and zeros for all other lags.

```
par(mfrow=c(1,2), mar=c(2,4,1,1))
acf(y); pacf(y)
```




Figure 6.2: ACF and PACF for AR(1) process

6.2 Lag Operator

To facilitate easy manipulation of lags, we introduce the lag operator:

$$Ly_t = y_{t-1}.$$

The AR(1) process can be written with the lag operator:

$$y_t = \phi Ly_t + \epsilon_t \implies (1 - \phi L)y_t = \epsilon_t.$$

The lag operator L can be manipulated just as polynomials. It looks weird, but it actually works. Do a few exercises to convince yourself.

$$L^2 y_t = L(Ly_t) = Ly_{t-1} = y_{t-2}.$$

$$\begin{aligned}(1 - L)^2 y_t &= (1 - L)(y_t - y_{t-1}) \\ &= (y_t - y_{t-1}) - (y_{t-1} - y_{t-2}) \\ &= y_t - 2y_{t-1} + y_{t-2} \\ &= (1 - 2L + L^2)y_t.\end{aligned}$$

We can even inverse a lag polynomial (provided $|\phi| < 1$),

$$\begin{aligned}(1 - \phi L)y_t &= \epsilon_t \\ \implies y_t &= (1 - \phi L)^{-1} \epsilon_t = \sum_{j=0}^{\infty} \phi^j L^j \epsilon_t = \sum_{j=0}^{\infty} \phi^j \epsilon_{t-j}.\end{aligned}$$

We reach the same conclusion as Equation 6.3 with the lag operator.

6.3 AR(p) Process

We now generalize the conclusions above to AR(p) processes. With the help of the lag operator, an AR(p) process can be written as

$$(1 - \phi_1 L - \phi_2 L^2 - \cdots - \phi_p L^p)y_t = \epsilon_t,$$

or even more parsimoniously,

$$\phi(L)y_t = \epsilon_t.$$

Note that we ignore the constant term, which can always be removed by redefine $\tilde{y}_t = y_t - \frac{\mu}{1-\phi_1-\phi_2-\dots-\phi_p}$.

To derive the MA representation, we need to figure out $\phi^{-1}(L)$. By the Fundamental Theorem of Algebra, we know the polynomial $\phi(z)$ has p roots in the complex space. So the lag polynomial can be factored as

$$(1 - \lambda_1 L)(1 - \lambda_2 L) \dots (1 - \lambda_p L)y_t = \epsilon_t,$$

where $z = \lambda_i^{-1}$ is the i -th root of $\phi(z)$. If the roots are outside the unit circle, $|\lambda_i| < 1$ means each of the left hand terms is invertible.

$$\begin{aligned} y_t &= \frac{1}{(1 - \lambda_1 L)(1 - \lambda_2 L) \dots (1 - \lambda_p L)} \epsilon_t \\ &= \left(\frac{c_1}{1 - \lambda_1 L} + \frac{c_2}{1 - \lambda_2 L} + \dots + \frac{c_p}{1 - \lambda_p L} \right) \epsilon_t \\ &= \sum_{j=0}^{\infty} (c_1 \lambda_1^j + c_2 \lambda_2^j + \dots + c_p \lambda_p^j) L^j \epsilon_t \\ &= \sum_{j=0}^{\infty} \theta_j \epsilon_{t-j}, \text{ where } \theta_j = c_1 \lambda_1^j + \dots + c_p \lambda_p^j. \end{aligned}$$

It follows that this process has constant mean and variance. For the covariances, given

$$y_t = \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + \epsilon_t,$$

Multiply both sides by y_t and take expectation,

$$\mathbb{E}[y_t^2] = \phi_1 \mathbb{E}[y_t y_{t-1}] + \phi_2 \mathbb{E}[y_t y_{t-2}] + \dots + \phi_p \mathbb{E}[y_t y_{t-p}],$$

$$\gamma_0 = \phi_1 \gamma_{-1} + \phi_2 \gamma_{-2} + \dots + \phi_p \gamma_{-p}.$$

Similarly, multiply both sides by y_{t-1}, \dots, y_{t-j} , we have

$$\begin{aligned}\gamma_1 &= \phi_1\gamma_0 + \phi_2\gamma_{-1} + \cdots + \phi_p\gamma_{-p+1}, \\ &\vdots \\ \gamma_j &= \phi_1\gamma_{j-1} + \phi_2\gamma_{j-2} + \cdots + \phi_p\gamma_{j-p}.\end{aligned}$$

This is called the Yule-Walker equation. The first p unknowns $\gamma_0, \dots, \gamma_{p-1}$ can be solved by the first p equations. The rest can then be solved iteratively.

It can be shown all of the covariances are invariant with time. Therefore, under the condition all $|\lambda_i| < 1$, the $AR(p)$ process is stationary.

For the PACF, a regression of y_t over its lags would recover p non-zero coefficients. Longer lags should have coefficients insignificantly different from zero.

```
y = arima.sim(list(ar=c(2.4, -1.91, 0.5)), n=3000)
par(mfrow=c(1,2), mar=c(2,4,1,1))
acf(y); pacf(y)
```




Figure 6.3: ACF and PACF for $AR(p)$ process

Proposition 6.2. *An $AR(p)$ process is stationary if all the roots of $\phi(z)$ are outside the unit circle.*

Proposition 6.3. *An $AR(p)$ process is characterized by (i) an ACF that is infinite in extend but tails off gradually; and (ii) a PACF that is (close to) zero for lags after p .*

7 MA Models

7.1 MA(1) Process

Again, let's start with the simplest moving average model. A first-order moving average process, or MA(1), is defined as

$$y_t = \mu + \epsilon_t + \theta\epsilon_{t-1}, \quad (7.1)$$

where $\{\epsilon_t\} \sim \text{WN}(0, \sigma^2)$ are uncorrelated innovations. The MA model says the current value y_t is a moving average of past innovations (in MA(1), the weight on ϵ_{t-1} is θ). MA models directly relate the observable variable to past innovations. If we know the past innovation ϵ_{t-j} , we can easily figure out its contribution to the outcome variable (unlike AR models where the effect of a past innovation is transmitted through y_{t-j}, \dots, y_{t-1}). So MA models are the preferred analytic tool in many applications, despite it looks odd from the eyes of regression modelers. You may wonder how it is possible to estimate such a model. We will put off the estimation techniques to the next chapter.

It is clear that y_t has a constant mean, $\mathbb{E}(y_t) = \mu$. We can omit the constant if we work with the demeaned series $\tilde{y}_t = y_t - \mu$. Without loss of generality, we assume for the rest $\{y_t\}$ has zero mean, so the model is simplified as

$$y_t = \epsilon_t + \theta\epsilon_{t-1}. \quad (7.2)$$

Let's compute its variance and covariances:

$$\begin{aligned} \gamma_0 &= \text{var}(\epsilon_t + \theta\epsilon_{t-1}) = \text{var}(\epsilon_t) + \theta^2\text{var}(\epsilon_{t-1}) = (1 + \theta^2)\sigma^2; \\ \gamma_1 &= \text{cov}(y_t, y_{t-1}) = \text{cov}(\epsilon_t + \theta\epsilon_{t-1}, \epsilon_{t-1} + \theta\epsilon_{t-2}) = \text{cov}(\theta\epsilon_{t-1}, \epsilon_{t-1} + \theta\epsilon_{t-2}) = \theta\sigma^2; \\ \gamma_2 &= \text{cov}(y_t, y_{t-2}) = \text{cov}(\epsilon_t + \theta\epsilon_{t-1}, \epsilon_{t-2} + \theta\epsilon_{t-3}) = 0; \\ &\vdots \\ \gamma_j &= 0 \text{ for } |j| \geq 2. \end{aligned}$$

It is clear that the MA(1) process is *stationary*. And the ACF cuts off after the first lag. Because more distant lags y_{t-k} are constituted by even more distant innovations $\epsilon_{t-k}, \epsilon_{t-k-1}, \dots$ which has no relevance for y_t given the MA(1) structure.

We have seen AR processes are equivalent to MA(∞) processes. Similar results hold for MA models. Rewrite the MA(1) process with the lag operator, assuming $|\theta| < 1$,

$$y_t = (1 + \theta L)\epsilon_t \Leftrightarrow (1 + \theta L)^{-1}y_t = \epsilon_t \Leftrightarrow \sum_{j=0}^{\infty} (-\theta)^j y_{t-j} = \epsilon_t.$$

That means an MA(1) is equivalent to an AR(∞) process if $(1 + \theta L)$ is *invertible*. This shows AR and MA are really the same family of models. The model AR or MA is chosen by parsimonious principle. For example, an AR model with many lags can possibly be modeled by a parsimonious MA model.

Since an MA(1) is equivalent to some AR(∞) process, the PACF of an MA(1) should tail off gradually.

```
y = arima.sim(list(ma=0.8), n=2000)
par(mfrow=c(1,2), mar=c(1,4,1,1))
acf(y); pacf(y)
```




Figure 7.1: ACF and PACF of MA(1) process

i Invertibility

If $|\theta| > 1$, $\theta(L)$ is not invertible. Define another MA(1) process,

$$y_t = \epsilon_t + \theta^{-1}\epsilon_{t-1}, \quad \epsilon_t \sim \text{WN}(0, \theta^2\sigma^2).$$

We can verify that its variance and covariances are exactly the same as Equation 7.2. For non-invertible MA process, as long as $\theta(L)$ avoids unit root, we can always find an invertible process that shares the same ACF. This means, for a stationary MA process, it makes no harm to just assume it is invertible.

7.2 MA(q) Process

A q -th order moving average, or MA(q) process, is written as

$$y_t = \mu + \epsilon_t + \theta_1 \epsilon_{t-1} + \cdots + \theta_q \epsilon_{t-q}, \quad (7.3)$$

where $\{\epsilon_t\} \sim \text{WN}(0, \sigma^2)$.

Proposition 7.1. *An MA(q) process is stationary.*

Proof. We will show that the mean, variance and covariances of MA(q) are all invariant with time.

$$\mathbb{E}(y_t) = \mu.$$

Assume for the rest, $\{y_t\}$ is demeaned.

$$\begin{aligned} \gamma_0 &= \mathbb{E}(y_t^2) = \mathbb{E}[(\epsilon_t + \theta_1 \epsilon_{t-1} + \cdots + \theta_q \epsilon_{t-q})^2] \\ &= \mathbb{E}[\epsilon_t^2] + \theta_1^2 \mathbb{E}[\epsilon_{t-1}^2] + \cdots + \theta_q^2 \mathbb{E}[\epsilon_{t-q}^2] \\ &= (1 + \theta_1^2 + \cdots + \theta_q^2) \sigma^2; \\ \gamma_1 &= \mathbb{E}[y_t y_{t-1}] = \mathbb{E}[(\epsilon_t + \theta_1 \epsilon_{t-1} + \cdots + \theta_q \epsilon_{t-q}) \\ &\quad (\epsilon_{t-1} + \cdots + \theta_{q-1} \epsilon_{t-q} + \theta_q \epsilon_{t-q-1})] \\ &= \theta_1 \mathbb{E}[\epsilon_{t-1}^2] + \theta_2 \theta_1 \mathbb{E}[\epsilon_{t-2}^2] + \cdots + \theta_q \theta_{q-1} \mathbb{E}[\epsilon_{t-q}^2] \\ &= (\theta_1 + \theta_2 \theta_1 + \cdots + \theta_q \theta_{q-1}) \sigma^2; \\ &\vdots \\ \gamma_j &= \mathbb{E}[y_t y_{t-j}] = \mathbb{E}[(\epsilon_t + \cdots + \theta_j \epsilon_{t-j} + \cdots + \theta_q \epsilon_{t-q}) \\ &\quad (\epsilon_{t-j} + \cdots + \theta_{q-j} \epsilon_{t-q} + \cdots + \theta_q \epsilon_{t-q-j})] \\ &= \theta_j \mathbb{E}[\epsilon_{t-j}^2] + \theta_{j+1} \theta_j \mathbb{E}[\epsilon_{t-j-1}^2] + \cdots + \theta_q \theta_{q-j} \mathbb{E}[\epsilon_{t-q}^2] \\ &= (\theta_j + \theta_{j+1} \theta_j + \cdots + \theta_q \theta_{q-j}) \sigma^2, \text{ for } j \leq q; \\ \gamma_j &= 0, \text{ for } j > q. \end{aligned}$$

□

Proposition 7.2. *An $MA(q)$ process is invertible iff the roots of $\theta(z)$ are outside the unit circle.*

Proposition 7.3. *An $MA(q)$ process is characterized by (i) an ACF that is (close to) zero after q lags; and (ii) a PACF that is infinite in extend but tails off gradually.*

7.3 $MA(\infty)$ Process

$MA(\infty)$ is a special case deserves attention. Partly because all ARMA processes can be reduced to $MA(\infty)$ processes. In addition to $MA(q)$ processes, we need more conditions for $MA(\infty)$ to be stationary. Consider the variance of

$$y_t = \sum_{j=0}^{\infty} \theta_j \epsilon_{t-j},$$

$$\gamma_0 = \mathbb{E}[y_t^2] = \mathbb{E} \left[\left(\sum_{j=0}^{\infty} \theta_j \epsilon_{t-j} \right)^2 \right] = \left(\sum_{j=0}^{\infty} \theta_j^2 \right) \sigma^2.$$

It only make sense if $\sum_{j=0}^{\infty} \theta_j^2 < \infty$. This property is called *square summable*.

Proposition 7.4. *An $MA(\infty)$ process is stationary if the coefficients $\{\theta_j\}$ are square summable.*

8 ARMA Models

8.1 ARMA(p,q)

ARMA(p, q) is a mixed autoregressive and moving average process.

$$y_t = \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \epsilon_t + \theta_1 \epsilon_{t-1} + \cdots + \theta_q \epsilon_{t-q},$$

or

$$\phi(L)y_t = \theta(L)\epsilon_t,$$

where $\{\epsilon_t\} \sim \text{WN}(0, \sigma^2)$.

The MA part is always stationary as shown in Proposition 7.1. The stationarity of an ARMA process solely depends on the AR part. The condition is the same as Proposition 6.2.

Assume $\phi^{-1}(L)$ exist, then the ARMA(p, q) process can be reduce to MA(∞) process:

$$y_t = \phi^{-1}(L)\theta(L)\epsilon_t = \psi(L)\epsilon_t,$$

where $\psi(L) = \phi^{-1}(L)\theta(L)$.

Exercise

Compute the MA equivalence for ARMA(1,1).

8.2 ARIMA(p,d,q)

ARMA(p,q) is used to model stationary time series. If y_t is not stationary, we can transform it to stationary and model it with an ARMA model. If the first-order difference $(1-L)y_t = y_t - y_{t-1}$ is stationary, then we say y_t is **integrated** of order 1. If it requires d -th order difference to be stationary, $(1-L)^d y_t$, we say it is integrated of order d . The ARMA model involves integrated time series is called ARIMA model:

$$\phi(L)(1-L)^d y_t = \theta(L)\epsilon_t.$$

9 Wold Theorem

9.1 Wold Decomposition

So far we have spent a lot of effort with ARMA models, which are the indispensable components of any time series textbook. The following theorem justifies its importance. The Wold Decomposition Theorem basically says every covariance-stationary process has an ARMA representation. Therefore, with long enough lags, any covariance-stationary process can be approximated arbitrarily well by ARMA models. This is a very bold conclusion to make. It sets up the generality of ARMA models, which makes it one of the most important theorems in time series analysis.

Theorem 9.1 (Wold Decomposition Theorem). *Every covariance-stationary time series y_t can be written as the sum of two time series, one deterministic and one stochastic. Formally,*

$$y_t = \eta_t + \sum_{j=0}^{\infty} b_j \epsilon_{t-j},$$

where $\eta_t \in I_{-\infty}$ is a deterministic time series (such as one represented by a sine wave); ϵ_t is an uncorrelated innovation sequence with $\mathbb{E}[\epsilon_t] = 0$, $\mathbb{E}[\epsilon_t \epsilon_{t-j}] = 0$ for $j \neq 0$; and $\{b_j\}$ are square summable, $\sum_{j=0}^{\infty} |b_j|^2 < \infty$.

Proof. We will prove the theorem by constructing the innovation sequence $\{e_t\}$ and showing it satisfies the conditions stated. Let $e_t = y_t - \hat{\mathbb{E}}(y_t | I_{t-1}) = y_t - a(L)y_{t-1}$, where $\hat{\mathbb{E}}(y_t | I_{t-1})$ is the best linear predictor (BLP) of y_t based on information set at $t-1$. $a(L)$ does not depend on t because y_t is covariance stationary. As the best linear predictor, $a(L)$ solves

$$\min_{\{a_j\}} \mathbb{E}(y_t - \sum_{j=1}^{\infty} a_j y_{t-j})^2.$$

The first-order conditions with respect to a_j gives

$$\begin{aligned} \mathbb{E}[y_{t-j}(y_t - \sum_{j=1}^{\infty} a_j y_{t-j})] &= 0, \\ \implies \mathbb{E}[y_{t-j}e_t] &= 0. \end{aligned}$$

We now verify that e_t satisfies the white noise conditions. Without loss of generality, we may assume $\mathbb{E}(y_t) = 0$, it follows that $\mathbb{E}(e_t) = 0$. $\text{var}(e_t) = \mathbb{E}(y_t - a(L)y_t)^2$ is a function of covariance of y_t and a_j , none of which varies with time. So $\text{var}(e_t) = \sigma^2$ is constant. Utilizing the first-order condition, $\mathbb{E}[e_t e_{t-j}] = \mathbb{E}[e_t(y_{t-j} - a(L)y_{t-j})] = 0$.

Repeatedly substituting for y_{t-k} gives

$$\begin{aligned}
y_t &= e_t + \sum_{k=1}^{\infty} a_k y_{t-k} \\
&= e_t + a_1(e_{t-1} + \sum_{k=1}^{\infty} a_k y_{t-1-k}) + \sum_{k=2}^{\infty} a_k y_{t-k} \\
&= e_t + a_1 e_{t-1} + \sum_{k=1}^{\infty} \tilde{a}_k y_{t-k-1} \\
&= e_t + a_1 e_{t-1} + \eta_t^1 \\
&\vdots \\
&= \sum_{j=0}^k c_j e_{t-j} + \eta_t^k,
\end{aligned}$$

where $\eta_t^k \in I_{t-k-1}$. As $k \rightarrow \infty$, we have $v_t = y_t - \sum_{j=0}^{\infty} c_j e_{t-j} \in I_{-\infty}$. □

Let's appreciate this theorem for a while. The property of stationarity can be loosely understood as having stable patterns over time. The Wold Theorem states that any such patterns can be captured by ARMA models. In other words, ARMA models are effective in modelling stable patterns repeated over time, in so far as only 2nd-order moments are of concern. Even if the time series is not entirely stationary, if we model it with ARMA, it can be thought as extracting the stationary patterns. Figure 9.1 demonstrates the ARIMA modelling of real data.

```

x <- BJsales # Sales data

# model with ARIMA(1,1,1)
mod <- arima(x, order = c(1,1,1))

# fitted values
fit <- fitted(mod)

# make plot
plot(x); lines(fit, col = "red")

```




Figure 9.1: Monthly sales modelled with ARIMA(2,0,1)

9.2 Causality and Invertibility*

We have seen that AR models can be rewritten as MA models and vice versa, suggesting the ARMA representation of a stochastic process is not unique. We have also seen that a non-invertible MA process can be equivalently represented by an invertible MA process. For example, the following MA(1) processes have the same ACF:

$$\begin{aligned} x_t &= w_t + \frac{1}{5}w_{t-1}, & w_t &\sim \text{WN}(0, 25); \\ y_t &= v_t + 5v_{t-1}, & v_t &\sim \text{WN}(0, 1). \end{aligned}$$

The same property holds for AR processes. In Chapter 6, we state that an AR(1) process is explosive if $|\phi| > 1$. This is not entirely rigorous. Consider an AR(1) process,

$$y_t = \phi y_{t-1} + \epsilon_t, \text{ where } |\phi| > 1.$$

Multiply both sides by ϕ^{-1} ,

$$\phi^{-1}y_t = y_{t-1} + \phi^{-1}\epsilon_t,$$

Rewrite it as an MA process,

$$\begin{aligned}
y_t &= \phi^{-1}y_{t+1} - \phi^{-1}\epsilon_{t+1} \\
&= \phi^{-1}(\phi^{-1}y_{t+2} - \phi^{-1}\epsilon_{t+2}) - \phi^{-1}\epsilon_{t+1} \\
&\vdots \\
&= \sum_{j=1}^{\infty} -\phi^{-j}\epsilon_{t+j}.
\end{aligned}$$

Given $|\phi^{-1}| < 1$, the process is stationary, expressed as discounted innovations in the future (despite this looks quite odd). In fact, for a non-causal AR process, we can find a causal AR process that generates the same ACF (remember the term *causal* means an AR process can be converted to an MA process with absolute summable coefficients).

The problem is given an ARMA equation, it is not enough to uniquely pin down a stochastic process. Both the explosive process and the stationary process can be a solution to $y_t = \phi y_{t-1} + \epsilon_t$. But for a stationary process expressed as an AR model with $|\phi| > 1$, we can always find an AR(1) process with $|\tilde{\phi}| < 1$ and a different white noise sequence $\{\tilde{\epsilon}_t\}$ that generate the same ACF.

The following theorems state the conditions for the existence of stationary solutions, and the possibility of rewriting non-causal or non-invertible ARMA representations as causal and invertible ones. Since it is always possible to do so, it loses nothing to stick with causal and invertible ARMA processes when modelling stationary time series.

Theorem 9.2. *A unique stationary solution to the ARMA process $\phi(L)y_t = \theta(L)\epsilon_t$ exists iff ϕ and θ have no common factors and the roots of $\phi(z)$ avoid the unit circle:*

$$|\phi(z)| = 1 \implies \phi(z) = 1 - \phi_1 z - \dots - \phi_p z^p \neq 0.$$

Theorem 9.3. *Let $\{y_t\}$ be a stationary ARMA process defined by $\phi(L)y_t = \theta(L)\epsilon_t$. If the roots of $\theta(z)$ avoid unit circle, then there are polynomials $\tilde{\phi}$ and $\tilde{\theta}$ and a white noise sequence $\tilde{\epsilon}$ such that $\{y_t\}$ satisfies $\tilde{\phi}(L)y_t = \tilde{\theta}(L)\tilde{\epsilon}_t$, and this is a causal and invertible ARMA process.*

Part III

Time Series Regression

10 Preliminaries

10.1 Chapter Overview

This chapter serves two purposes. One is to introduce the techniques for estimating time series models. The other is to explain the concept of dynamic causal effect. We join the two topics in one chapter because both of them can be done via a regression framework. Maximum likelihood estimation plays a pivotal role in estimating time series models. Nonetheless, starting with OLS always make things easier. We start with a quick review of the basic OLS concepts that are familiar to any students in econometrics, that is the regressions applied to cross-sectional *iid* observations. We then extend it to time series data. We will see it is not as straightforward as one might expect, as intertemporal dependencies between observation need additional treatment. In the second half of the chapter, we will explain the concept of dynamic causal effect, that is the causal effect of an intervention on outcome variables. Similar to cross-sectional studies, we need to define the causal effect relative to counterfactuals. With time series data, the counterfactuals have to be defined across time rather across individuals.

10.2 Asymptotic Theorems for i.i.d Random Variables

Theorem 10.1 (Law of Large Numbers). *Let $\{x_i\}$ be iid random variables with $\mathbb{E}(x_i) = \mu$ and $\text{Var}(x_i) = \sigma^2 < \infty$. Define $\bar{x}_n = \frac{1}{n} \sum_{i=1}^n x_i$. Then $\bar{x}_n \xrightarrow{p} \mu$ as $n \rightarrow \infty$.*

Proof. We will give an non-rigorous proof, but nonetheless shows the tenets. It is easy to see $\mathbb{E}(\bar{x}_n) = \mu$. Consider the variance,

$$\text{Var}(\bar{x}_n) = \text{Var}\left(\frac{1}{n} \sum_{i=1}^n x_i\right) \stackrel{iid}{=} \frac{1}{n^2} \sum_{i=1}^n \text{Var}(x_i) = \frac{\sigma^2}{n} \rightarrow 0.$$

That is \bar{x}_n converges to μ with probability 1 as $n \rightarrow \infty$. Note that we can move the variance inside the summation operator because x_i are *iid*, in which all the covariance terms are 0. \square

Theorem 10.2 (Central Limit Theorem). *Let $\{x_i\}$ be iid random variables with $\mathbb{E}(x_i) = \mu$ and $\text{Var}(x_i) = \sigma^2 < \infty$. Define $\bar{x}_n = \frac{1}{n} \sum_{i=1}^n x_i$. Then*

$$\frac{\bar{x}_n - \mu}{\sigma/\sqrt{n}} \xrightarrow{d} N(0, 1).$$

Proof. Without loss of generality, assume x_i is demeaned and standardized to have standard deviation 1. It remains to show $\sqrt{n}\bar{x}_n \rightarrow N(0, 1)$. Define the moment generating function (MGF) for $\sqrt{n}\bar{x}_n$:

$$M_{\sqrt{n}\bar{x}_n}(t) = \mathbb{E}[e^{(\sqrt{n}^{-1} \sum_{i=1}^n x_i)t}] \stackrel{iid}{=} \{\mathbb{E}[e^{(n^{-1/2}x_i)t}]\}^n.$$

Evaluate the MGF for each x_i :

$$\mathbb{E}[e^{(n^{-1/2}x_i)t}] = 1 + \mathbb{E}(n^{-1/2}x_i)t + \mathbb{E}(n^{-1}x_i^2)t^2 + \dots = 1 + \frac{t^2}{2n} + o(n^{-1}).$$

Substituting back,

$$M_{\sqrt{n}\bar{x}_n}(t) = \left[1 + \frac{t^2}{2n} + o(n^{-1})\right]^n = \left[\left(1 + \frac{t^2}{2n}\right)^{\frac{2n}{t^2}}\right]^{\frac{t^2}{2}} \rightarrow e^{\frac{t^2}{2}}.$$

Note that we drop the $o(n^{-1})$ because it converges faster than $\frac{1}{n}$. $e^{\frac{t^2}{2}}$ is the MGF for standard normal distribution. Hence, the theorem is proved. \square

10.3 OLS for i.i.d Random Variables

We now give a very quick review of OLS with *iid* random variables. These materials are assumed familiar to the readers. We do not intend to introduce them in any detail. This section is a quick snapshot of some key concepts, so that we could contrast them with the time series regression introduced in the next section.

A linear regression model postulates the joint distribution of (y_i, x_i) follows a linear relationship,

$$y_i = x_i' \beta + \epsilon_i.$$

Expressed in terms of data matrix,

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} x_{11}, x_{12}, \dots, x_{1p} \\ x_{21}, x_{22}, \dots, x_{2p} \\ \ddots \\ x_{n1}, x_{n2}, \dots, x_{np} \end{bmatrix}' \begin{bmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{bmatrix}.$$

From the perspective of dataset, the matrix matrix is fixed in the sense that they are just numbers in the dataset. But for statistical analysis, we view each entry in the matrix as random, that is as a realization of a random process.

To estimate the parameter β from sample data, OLS seeks to minimize the squared residuals

$$\min_{\beta} \sum_{i=1}^n (y_i - x_i' \beta)^2.$$

The first-order condition implies,

$$\begin{aligned} \sum_i x_i (y_i - x_i' \beta) &= 0, \\ \sum_i x_i y_i - \sum_i x_i x_i' \beta &= 0, \\ \hat{\beta} &= \left(\sum_i x_i x_i' \right)^{-1} \left(\sum_i x_i y_i \right) \\ &= \beta + \left(\sum_i x_i x_i' \right)^{-1} \left(\sum_i x_i \epsilon_i \right). \end{aligned}$$

Under the Gauss-Markov assumptions, particularly $\mathbb{E}(\epsilon_i | x_j) = 0$ and $\text{var}(\epsilon | X) = \sigma^2 I$ (homoskedasticity and nonautocorrelation), the OLS estimator is **BLUE** (Best Linear Unbiased Estimator).

Under the assumption of *iid* random variables and homoskedasticity, we invoke the LLN and CLT to derive the asymptotic distribution for the OLS estimator,

$$\begin{aligned} \sqrt{n}(\hat{\beta} - \beta) &= \left(\frac{1}{n} \sum_i x_i x_i' \right)^{-1} \left(\sqrt{n} \frac{1}{n} \sum_i x_i \epsilon_i \right) \\ &\rightarrow [\mathbb{E}(x_i x_i')]^{-1} N(0, \mathbb{E}(x_i \epsilon_i \epsilon_i' x_i')) \\ &\rightarrow N(0, \sigma^2 [\mathbb{E}(x_i x_i')]^{-1}). \end{aligned}$$

Note how the *iid* assumption is required throughout the process. The following section will show how to extend the OLS to non-*iid* random variables and how it leads to modification of the results.

11 OLS for Time Series

11.1 Asymptotic Theorems for Dependent Random Variables

The asymptotic theorems and regressions that work for *iid* random variable do not immediately apply to time series. Consider the proof for Theorem 10.1, without the *iid* assumption we have

$$\begin{aligned}
 \text{var} \left(\frac{1}{n} \sum_{i=1}^n x_i \right) &= \frac{1}{n^2} \sum_{i=1}^n \sum_{j=1}^n \text{cov}(x_i, x_j) \\
 &= \frac{1}{n^2} [\text{cov}(x_1, x_1) + \text{cov}(x_1, x_2) + \cdots + \text{cov}(x_1, x_n) + \\
 &\quad \text{cov}(x_2, x_1) + \text{cov}(x_2, x_2) + \cdots + \text{cov}(x_2, x_n) + \\
 &\quad \vdots \\
 &\quad \text{cov}(x_n, x_1) + \text{cov}(x_n, x_2) + \cdots + \text{cov}(x_n, x_n)] \\
 &= \frac{1}{n^2} [n\gamma_0 + 2(n-1)\gamma_1 + 2(n-2)\gamma_1 + 2(n-2)\gamma_2 + \dots] \\
 &= \frac{1}{n} \left[2 \sum_{k=1}^n \gamma_k \left(1 - \frac{k}{n} \right) + \gamma_0 \right].
 \end{aligned}$$

The argument for the *iid* does not work with the presence of serial correlations. If we assume absolute summability, $\sum_{j=-\infty}^{\infty} |\gamma_j| < \infty$, then

$$\lim_{n \rightarrow \infty} \frac{1}{n} \left[2 \sum_{k=1}^n \gamma_k \left(1 - \frac{k}{n} \right) + \gamma_0 \right] = 0.$$

In this case, we still have the LLN holds. Otherwise, as the variance may not converge. Remember Theorem 4.1, absolute summability implies the series is ergodic.

Proposition 11.1. *If x_t is a covariance stationary time series with absolutely summable auto-covariances, then a Law of Large Numbers holds.*

From the new proof of LLN one can guess that the variance in a Central Limit Theorem should also change. The serially correlated x_t , the limiting variance is given by

$$\begin{aligned}\text{var} \left(\frac{1}{\sqrt{n}} \sum_{i=1}^n x_i \right) &= 2 \sum_{k=1}^n \gamma_k \left(1 - \frac{k}{n} \right) + \gamma_0 \\ &\rightarrow 2 \sum_{k=1}^{\infty} \gamma_k + \gamma_0 = \sum_{k=-\infty}^{\infty} \gamma_k = S.\end{aligned}$$

We call S the *long-run variance*. There are many CLTs for serially correlated observations. We give the two mostly commonly cited versions: one applies to $\text{MA}(\infty)$ processes, the other one is more general.

Theorem 11.1. *Let y_t be an MA process: $y_t = \mu + \sum_{j=0}^{\infty} c_j \epsilon_{t-j}$ where ϵ_t is independent white noise and $\sum_{j=0}^{\infty} |c_j| < \infty$ (this implies ergodic), then*

$$\sqrt{T} \bar{y}_t \xrightarrow{d} N(0, S),$$

where $S = \sum_{k=-\infty}^{\infty} \gamma_k$ is the long-run variance.

Theorem 11.2 (Gordin's CLT). *Assume we have a strictly stationary and ergodic series $\{y_t\}$ with $\mathbb{E}(y_t^2) < \infty$ satisfying: $\sum_j \{\mathbb{E}[\mathbb{E}[y_t | I_{t-j}] - \mathbb{E}[y_t | I_{t-j-1}]]^2\}^{1/2} < \infty$ and $\mathbb{E}[y_t | I_{t-j}] \rightarrow 0$ as $j \rightarrow \infty$, then*

$$\sqrt{T} \bar{y}_t \xrightarrow{d} N(0, S),$$

where $S = \sum_{k=-\infty}^{\infty} \gamma_k$ is the long-run variance.

The Gordin's conditions are intended to make the dependence between distant observations to decrease to 0. ARMA process is a special case of Gordin series. The essence of these theorems is that we need some restrictions on dependencies for LLN and CLT to hold. We allow serial correlations as long as they are not too strong. If the observations become almost independent as they are far away in time, they can still apply the asymptotic theorems.

11.2 OLS for Time Series

Definition 11.1. Given a time series regression model

$$y_t = x_t' \beta + \epsilon_t,$$

x_t is **weakly exogenous** if

$$\mathbb{E}(\epsilon_t | x_t, x_{t-1}, \dots) = 0;$$

x_t is **strictly exogenous** if

$$\mathbb{E}(\epsilon_t | \{x_t\}_{t=-\infty}^{\infty}) = 0.$$

Strictly exogeneity requires innovations being exogenous from all past and future regressors; while weakly exogeneity only requires being exogenous from past regressors. In practice, strict exogeneity is too strong as an assumption. The weak exogenous is more practical and it is enough to ensure the consistency of the OLS estimator.

The OLS estimator is as usual:

$$\hat{\beta} = \beta + \left(\frac{1}{n} \sum_t x_t x_t' \right)^{-1} \left(\frac{1}{n} \sum_t x_t \epsilon_t \right).$$

Assuming LLN holds and x_t is weakly exogenous, we have

$$\begin{aligned} \frac{1}{n} \sum_t x_t x_t' &\rightarrow \mathbb{E}(x_t x_t') = Q, \\ \frac{1}{n} \sum_t x_t \epsilon_t &\rightarrow \mathbb{E}(x_t \epsilon_t) = \mathbb{E}[x_t \mathbb{E}[\epsilon_t | x_t]] = 0. \end{aligned}$$

Therefore, $\hat{\beta} \rightarrow \beta$. The OLS estimator is *consistent*.

Assuming the Gordin's conditions hold for $z_t = x_t \epsilon_t$, the CLT gives

$$\frac{1}{\sqrt{n}} \sum_t x_t \epsilon_t \rightarrow N(0, S),$$

where $S = \sum_{-\infty}^{\infty} \gamma_j$ is the long-run variance for z_t . Thus, we have the asymptotic normality for the OLS estimator

$$\sqrt{T}(\hat{\beta} - \beta) \rightarrow N(0, Q^{-1}SQ^{-1}).$$

Note how the covariance matrix S is different from the one in the *iid* case where $S = \sigma^2 \mathbb{E}(x_i x_i')$. The long-run variance S takes into account the auto-dependencies between observations. The auto-dependencies usually arise from the serially correlated error terms. It may also arise from x_t being autocorrelated and from conditional heteroskedasticity of the error terms. Because of the auto-covariance structure, S cannot be estimated in the same way as in the *iid* case. The estimator for S is called HAC (heteroskedasticity autocorrelation consistent) standard errors.

11.3 HAC Standard Errors

S can be estimated with truncated autocovariances,

$$\hat{S} = \sum_{j=-h(T)}^{h(T)} \hat{\gamma}_j.$$

$h(T)$ is a function of T and $h(T) \rightarrow \infty$ as $T \rightarrow \infty$, but more slowly. Because we don't want to include too many imprecisely estimated covariances. Another problem is the estimated \hat{S} might be negative. The solution is weight the covariances in a way to ensure positiveness:

$$\hat{S} = \sum_{j=-h(T)}^{h(T)} k_T(j) \hat{\gamma}_j.$$

$k_T(\cdot)$ is called a kernel. The weights are chosen to guarantee positive-definiteness by weighting down high lag covariances. Also we need $k_T(\cdot) \rightarrow 1$ for consistency.

A popular HAC estimator is the Newey-West variance estimator, in which $h(T) = 0.75T^{1/3}$ and $k_T(j) = \frac{h-j}{h}$, so that

$$\hat{S} = \sum_{j=-h}^h \left(\frac{h-j}{h} \right) \hat{\gamma}_j.$$

11.4 Example

Note that all of our discussions in this chapter apply only to stationary time series. Without stationarity, even the autocovariance γ_j might not be well-defined. In the following example, we generate artificial data from an AR(2) process, and recover the parameters by regression y_t on its lags.

```
library(lmtest)
y = arima.sim(list(ar = c(0.5, 0.3)), n = 1000)
mod = lm(y ~ ., data = cbind(y, lag(y,-1), lag(y,-2)))
coeftest(mod, vcov. = sandwich::NeweyWest(mod))
```

t test of coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.029790	0.031865	-0.9349	0.3501
`lag(y, -1)`	0.523868	0.036742	14.2580	< 2.2e-16 ***
`lag(y, -2)`	0.275884	0.035991	7.6654	4.227e-14 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

12 MLE for ARMA Models

OLS can only be used to estimate AR models, but not MA models. MA models or ARMA models in general can be estimated using maximum likelihood approach. Maximum likelihood estimation (MLE) starts with an assumed distribution of the random variables. The parameters are chosen to maximize the likelihood of observing the data under the distribution.

Consider an ARMA(p, q) model

$$y_t = \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} + u_t + \theta_1 u_{t-1} + \cdots + \theta_q u_{t-q}$$

Write in the form of data matrix:

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_T \end{bmatrix} = \underbrace{\begin{bmatrix} y_0 & y_{-1} & \cdots & y_{1-p} \\ y_1 & y_0 & \cdots & y_{2-p} \\ y_2 & y_1 & \cdots & y_{3-p} \\ \vdots & \vdots & \ddots & \vdots \\ y_T & y_{T-1} & \cdots & y_{T-p} \end{bmatrix}}_{\mathbf{X}} \begin{bmatrix} \phi_1 \\ \phi_2 \\ \phi_3 \\ \vdots \\ \phi_p \end{bmatrix} + \underbrace{\begin{bmatrix} 1 & 0 & 0 & \cdots & 0 \\ \theta_1 & 1 & 0 & \cdots & 0 \\ \theta_2 & \theta_1 & 1 & \cdots & 0 \\ \vdots & \vdots & & \ddots & \vdots \\ 0 & \cdots & \theta_2 & \theta_1 & 1 \end{bmatrix}}_{\mathbf{\Gamma}} \begin{bmatrix} u_1 \\ u_2 \\ u_3 \\ \vdots \\ u_T \end{bmatrix}$$

Or compactly,

$$\mathbf{y} = \mathbf{X}\boldsymbol{\phi} + \mathbf{\Gamma}\mathbf{u}.$$

We assume the innovations are jointly normal $\mathbf{u} \sim N(0, \sigma^2 \mathbf{I})$. We also assume the first p observations are known initial values $y_0, y_{-1}, \dots, y_{1-p}$ and $u_0 = u_{-1} = \cdots = u_{1-q} = 0$. Therefore, the observed data are jointly normal given the initial condition,

$$\mathbf{y}|\mathbf{y}_0 \sim N(\mathbf{X}\boldsymbol{\phi}, \sigma^2 \mathbf{\Gamma}\mathbf{\Gamma}').$$

The probability density function for multivariate normal is

$$f(\mathbf{y}|\mathbf{y}_0, \boldsymbol{\phi}, \mathbf{\Gamma}, \sigma^2) = (2\pi)^{-T/2} |\sigma^2 \mathbf{\Gamma}\mathbf{\Gamma}'|^{-1/2} \exp \left(-\frac{1}{2} (\mathbf{y} - \mathbf{X}\boldsymbol{\phi})' (\sigma^2 \mathbf{\Gamma}\mathbf{\Gamma}')^{-1} (\mathbf{y} - \mathbf{X}\boldsymbol{\phi}) \right)$$

To simplify computation, take logarithm to get the log-likelihood function

$$\ell(\phi, \Gamma, \sigma^2 | \mathbf{y}, \mathbf{y}_0) = -\frac{T}{2} \ln(2\pi) - \frac{1}{2} \ln |\sigma^2 \Gamma \Gamma'| - \frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{X}\phi)' (\Gamma \Gamma')^{-1} (\mathbf{y} - \mathbf{X}\phi).$$

The parameters are then chosen to maximize this log-likelihood function, i.e. the probability of observing the data under the assumed distribution. This can be done by conducting a grid search over the parameter space using a computer. To reduce the search dimensions, we may concentrate the log-likelihood by computing the first-order conditions:

$$\frac{\partial \ell}{\partial \phi} = 0 \implies \hat{\phi} = (\mathbf{X}'(\hat{\Gamma}\hat{\Gamma}')^{-1}\mathbf{X})^{-1}\mathbf{X}'(\hat{\Gamma}\hat{\Gamma}')^{-1}\mathbf{y}$$

$$\frac{\partial \ell}{\partial \sigma^2} = 0 \implies \hat{\sigma}^2 = \frac{1}{T}(\mathbf{y} - \mathbf{X}\hat{\phi})'(\hat{\Gamma}\hat{\Gamma}')^{-1}(\mathbf{y} - \mathbf{X}\hat{\phi})$$

This allows us to focus our search only on ϕ .

```
#' -----
#'  MLE Estimation for MA(1)
#'  -----

# Simulate an MA(1) process
y = arima.sim(list(ma=0.5), n=200)

# Negative log-likelihood function for MA(1)
neg_log_likelihood <- function(params, y) {
  theta <- params[1]
  sigma <- abs(params[2]) # Ensure sigma > 0
  n <- length(y)

  # Initialize residuals
  e_hat <- numeric(n)
  e_hat[1] <- y[1] #/ (1 + theta^2)

  for (t in 2:n) {
    e_hat[t] <- y[t] - theta * e_hat[t - 1]
  }

  ll <- -n/2 * log(2*pi*sigma^2) - sum(e_hat^2) / (2*sigma^2)

  return(-ll) # Return negative log-likelihood
}
```

```
# Initial guesses for theta and sigma
init_params <- c(0.1, 1)

# Run MLE using optim
result <- optim(par = init_params,
               fn = neg_log_likelihood,
               y = y)

# Estimated parameters
theta_hat <- result$par[1]
sigma_hat <- abs(result$par[2])

cat("Estimated theta:", theta_hat)
```

Estimated theta: 0.5226251

```
cat("Estimated sigma:", sigma_hat)
```

Estimated sigma: 1.013491

13 Forecasting

So far we have introduced basic univariate time series models and their estimation. One common application of univariate time series analysis is forecasting. Forecasting is a rather complex topic, with a wide range of techniques from basic ARMA models to machine learning. This book is not specialized in forecasting. We only devote this section to briefly cover forecasting based on ARMA models. We will start with some intuition. Then justify the intuition with a bit of formal theory.

13.1 Intuitive Approach

Suppose we have an AR(1) process,

$$y_t = \phi y_{t-1} + \epsilon_t, \quad \epsilon_t \sim \text{WN}(0, \sigma^2).$$

What would be the reasonable forecast for y_{T+1} given y_1, \dots, y_T ? It seems sensible to simply drop the white noise, as it is something completely unpredictable and it has mean zero. Thus,

$$\hat{y}_{T+1|T} = \phi y_T.$$

This is 1-period ahead forecast. But how do we forecast k -period ahead? Heuristically, we can simply iterate over to the future:

$$\begin{aligned} \hat{y}_{T+2|T} &= \phi \hat{y}_{T+1|T} = \phi^2 y_T, \\ \hat{y}_{T+h|T} &= \phi \hat{y}_{T+h-1|T} = \dots = \phi^h y_T. \end{aligned}$$

We will leave the heuristic solutions here and justify them later. If we accept this heuristic approach, we can easily generalize it to AR(p) processes:

$$\begin{aligned}
y_t &= \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \epsilon_t. \\
\hat{y}_{T+1|T} &= \phi_1 y_T + \phi_2 y_{T-1} + \cdots + \phi_p y_{T-p+1}, \\
\hat{y}_{T+2|T} &= \phi_1 \hat{y}_{T+1|T} + \phi_2 y_T + \cdots + \phi_p y_{T-p+2}, \\
&\vdots \\
\hat{y}_{T+h|T} &= \phi_1 \hat{y}_{T+h-1|T} + \phi_2 \hat{y}_{T+h-2|T} + \cdots + \phi_p y_{T-p+h}.
\end{aligned}$$

For MA(q) processes

$$y_t = \epsilon_t + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \cdots + \theta_q \epsilon_{t-q}$$

Suppose we know the past innovations until T : $\epsilon_T, \epsilon_{T-1}, \dots$. The best way to forecast $\hat{y}_{T+h|T}$ looks to simply discard $\epsilon_{T+1}, \dots, \epsilon_{T+h}$. Since we have no knowledge about future innovations given the information at time T . Therefore,

$$\begin{aligned}
\hat{y}_{T+1|T} &= \theta_1 \epsilon_T + \theta_2 \epsilon_{T-1} + \theta_3 \epsilon_{T-2} + \cdots \\
\hat{y}_{T+2|T} &= \theta_2 \epsilon_T + \theta_3 \epsilon_{T-1} + \cdots \\
&\vdots \\
\hat{y}_{T+h|T} &= \theta_h \epsilon_T + \theta_{h+1} \epsilon_{T-1} + \cdots
\end{aligned}$$

13.2 Best Linear Predictor

We now justify our heuristic solutions by the theory of best linear predictor. Suppose we want to forecast y given the information set X .

Definition 13.1. The best linear predictor (BLP) is defined as

$$\mathcal{F}(y|X) = x' \beta^*$$

which is a linear function of $X = (x_1, x_2, \dots, x_p)$ such that

$$\beta^* = \operatorname{argmin} \mathbb{E}(y - x' \beta)^2.$$

Taking first-order condition with respect to β gives

$$\beta^* = [\mathbb{E}(xx')]^{-1} \mathbb{E}(xy).$$

Therefore, the BLP is given by

$$\hat{y} = \mathcal{F}(y|X) = x'\beta^* = x'[\mathbb{E}(xx')]^{-1}\mathbb{E}(xy).$$

The prediction error is

$$r_{y|X} = y - \hat{y} = y - x'[\mathbb{E}(xx')]^{-1}\mathbb{E}(xy).$$

The BLP is the linear projection of y onto X . Because $\mathbb{E}[x(y - x'\beta)] = 0$. The forecast error is orthogonal to X .

Proposition 13.1. *BLP has the following properties:*

1. $\mathcal{F}[ax + by|z_1 \dots z_k] = a\mathcal{F}[x|z_1 \dots z_k] + b\mathcal{F}[y|z_1 \dots z_k]$;
2. If $x = a_1z_1 + \dots + a_kz_k$ is already a linear combination of $z_1 \dots z_k$, then $\mathcal{F}[x|z_1 \dots z_k] = x$;
3. If for all $1 \leq j \leq k$, $\text{cov}(x, z_j) = \mathbb{E}(xz_j) = 0$, then $\mathcal{F}[x|z_1 \dots z_k] = 0$.

13.3 Forecasting with ARMA Models

ARMA model is a basic yet powerful tool for forecasting. Given all stationary time series can be approximated by ARMA processes, it makes sense to model a stationary time series with ARMA, and then make forecast based on that model. We will see our heuristic solutions in the first part can be easily justified with the theory of BLP.

13.3.1 Forecasting with AR(p)

We have said that, for an AR(p) process

$$y_t = \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + \epsilon_t,$$

The one-step-ahead forecast is simply

$$\hat{y}_{T+1|T} = \phi_1 y_T + \phi_2 y_{T-1} + \dots + \phi_p y_{T-p+1}.$$

This is the BLP immediately from Property 2 of Proposition 13.1. We can also justify the iterated h -step-ahead forecast by Property 1 (assuming $h < p$):

$$\begin{aligned}
\hat{y}_{T+h|T} &= \phi_1 \hat{y}_{T+h-1|T} + \phi_2 \hat{y}_{T+h-2|T} + \cdots + \phi_h y_T + \cdots + \phi_p y_{T+h-p} \\
&= \phi_1 \mathcal{F}[y_{T+h-1}|y_T, y_{T-1}, \dots] + \cdots + \phi_p \mathcal{F}[y_{T+h-p}|y_T, y_{T-1}, \dots] \\
&= \mathcal{F}[\phi_1 y_{T+h-1} + \cdots + \phi_p y_{T+h-p} | y_T, y_{T-1}, \dots] \\
&= \mathcal{F}[y_{T+h} | y_T, y_{T-1}, \dots]
\end{aligned}$$

This is assuming all the forecast before h are BLPs, which can be justified recursively. Also note that for the values readily observed: y_T, y_{T-1}, \dots , the BLP is the value itself.

13.3.2 Forecasting with MA(q)

For the MA(q) process

$$y_t = \epsilon_t + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \cdots + \theta_q \epsilon_{t-q}$$

The BLP for h -step-ahead forecast is (assuming $h < q$)

$$\begin{aligned}
\hat{y}_{T+h|T} &= \mathcal{F}(y_{T+h} | \epsilon_T, \epsilon_{T-1}, \dots) \\
&= \mathcal{F}(\epsilon_{T+h} | \epsilon_T, \epsilon_{T-1}, \dots) + \theta_1 \mathcal{F}(\epsilon_{T+h-1} | \epsilon_T, \epsilon_{T-1}, \dots) + \cdots + \theta_q \mathcal{F}(\epsilon_{T+h-q} | \epsilon_T, \epsilon_{T-1}, \dots) \\
&= 0 + \cdots + 0 + \theta_h \epsilon_T + \cdots + \theta_q \epsilon_{T+h-q}
\end{aligned}$$

We make use of Property 3 of Proposition 13.1 with the knowledge that $\text{cov}(\epsilon_i, \epsilon_j) = 0$ for $i \neq j$. This result is also consistent with our intuition, because we have no knowledge of future innovations, the best thing we can do is assuming they are zeros. If $h > q$, then all $\mathcal{F}(\epsilon_{T+h-j} | \epsilon_T, \epsilon_{T-1}, \dots)$ are zero, which yields $\hat{y}_{T+h|T} = 0$.

In practice, we do not observe $\{\epsilon_t\}$. If we have an estimated MA model and we want to make forecast based on the model, we need to back out $\{\epsilon_t\}$ from $\{y_t\}$ by inverting the MA process: $\epsilon_t = \theta^{-1}(L)y_t$.

With the MA specification, we can easily compute the **Mean Squared Forecast Error (MSFE)** as follows

$$Q_{T+h|T} = \mathbb{E}(y_{T+h} - \hat{y}_{T+h|T})^2 = \mathbb{E} \left(\sum_{j=0}^{h-1} \theta_j \epsilon_{T+h-j} \right)^2 = \sigma^2 \sum_{j=0}^{h-1} \theta_j^2.$$

13.3.3 Forecasting with ARMA(p,q)

Consider the ARMA(p, q) process

$$(1 - \phi_1 L - \dots - \phi_p L^p)y_t = (1 + \theta_1 L + \dots + \theta_q L^q)\epsilon_t$$

We assume the process is causal and invertible. We can transform it to an AR(∞) process or MA(∞) process.

Causal form:

$$y_t = \phi^{-1}(L)\theta(L)\epsilon_t = \sum_{j=0}^{\infty} \psi_j \epsilon_{t-j}$$

Invertible form:

$$\epsilon_t = \theta^{-1}(L)\phi(L)y_t = \sum_{j=0}^{\infty} \pi_j y_{t-j}$$

or

$$y_t = - \sum_{j=1}^{\infty} \pi_j y_{t-j} + \epsilon_t$$

As we have seen so far, it is relatively easier to compute the mean forecast with AR models, and the MSFE with MA models. So we make forecast with the AR representation:

$$\hat{y}_{T+h|T} = - \sum_{j=1}^{h-1} \pi_j \hat{y}_{T+h-j} - \sum_{j=h}^{\infty} \pi_j y_{T+h-j}$$

However, we do not observe infinite past values in real world. We can only use the truncated values, discarding past values that we do not observe $y_0, y_{-1}, y_{-2}, \dots$

$$\hat{y}_{T+h|T} = - \sum_{j=1}^{h-1} \pi_j \hat{y}_{T+h-j} - \sum_{j=h}^{T+h-1} \pi_j y_{T+h-j}$$

We compute the MSFE with the MA representation:

$$Q_{T+h|T} = \mathbb{E}(y_{T+h} - \hat{y}_{T+h})^2 = \sigma^2 \sum_{j=0}^{h-1} \psi_j^2$$

If we can compute the prediction interval if we assume some probability distributions for the innovations. If we assume $\epsilon_t \stackrel{iid}{\sim} N(0, \sigma^2)$, then $(y_1, \dots, y_{T+h})'$ is jointly normal. Therefore,

$$y_{T+h} - \hat{y}_{T+h|T} \sim N(0, Q_{T+h|T})$$

The prediction interval is thus given by $\hat{y}_{T+h|T} \pm z_{\alpha/2} \sqrt{Q_{T+h|T}}$.

13.4 Applications

The following examples use ARMA models to forecast inflation rate and stock market index. The parameters of the ARMA models are chosen automatically. We can see for the inflation rate, the model produces some patterns in the forecast. But for the stock market index, the forecast is an uninformative flat line, indicating there is no useful patterns in the past data can be extrapolated by the ARMA model.

```
library(forecast)

# ARIMA is good for data with covariance patterns
AirPassengers |>
  arima(order = c(2,1,1), seasonal = c(0,1,0)) |>
  forecast(h=36) |>
  autoplot()
```





Figure 13.1: ARMA forecast for monthly inflation

```
# ARIMA provides little information for random walk process
EuStockMarkets[, 'DAX'] |>
  window(start = 1997) |>
  arima(order = c(2,1,1)) |>
  forecast(h=30) |>
  autoplot()
```




Figure 13.2: ARMA forecast for stock market index

14 Dynamic Causal Effect

As in all fields of science, we are perpetually interested in understanding the causal effect of one thing on another. In economics, we want to understand how monetary policy affects output and inflation, how exchange rate affects import and export, and so on. However, causality is something much easier said than done. In reality, there are multiple forces at work simultaneously that leads to the consequences we observed. It is challenging both conceptually and statistically to isolate the causality of a variable of particular interest.

In cross-sectional analysis, causality is defined counterfactually. That is, the causal effect of a treatment is defined as the difference between the treated outcome and the untreated outcome assuming that they would be otherwise the same without the treatment. In practice, that involves working with a large number of *iid* observations that are similar on average only differentiated by the status of the treatment. This approach, however, does not work well with many macroeconomic studies. For example, suppose we want to figure out the causal effect of monetary policy on inflation rate. The cross-sectional approach would entail finding a large number of almost identical countries, each with independent monetary policy. And a random subset of them tighten their monetary policies while others do not. Then we work out the different economic outcomes between these two groups. This is clearly infeasible. The question we posed concerns only one country with inflation and interest rates observed through time. We would need a definition of causal effect that encompasses observations over time not across individuals.

Suppose ϵ_t denote a random treatment happened at time t . Then the causal effect on an outcome variable y_{t+h} , h periods ahead, of a unit shock in ϵ is defined as

$$\mathbb{E}[y_{t+h}|\epsilon_t = 1] - \mathbb{E}[y_{t+h}|\epsilon_t = 0]. \quad (14.1)$$

We require the randomness of the treatment ϵ_t in a sense that it is uncorrelated with any other variables that could possible have an impact on the outcome. Therefore, ϵ_t happens or not does not affect other forces that shape the outcome. The difference in the outcomes is solely attributable to ϵ_t . It is this randomness that guarantees a causal interpretation.

Our example of monetary policy above clearly does not meet this requirement. The monetary authority does not set the interest rate randomly, but based on the economic conditions of the time, which makes it correlated with other economic variables that could also have an impact on inflation. A qualified random shock may be a change in weather conditions. Weather has huge impact on agricultural production, but it is determined independent of any human activity.

If ϵ_t denotes a rainy day at time t , and y_{t+h} be the agricultural production, Equation 14.1 could be a plausible causal effect. However, most variables of interest in economics are endogenously determined. How to estimate the causal effect in such cases is an art in itself. We will come back to this point later.

The conceptual definition of Equation 14.1 can not be computed directly as the counterfactual is not observed. What we have is a sample of experiments over time, in which the treatment happens randomly at some points but not others, $\{\epsilon_1 = 0, \epsilon_2 = 1, \epsilon_3 = 0, \dots\}$. We could envision that if we have long enough observations, by comparing the outcomes when the shock happens and when it does not, it gives us an reasonable estimation of the causal effect because all other factors that contributing to the outcome, despite they are changing over time, would be averaged out provided the randomness of the treatment.

Assuming linearity and stationarity, the causal effect of Equation 14.1 can be effectively captured by a regression framework,

$$y_{t+h} = \theta_h \epsilon_t + u_{t+h},$$

where u_{t+h} represents all other factors contributing to the outcome variable. Since ϵ_t is random, it holds that $\mathbb{E}(u_{t+h}|\epsilon_t) = 0$. Therefore,

$$\theta_h = \mathbb{E}(y_{t+h}|\epsilon_t = 1) - \mathbb{E}(y_{t+h}|\epsilon_t = 0).$$

Thus, θ_h captures the causal effect of one unit shock of ϵ_t on y_{t+h} . The path of the causal effects mapped out by $\{\theta_0, \theta_1, \theta_2, \dots\}$ is called the **dynamic causal effect**, in a sense that it is the causal effects through time.

15 The Structural Shock Framework

The counterfactual framework introduced in the last section defines the dynamic causal effect of any variable on another. As economists, we are more interested in understanding the causal relationships between important forces that drive the economy. We now introduce the **structural shock framework**, or the **Slutsky-Frisch paradigm**. This paradigm is explicitly or implicitly embedded in virtually every mainstream macroeconomic models or econometric models. It is not an essential component of time series analysis. But, as we would like to approach the topic from an economist's perspective, it is good to have this framework in mind for many of our applications.

The structural shock framework envisions our economy as a complex system driven by a set of fundamental structural forces and coordinated by numerous price signals that automatically balance the demand and supply of all goods and services. The structural forces could be technology progress, climate change, policy changes and so on. These structural shocks are the primitive forces underlying our economy. When a structural shock happens, it triggers a reallocation of economic resources guided by market forces. In theoretical works, we are interested in modelling the system as a whole, particularly how resources are allocated optimally by market forces. In empirical works, we are interested how to recover the underlying structural shocks and estimate their causal effect on other economic variables.

In the language of time series analysis, we can envision our economy as an MA process, in which the observable variables (output, employment, inflation, etc) are the outcomes of accumulated past and current structural shocks:

$$\mathbf{y}_t = \Theta(L)\boldsymbol{\epsilon}_t,$$

or

$$\begin{bmatrix} y_{1t} \\ y_{2t} \\ \vdots \\ y_{nt} \end{bmatrix} = \sum_{j=0}^{\infty} \begin{bmatrix} \theta_{j,11} & \theta_{j,12} & \cdots & \theta_{j,1m} \\ \theta_{j,21} & \theta_{j,22} & \cdots & \theta_{j,2m} \\ \vdots & \vdots & \ddots & \vdots \\ \theta_{j,n1} & \theta_{j,n2} & \cdots & \theta_{j,nm} \end{bmatrix} \begin{bmatrix} \epsilon_{1,t-j} \\ \epsilon_{2,t-j} \\ \vdots \\ \epsilon_{m,t-j} \end{bmatrix}.$$

\mathbf{y}_t represents the vector of economic variables of concern. The space of \mathbf{y}_t are spanned by m structural shocks (current and past): $\{\boldsymbol{\epsilon}_{t-j}\}_{j=0}^{\infty}$.

Structural shocks are conceptual constructions that are primitive, unforeseeable, and uncorrelated underlying forces. Whether structural shocks do exist or not is an open question. But they are useful constructions that enable econometricians to disentangle different driving forces of the outcome variable.

In reality, almost every economic variable is endogenous. For example, monetary policy (interest rate) is set by the monetary authority based on their assessment of the economic conditions. But we can also imagine, there is a “genuine” component of the monetary policy, which may come from the personality of the policymaker and his mental conditions when he make the decision, that is not predictable from other variables. This genuine component is what we deem as the “monetary policy shock”. It is a shock in a sense that it is not predictable. It speaks for its own sake and contribute to the economic outcomes independently.

We do not observe the structural shocks directly. The observable variables are linear combinations of the structural shocks. For example, we may think of the observed interest rate as a linear combination of the monetary policy shock together with supply-side shocks and others.

$$i_t = \theta_1(L)\epsilon_t^{\text{MP}} + \theta_2(L)\epsilon_t^{\text{SS}} + \dots$$

Therefore, regressing inflation or output on interest rate will not give the causal effect of the monetary policy. Because interest rate does not represent the “genuine” monetary policy shock. It is determined by other economic variables and there are multiple structural forces at work. There are numerous literature that works on methods to isolate the “monetary policy shock” from the observed interest rates. Such way of constructing the structural shocks is not only a conceptual idea, but also a prerequisite for meaningful interpretation of the coefficients of econometric models.

16 Estimating Dynamic Multipliers

This section will cover the specifications commonly used to estimate dynamic causal effect. Just like in cross-sectional analysis, regression techniques can always be applied as long as the time series are covariance stationary without an emphasis on causality. However, we pay special attention to causal inferences, as we are more interested in understanding the causality rather than mere correlations in most empirical researches. We start with the case where the structural shock is directly observed and move on to the cases where the structural shocks need to be constructed.

16.1 Distributed Lags

The easiest approach to estimate dynamic causal effect is to include lags in the specification:

$$y_t = \beta_0 \epsilon_t + \beta_1 \epsilon_{t-1} + \cdots + \beta_p \epsilon_{t-p} + u_t,$$

where ϵ_t is the structural shock, u_t is everything that otherwise influences y_t . Since ϵ_t happens randomly, we have $\mathbb{E}(u_t | \epsilon_{t-j}) = 0$. Thus, the β s, which capture the dynamic causal effect, would be consistently estimated by OLS.

Note that we call it a specification, in a sense that the joint distribution of the random variables is unknown, which distinguishes itself from the DGP model in Chapter 6. But it does not stop us from uncovering the causal effect, as long as the exogenous condition holds.

The effect of a unit change in ϵ on y after h periods, which is β_h , is also called the h -period **dynamic multiplier**. Sometimes, we are interested in the accumulated effect over time, $\beta_0 + \beta_1 + \cdots + \beta_h$, which is called **cumulative dynamic multiplier**.

Because u_t is the linear combination of all other current and past shocks, it is likely serially correlated. So HAC standard errors are required for robust inferences.

Proposition 16.1. *Assumptions for a consistent estimation of dynamic causal effects with distributed lag models:*

1. ϵ is an exogenous shock, $\mathbb{E}(u_t | \epsilon_t, \epsilon_{t-1}, \dots) = 0$;
2. All variables are stationary;
3. Regular conditions for OLS to work.

To reduce the serial correlations $\{u_t\}$, and also allow for slow adjustment of y_t , we can also include lagged dependent variables in the specification, which becomes an **autoregressive distributed lag (ADL)** specification:

$$y_t = \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} + \beta_0 \epsilon_t + \beta_1 \epsilon_{t-1} + \cdots + \beta_p \epsilon_{t-p} + u_t,$$

or

$$\phi(L)y_t = \beta(L)\epsilon_t + u_t.$$

When lags of the dependent variable are included as regressors, strict exogeneity fails for sure, because $X = \{y_{t-1}, \dots, \epsilon_t, \epsilon_{t-1}, \dots\}$ is correlated with past errors u_{t-1} , despite it is uncorrelated with the contemporary error u_t . The OLS is consistent so long as $\{u_t\}$ are not serially correlated. Otherwise, u_t would be correlated with X through u_{t-1} . The serial correlation can be tested with Durbin-Watson test or Breusch-Godfrey test.

The dynamic causal effect is more convoluted with the ADL specification though,

$$\hat{\theta}(L) = \hat{\phi}^{-1}(L)\hat{\beta}(L).$$

ADL also require truncated lags. p and q are chosen as an increasing function of the sample size. In general, choosing p and q to be of order $T^{1/3}$ would be sufficient for consistency.

16.2 Local Projections

Dynamic causal effect can also be estimated by projecting future outcomes directly on the shock. Jordà (2005) named it **local projections (LP)**.

$$y_{t+h} = \theta_h \epsilon_t + u_{t+h}.$$

By assumption, $\mathbb{E}(u_{t+h}|\epsilon_t) = 0$. So $\hat{\theta}_h$ is a consistent estimate of the h -period dynamic multiplier. HAC standard errors are also required in local projections, as u_{t+h} in are usually serially correlated.

Readers may wonder, since ADL and LP both give consistent estimates of the dynamic multipliers, what is the difference between them. There are two obvious differences:

1. Lagged shocks do not appear in LP specifications as they do in distributed lag specifications.
2. The LP method requires running separate regressions for each h . The dynamic response $\{\theta_0, \theta_1, \theta_2, \dots\}$ are estimated through multiple regressions rather than one.

The error structure is also different. To see this, suppose the DGP is an $MA(\infty)$ process

$$y_t = \epsilon_t + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \dots$$

If we estimate it with a DL specification with two lags,

$$y_t = \beta_0 \epsilon_t + \beta_1 \epsilon_{t-1} + u_t,$$

where $u_t = \sum_{j=2}^{\infty} \theta_j \epsilon_{t-j}$. Exogeneity would ensure $\hat{\beta}_1 \rightarrow \theta_1$.

We can also estimate it with a local projection (suppose we are interested in the one-step-ahead dynamic multiplier):

$$y_{t+1} = \psi_1 \epsilon_t + u_{t+1}.$$

Again, we have consistency $\hat{\psi}_1 \rightarrow \theta_1$. But the error structure is different $u_{t+1} = \epsilon_{t+1} + \sum_{j=2}^{\infty} \theta_j \epsilon_{t-j}$.

Both the DL and LP specifications may include additional control variables, which can reduce the variance of the residuals and improve the efficiency of the estimators.

16.3 Example of Observable Exogenous Shocks

Directly observable exogenous shocks are rare. Here we use an example from Stock and Watson's textbook, which explores the dynamic causal effect of cold weather on orange juice prices. Cold weather is bad for orange production. Orange trees cannot withstand freezing temperatures that last for more than a few hours. Florida accounts for more than 98 percent of U.S. production of frozen concentrated orange juice. Therefore, the frozen weather in Florida would reduce the supply and orange juice and raise the price. The dataset includes the number of freezing degree days in Florida and the average producer price for orange juice. Cold weather is plausibly exogenous, which allows us to utilize the regression framework above to estimate the dynamic causal effect.

```
library(AER)
library(dynlm)
library(lmtest)

data("FrozenJuice") # load data

# compute percentage change on price
pchg = 100*diff(log(FrozenJuice[, 'price']))
```

```

sample = ts.union(fdd = FrozenJuice[, 'fdd'], pchg)

# distributed lag model
mod = dynlm(pchg ~ L(fdd, 0:6), data = sample)

# compute Newey-West standard error
coeftest(mod, vcov. = NeweyWest)

```

t test of coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.418188	0.227219	-1.8405	0.0661952 .
L(fdd, 0:6)0	0.484907	0.135850	3.5694	0.0003865 ***
L(fdd, 0:6)1	0.139460	0.084986	1.6410	0.1013277
L(fdd, 0:6)2	0.057297	0.057063	1.0041	0.3157410
L(fdd, 0:6)3	0.068637	0.043709	1.5703	0.1168703
L(fdd, 0:6)4	0.035142	0.028405	1.2372	0.2165036
L(fdd, 0:6)5	0.048680	0.028938	1.6822	0.0930485 .
L(fdd, 0:6)6	0.040545	0.045168	0.8976	0.3697343

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

# To plot the dynamic multiplier, the following code creates
# a custom function that manually plots the coefficients and
# the confidence intervals
plotDM <- function(mod,
                    horizon = 0:5,
                    vcov = "NeweyWest",
                    col = "red") {

  # only Newey-West standard error is supported
  if (vcov == "NeweyWest") {
    ci = coefci(mod, vcov. = sandwich::NeweyWest)
  }

  # extract coefficients of the lagged regressors
  plot(horizon, mod$coefficients[-1], type = "l",
       col = col, ylim = c(min(ci[-1,]), max(ci[-1,])),
       xlab = "Lags", ylab = "Dynamic Multiplier")
}

```


```

# horizontal zero
abline(h = 0, col = "gray")

# confidence intervals
lines(horizon, ci[-1,1], lty = 2, col = col)
lines(horizon, ci[-1,2], lty = 2, col = col)
}

plotDM(mod, horizon = 0:6)

```



We can also use local projections. Note that local projections require estimating multiple regressions. The coefficients from each of the regressions constitute the dynamic multiplier.

```

# apply local projection for horizons 0-6
lps = sapply(0:6, function(h) {

  # regress future price change on fdd
  lp = dynlm(L(pchg, -h) ~ fdd, data = sample)

  # Newey-West confidence interval
  ci = coefci(lp, vcov. = NeweyWest)

  # extract coefficients and CIs


```

```

    c(lp$coefficients[-1], ci[-1,]) # remove intercept
}) |> t() # transpose it

# plot the LP coefficients
{
  plot(0:6, lps[, 'fdd'], type = "l", col = 2, ylim = c(-0.05, 0.75), ann = F)
  abline(h = 0, col = "gray")
  lines(0:6, lps[, '2.5 %'], lty = 2, col = 2)
  lines(0:6, lps[, '97.5 %'], lty = 2, col = 2)
}

```



16.4 Example of Constructed Structural Shocks

Most structural shocks in economics are not directly observed, such as monetary policy shocks, or fiscal policy shocks, yet they are of profound interest of researchers. As we have explained before, regressing output or inflation on interest rate does not give a plausible estimation of the causal effect of monetary policy, due to the endogeneity problem. Thus, we need to isolate the exogenous part of the monetary policy from observed variables. The method to achieve this is an active research field in itself.


We may utilize the high-frequency price changes of fed fund futures around the window of monetary policy announcement to approximate the monetary policy shock. The rationale of

this construction is that, the price of the financial instrument reflects the expected interest rate by market participants based on the economic conditions. Therefore, the sudden change of the price in the tiny window of monetary policy announcement captures the unexpected part of the monetary policy.

```
# load US monthly data
data = read.csv("data/usmd.csv")

# convert to time series
ffr = ts(data[, 'FEDFUNDS'], start = c(1990,1), frequency = 12)
mps = ts(data[, 'MPS'], start = c(1990,1), frequency = 12)

# interest rate changes and MP shocks are different
plot(cbind(ffr, mps), main = "Monetary Policy Shocks")
```




We estimate the dynamic causal effect of monetary policy shocks on stock market returns. We employ the simplest regression specification. Note that there is quite a difference when using changes in the FF rate directly versus using identified monetary policy shocks. In the regression using the FF rate, the coefficient is not statistically significant. One might tempt one to conclude that monetary policy has no impact on stock market returns based this result. However, this conclusion is misleading because common factors, such as overall economic conditions, influence both stock market returns and the FF rate. In our second experiment—where we regress stock market returns on the identified monetary policy shocks—we observe a much larger coefficient and a significant negative impact, especially in the month immediately following the shock.

```
# monthly S&P 500 returns
ret = data[, 'SP500'] |>
  ts(start = c(1990,1), frequency = 12) |>
  log() |>
  diff()


# regression on FF rate changes
mod_1 = dynlm(ret ~ L(d(ffr), 0:6))

# plot dynamic multiplier
plotDM(mod_1, horizon = 0:6)
```



```
# regression with MP shocks
mod_2 = dynlm(ret ~ L(mps, 0:6))

# plot dynamic multiplier
plotDM(mod_2, horizon = 0:6, col = "blue")
```



17 Instrument Variables

If a structural shock is not directly observable, neither can it be constructed through observable variables, we can identify it using an instrument variable approach if an instrument is available.

Suppose our observable space $\mathbf{y} = (y_1, y_2, \dots)'$ is spanned by multiple structural shocks $\epsilon = (\epsilon_1, \epsilon_2, \dots)'$. We want to identify the causal effect of structural shock ϵ_1 . An instrument variable z satisfies the following conditions:

1. $\mathbb{E}(\epsilon_{1t}z_t) = \alpha \neq 0$ (relevance);
2. $\mathbb{E}(\epsilon_{2:n}z_t) = 0$ (contemporaneous exogeneity);
3. $\mathbb{E}(\epsilon_{t+j}z_t) = 0$ for $j \neq 0$ (lead-lag exogeneity).

$\epsilon_{2:n}$ denotes all other structural shocks except ϵ_1 . The lead-lag exogeneity is unique to time series. To understand this, consider an local projection: $y_{t+h} = \theta_h \epsilon_t + u_{t+h}$. As illustrated in the last section, u_{t+h} is a linear combination of the entire history of structural shocks. If z_t is to identify the causal effect of shock ϵ_{1t} alone, it must be uncorrelated with all leads and lags. The requirement that z_t be uncorrelated with future ϵ 's is generally not restrictive — by definition, future shocks are unanticipated. To the contrary, the requirement that z_t be uncorrelated with past ϵ 's is more restrictive and hard to meet.

Suppose we want to estimate the causal effect of $\epsilon_{1,t}$ on $y_{2,t+h}$, where $\epsilon_{1,t}$ is only observable through $y_{1,t}$. Suppose we have an instrument variable z_t that satisfies the above conditions. The local projection

$$y_{2,t+h} = \theta_{h,21}y_{1,t} + u_{t+h}$$

cannot be consistently estimated because $y_{1,t}$ and u_{t+h} are correlated. However, with the help with z_t as an instrument, we can consistently estimate the dynamic multiplier $\theta_{h,21}$:

$$\begin{aligned} \beta_{\text{LP-IV}} &= \frac{\mathbb{E}(y_{2,t+h}z_t)}{\mathbb{E}(y_{1,t}z_t)} \\ &= \frac{\mathbb{E}[(\theta_{h,21}y_{1,t} + u_{t+h})z_t]}{\mathbb{E}(y_{1,t}z_t)} \\ &= \frac{\theta_{h,21}\alpha}{\alpha} = \theta_{h,21}. \end{aligned}$$

Lead-lag exogeneity implies z_t being unforecastable in a regression of z_t on lags of y_t . If the exogeneity fails, LP-IV is not consistent. This problem can be partially addressed by including control variables in the regression:

$$y_{2,t+h} = \theta_{h,21}y_{1,t} + \gamma'_h \mathbf{w}_t + u_{t+h}^\perp.$$

We could also include lagged values of y_t or other lagged variables. The IV estimator is consistent if \mathbf{w}_t absorbs all past shocks that could potentially correlated with z_t . In a broad sense, the validity of the instrument variable with additional controls requires that the controls span the space of all structural shocks.

```
library(lpirfs)

# Load data
ag_data <- na.omit(ag_data)

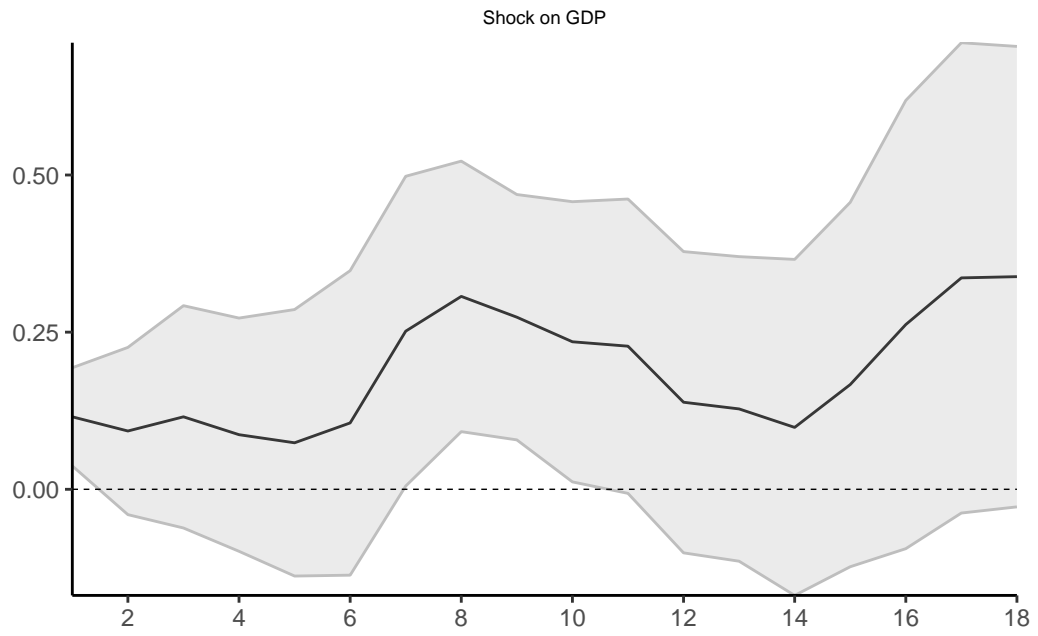
# Endogenous data (government spending, tax, GDP)
endog <- ag_data[,3:5]

# Variable to shock with (government spending)
gov <- ag_data[,3]

# Government spending shock identified by
# Ramey and Zubairy (2018) using military news
gov_shock <- ag_data[,7]

# Estimate linear model via 2SLS
results <- lp_lin_iv(endog, lags_endog_lin = 4,
                     shock = gov, instrum = gov_shock, use_twosls = T,
                     trend = 0, confint = 1.96, hor = 18)

# Show all responses
plot_lin(results)[[3]]
```



Part IV

Nonstationary Time Series

18 Spurious Regression

It is said, all stationary series are alike, but each non-stationary series is non-stationary in its own way (remember Leo Tolstoy's famous quote: *all happy families are alike; each unhappy family is unhappy in its own way.*)

In all previous chapters, we have been working on stationary processes. We have shown that similar regression techniques and asymptotic results hold for stationary processes as for *iid* observations, albeit not exactly the same. If a time series is not stationary, we transform it to stationary by taking differences.

This chapter is devoted to study non-stationary time series. Special attention is given to unit root processes. We will see the theories involving non-stationary processes are entirely different from those applied to stationary processes. This makes unit root analysis an rather independent topic.

The obsession with unit root in academia have faded away in recent decades (I do not know if this assessment is accurate). Despite the topic posses immense theoretical interest, it does not seem to provide proportionate value for applied studies. Nonetheless, the topic is indispensable for a comprehensive understanding of time series analysis.

We start by pointing out that, it is very dangerous to blindly include non-stationary variables in a regression. To illustrate this, consider two random walks:

$$\begin{aligned}x_t &= x_{t-1} + \epsilon_t, & \epsilon_t &\stackrel{iid}{\sim} N(0, \sigma_X^2) \\ y_t &= y_{t-1} + \eta_t, & \eta_t &\stackrel{iid}{\sim} N(0, \sigma_Y^2)\end{aligned}$$


ϵ_t and η_t are independent to each other.

```
set.seed(1)

# generate two random walk processes
x = cumsum(rnorm(200))
y = cumsum(rnorm(200))

plot.ts(cbind(x,y), main="Two Random Walks")
```

Two Random Walks



We would expect the two series completely uncorrelated, as they are two independent random processes. However, if we regress y_t on x_t , we would likely find a very strong correlation. This is called a **spurious regression**.

$$y_t = \alpha + \beta x_t + u_t$$

```
lm(y ~ x) |> summary()
```

Call:

```
lm(formula = y ~ x)
```

Residuals:

Min	1Q	Median	3Q	Max
-6.831	-3.904	1.021	3.223	9.731

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.23975	0.56584	5.726	3.78e-08 ***
x	0.15324	0.06886	2.225	0.0272 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.049 on 198 degrees of freedom
Multiple R-squared: 0.0244, Adjusted R-squared: 0.01948
F-statistic: 4.953 on 1 and 198 DF, p-value: 0.02718

Note that if we difference the two series to stationary, the spurious correlation disappears.

```
lm(diff(y) ~ diff(x)) |> summary()
```

Call:

```
lm(formula = diff(y) ~ diff(x))
```

Residuals:

Min	1Q	Median	3Q	Max
-2.93078	-0.55055	-0.02841	0.67557	2.58926

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.03963	0.07203	0.55	0.583
diff(x)	-0.02171	0.07756	-0.28	0.780

Residual standard error: 1.015 on 197 degrees of freedom
Multiple R-squared: 0.0003974, Adjusted R-squared: -0.004677
F-statistic: 0.07831 on 1 and 197 DF, p-value: 0.7799

If we simulate many random walks, we would observe that a large percentage of the regressions report statistically significant relationships even though the series are independent.

```
# Simulation parameters
n <- 100          # Number of observations in each random walk
num_sim <- 1000   # Number of simulations
sig_count <- 0    # Counter for significant p-values
pvals <- numeric(num_sim) # To store p-values

# Loop over simulations
for (i in 1:num_sim) {
  # Generate two independent random walks of length n
  x <- cumsum(rnorm(n))
  y <- cumsum(rnorm(n))
```

```

# Run linear regression of y on x
model <- lm(y ~ x)
pval <- summary(model)$coefficients[2, 4]
pvals[i] <- pval
}

# Calculate percentage of simulations with p-value < 0.05
sig_percent <- mean(pvals < 0.05) * 100
print(paste("Percentage of significant regressions:", sig_percent, "%"))

```

```
[1] "Percentage of significant regressions: 77 %"
```

This example gives you a quantitative feel for how frequently spurious results can occur. What's more interesting, however, is that we can eliminate this spurious strong correlation by including lags of dependent and independent variables —

```

# Compute lags for x and y
y_lag = dplyr::lag(y)
x_lag = dplyr::lag(x)

# Regression with lags
lm(y ~ y_lag + x + x_lag) |> summary()

```

Call:

```
lm(formula = y ~ y_lag + x + x_lag)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.6094	-0.6136	0.0433	0.7026	2.4059

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.34804	0.32558	1.069	0.288
y_lag	0.91403	0.03597	25.413	<2e-16 ***
x	0.12296	0.11279	1.090	0.278
x_lag	-0.05127	0.11040	-0.464	0.643

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.076 on 95 degrees of freedom

(1 observation deleted due to missingness)
Multiple R-squared: 0.9206, Adjusted R-squared: 0.9181
F-statistic: 367.3 on 3 and 95 DF, p-value: < 2.2e-16

Non-stationary time series can cause troubles, but they are also fascinating topics to explore. In what follows, we will focus on two types of non-stationary processes: trend-stationary processes and unit root processes, which are the most common types of non-stationary series we would encounter in economic and finance. Non-stationary series with exponential growth can be transformed into linear trend, hence is not of particular interest. We will start with the relatively easy trend-stationary processes, and spend most of the paragraphs on unit root processes.

19 Trend Stationary

Trend-stationary process is a stationary process round a deterministic trend:

$$y_t = \alpha + \delta t + \psi(L)\epsilon_t,$$

where δt is a deterministic linear time trend, $\psi(L)\epsilon_t$ is a stationary process. After de-trending $-(\alpha + \delta t)$, the result is a stationary process.

Trend stationary vs stochastic trend

Trend-stationary processes must be distinguished from *stochastic trend process* (unit root with a drift):

$$y_t = \delta + y_{t-1} + \epsilon_t = y_0 + \delta t + \sum_{j=1}^t \epsilon_j.$$

Both of them have a time trend component. But in the latter model, the stochastic component is not stationary. In other words, an innovation in a trend-stationary model does not have long-lasting effect, whereas the effect is persistent in a stochastic trend model.

The difference becomes clearer by comparing the variances. The variance of the trend-stationary process

$$\text{var}(y_t) = \psi^2(L)\sigma^2$$

is constant which does not depend on time. However, the variance of the stochastic trend process

$$\text{var}(y_t) = \text{var}\left(\sum_{j=1}^t \epsilon_j\right) = \sigma^2 t$$

is increasing over time. Therefore, stochastic trend process fluctuates more widely as time goes by.

Unlike unit root processes, trend-stationary processes can be safely estimated by OLS. The usual t and F statistics have the same asymptotic distribution as they are for stationary

processes. But they converge at a different speed, due to the presence of the trend. To see this, rewrite the regression in vector form

$$y_t = \alpha + \delta t + \epsilon_t = \begin{bmatrix} 1 & t \end{bmatrix} \begin{bmatrix} \alpha \\ \delta \end{bmatrix} + \epsilon_t = \mathbf{x}'_t \boldsymbol{\beta} + \epsilon_t;$$

For simplicity, assume $\epsilon_t \sim IID(0, \sigma^2)$ for the following computation. The result can be generalized to ϵ_t being stationary. The OLS estimator is given by

$$\begin{aligned} \hat{\boldsymbol{\beta}} &= \begin{bmatrix} \hat{\alpha} \\ \hat{\delta} \end{bmatrix} = \left(\sum_t x_t x'_t \right)^{-1} \left(\sum_t x_t y_t \right), \\ \sqrt{T}(\hat{\boldsymbol{\beta}} - \boldsymbol{\beta}) &= \left(\frac{1}{T} \sum_t x_t x'_t \right)^{-1} \left(\frac{1}{\sqrt{T}} \sum_t x_t \epsilon_t \right). \end{aligned}$$

The usual asymptotic results are

$$\begin{aligned} \frac{1}{T} \sum_t x_t x'_t &\rightarrow Q \\ \frac{1}{\sqrt{T}} \sum_t x_t \epsilon_t &\rightarrow N(0, \sigma^2 Q) \\ \sqrt{T}(\hat{\boldsymbol{\beta}} - \boldsymbol{\beta}) &\rightarrow N(0, \sigma^2 Q^{-1}) \end{aligned}$$

But this is not the case with deterministic trend if we do the computation:

$$\frac{1}{T} \sum_t x_t x'_t = \begin{bmatrix} 1 & \frac{1}{T} \sum t \\ \frac{1}{T} \sum t & \frac{1}{T} \sum t^2 \end{bmatrix}$$

does not converge. Because $\sum_{t=1}^T t = \frac{T(T+1)}{2}$, and $\sum_{t=1}^T t^2 = \frac{T(T+1)(2T+1)}{6}$. It requires stronger divider to make them converge, $T^{-2} \sum_{t=1}^T t \rightarrow \frac{1}{2}$, $T^{-3} \sum_{t=1}^T t^2 \rightarrow \frac{1}{3}$. In general,

$$\frac{1}{T^{v+1}} \sum_{t=1}^T t^v \rightarrow \frac{1}{v+1}.$$

Dividing by T^3 will make the convergence

$$\frac{1}{T^3} \sum_t x_t x'_t \rightarrow \begin{bmatrix} 0 & 0 \\ 0 & \frac{1}{3} \end{bmatrix}$$

However, this matrix is not invertible. We need different rates of convergence for $\hat{\alpha}$ and $\hat{\delta}$.

Define

$$\gamma_T = \begin{bmatrix} \sqrt{T} & 0 \\ 0 & T^{3/2} \end{bmatrix}$$

Multiply this matrix with the coefficient vector would apply different convergence speed to different coefficients:

$$\begin{aligned} \begin{bmatrix} \sqrt{T}(\hat{\alpha} - \alpha) \\ T^{3/2}(\hat{\delta} - \delta) \end{bmatrix} &= \gamma_T \left(\sum_t x_t x_t' \right)^{-1} \left(\sum_t x_t \epsilon_t \right) \\ &= \left[\gamma_T^{-1} \left(\sum_t x_t x_t' \right) \gamma_T^{-1} \right]^{-1} \left[\gamma_T^{-1} \left(\sum_t x_t \epsilon_t \right) \right] \end{aligned}$$

in which

$$\begin{aligned} \gamma_T^{-1} \left(\sum_t x_t x_t' \right) \gamma_T^{-1} &= \begin{bmatrix} T^{-1/2} & \\ & T^{-3/2} \end{bmatrix} \begin{bmatrix} \sum 1 & \sum t \\ \sum t & \sum t^2 \end{bmatrix} \begin{bmatrix} T^{-1/2} & \\ & T^{-3/2} \end{bmatrix} \\ &= \begin{bmatrix} T^{-1} \sum 1 & T^{-2} \sum t \\ T^{-2} \sum t & T^{-3} \sum t^2 \end{bmatrix} \rightarrow \begin{bmatrix} 1 & \frac{1}{2} \\ \frac{1}{2} & \frac{1}{3} \end{bmatrix} = Q. \end{aligned}$$

Turning to the second term:

$$\gamma_T^{-1} \left(\sum_t x_t \epsilon_t \right) = \begin{bmatrix} T^{-1/2} & \\ & T^{-3/2} \end{bmatrix} \begin{bmatrix} \sum \epsilon_t \\ \sum t \epsilon_t \end{bmatrix} = \begin{bmatrix} T^{-1/2} \sum \epsilon_t \\ T^{-1/2} \sum \frac{t}{T} \epsilon_t \end{bmatrix}$$

$T^{-1/2} \sum \epsilon_t \rightarrow N(0, \sigma^2)$ by standard CLT. Observe that $\zeta_t = \frac{t}{T} \epsilon_t$ is not serially correlated,

$$\mathbb{E}(\zeta_t \zeta_{t-j}) = \frac{t(t-j)}{T^2} \mathbb{E}(\epsilon_t \epsilon_{t-j}) = 0$$

with stabilized variance

$$\text{var}(T^{-1/2} \sum \zeta_t) = \frac{1}{T} \sum \text{var} \left(\frac{t}{T} \epsilon_t \right) = \frac{\sigma^2}{T^3} \sum t^2 \rightarrow \frac{\sigma^2}{3}$$

Therefore, $T^{-1/2} \sum \frac{t}{T} \epsilon_t \rightarrow N(0, \frac{\sigma^2}{3})$. We also need to consider the covariance,

$$\text{cov}(T^{-1/2} \sum \epsilon_t, T^{-1/2} \sum \frac{t}{T} \epsilon_t) = \frac{1}{T} \mathbb{E} \left(\sum \epsilon_t \sum \frac{t}{T} \epsilon_t \right) = \frac{\sigma^2}{T^2} \sum t \rightarrow \frac{\sigma^2}{2}$$

Therefore, we have

$$\begin{bmatrix} T^{-1/2} \sum \epsilon_t \\ T^{-1/2} \sum \frac{t}{T} \epsilon_t \end{bmatrix} \rightarrow \begin{bmatrix} \sigma^2 & \frac{\sigma^2}{2} \\ \frac{\sigma^2}{2} & \frac{\sigma^2}{3} \end{bmatrix} = \sigma^2 Q$$

Finally, putting everything together,

$$\gamma_T(\hat{\beta} - \beta) \rightarrow N(0, \sigma^2 Q^{-1}).$$

This means the usual OLS t -test and F -test are asymptotically valid, despite at different convergence rates. After all, trend-stationary process is stationary after de-trending. But unit root process is a totally different species.

! Key Point Summary

1. Trend-stationary process vs stochastic-trend process;
2. Applying different convergence rates to OLS estimator;
3. Usual t -test and F -test are still valid.

20 Unit Root Process

A unit root process is characterized by the presence of unit roots in the character equation of its ARMA representation. The simplest unit root process is an AR(1) process with $\phi = 1$:

$$y_t = \phi y_{t-1} + \epsilon_t$$

When $\phi = 1$, it makes each innovation persistent. The effect of past innovations do not fade away no matter how distant they are.

$$y_t = \sum_{j=0}^{\infty} \epsilon_{t-j}$$

The persistence makes the behavior of unit root processes drastically different from stationary processes. Unit root processes hold particular significance among non-stationary processes due to the prevalence of similar behavior in economic or financial time series. For example, stock prices behave a lot like unit root processes (the Random Walk Hypothesis).

The particularity of unit root process makes it a unique class in itself in terms of analytic techniques. The usual OLS estimator and asymptotic normality does not work with unit root processes. If we regress a unit root process on its lags, the OLS estimator is given by

$$\hat{\phi} = \frac{\sum_t y_{t-1} y_t}{\sum_t y_{t-1}^2}$$

We would expect $\sqrt{T}(\hat{\phi} - 1) \rightarrow N(0, \omega^2)$. However, this is not the case. To see this, consider

$$T(\hat{\phi} - 1) = \frac{T^{-1} \sum_t y_{t-1} \epsilon_t}{T^{-2} \sum_t y_{t-1}^2}$$

Assuming Gaussian innovation $u_t \sim N(0, \sigma^2)$, we have

$$y_t = \epsilon_t + \epsilon_{t-1} + \cdots + \epsilon_1 \sim N(0, \sigma^2 t)$$

Therefore, $z_t = \frac{y_t}{\sigma\sqrt{t}} \sim N(0, 1)$. Consider the numerator,




Figure 20.1: Stock market index (black) vs Simulation of a unit root process (red)


$$\begin{aligned}
\frac{1}{T} \sum_{t=1}^T y_{t-1} \epsilon_t &= \frac{1}{T} \sum_{i=1}^T (\epsilon_1 + \cdots + \epsilon_{t-1}) \epsilon_t = \frac{1}{T} \sum_{s < t}^T \epsilon_s \epsilon_t \\
&= \frac{1}{2T} \left[(\epsilon_1 + \cdots + \epsilon_T)^2 - \sum_{t=1}^T \epsilon_t^2 \right] \\
&= \frac{1}{2T} y_T^2 - \frac{1}{2T} \sum_{t=1}^T \epsilon_t^2 \\
&= \frac{\sigma^2}{2} \left(\frac{y_T}{\sigma \sqrt{T}} \right)^2 - \frac{1}{2T} \sum_{t=1}^T \epsilon_t^2 \\
&= \frac{\sigma^2}{2} z_T^2 - \frac{1}{2T} \sum_{t=1}^T \epsilon_t^2
\end{aligned}$$

Since $z_T \sim N(0, 1)$, $z_T^2 \sim \chi^2(1)$. By the LLN, $\frac{1}{T} \sum_{t=1}^T \epsilon_t^2 \rightarrow \mathbb{E}(\epsilon_t^2) = \sigma^2$. Thus,

$$\frac{1}{T} \sum_{t=1}^T y_{t-1} \epsilon_t \rightarrow \frac{\sigma^2}{2} (\chi^2(1) - 1).$$

So the asymptotic distribution of $\hat{\phi}$ is non-Gaussian. The conventional statistical inference no longer make sense. If we simulate the distribution of $\hat{\phi}$ by Monte Carlo, we see it is left-skewed. The negative values are almost twice as likely as positive values, meaning two thirds of the time, the estimated $\hat{\phi}$ will be less than the true value 1. Therefore, the OLS estimate of a unit root process is biased.

```
# Monte Carlo simulation
beta = sapply(1:1000, function(i) {
  y = cumsum(rnorm(200))
  x = dplyr::lag(y)
  coef(lm (y ~ x))[2]
})
hist(beta, freq = FALSE)
```



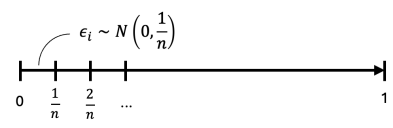
To derive the asymptotic distribution for the unit root process, we need further knowledge of Brownian motions. The idea is derive a continuous version of the unit root process, where each innovation is infinitesimally small. This is the topic of the next section.

21 Brownian Motion

21.1 Continuous random walk

Brownian motion, or **Wiener process**, is the continuous-time extension of a discrete-time random walk. To define the continuous version of a random walk, we cannot simply sum up infinite number of white noises, which will certainly explode. Instead, we chop up a finite interval into infinitely many small intervals, each one corresponding to a tiny Gaussian white noise. The following table shows how a discrete random walk extends to a continuous function.

Table 21.1: Random Walk and Brownian Motion

	Random Walk	Brownian Motion
Innovation	$\epsilon_i \sim \text{WN}(0, 1)$	
Stochastic process	$y_t = \epsilon_1 + \dots + \epsilon_t$	$W(t)_{t \in [0,1]} = \lim_{n \rightarrow \infty} \sum_{i=1}^{nt} \epsilon_i$
Expectation	$\mathbb{E}[y_t] = 0$	$\mathbb{E}[W(t)] = 0$
Variance	$\text{Var}[y_t] = t \text{Var}[\epsilon_i] = t$	$\text{Var}[W(t)] = nt \text{Var}[\epsilon_i] = t$
Quadratic variation	$\mathbb{E} \sum_{i=1}^t (y_i - y_{i-1})^2 = t$	$\int_0^t (dW)^2 = t$

Brownian motion $W(t)$ is a stochastic function. Its realized path is different each time we take a draw from it. But every piece of it follows a tiny Gaussian process. **The function is continuous, but nowhere differentiable.** It is hard to imagine such a function at first glance. But as we proceed, we will appreciate its amazing properties. Despite its path is random, the area under the curve integrates to a well-defined probability distribution. The squared changes (quadratic variation) even sum up to a deterministic constant.

21.2 Some properties

The **quadratic variation** is one of the most important properties of Brownian motions. Intuitively, it says the squared tiny changes sum up to a constant with probability 1 no




Figure 21.1: Multiple Realizations of Brownian Motion

matter which realized path it takes. Smooth functions will not have such property, because $(dX)^2$ diminishes much faster than dX , surely we get $\int_0^t (dX)^2 = 0$. Because Brownian motion fluctuates too much, the squared changes do not diminish away. To see this, imagine summing up the squares of infinitely many small increments up to t :

$$\lim_{n \rightarrow \infty} \sum_{i=1}^{nt} \left[X\left(\frac{i}{n}\right) - X\left(\frac{i-1}{n}\right) \right]^2 = \lim_{n \rightarrow \infty} \sum_{i=1}^{nt} \epsilon_i^2,$$

with $\epsilon_i \sim N\left(0, \frac{1}{n}\right)$. Therefore, $\mathbb{E}(\epsilon_i^2) = \frac{1}{n}$. Let $z_i = \epsilon_i^2$. The sum above can be rewritten as

$$\sum_{i=1}^{nt} z_i = nt \left(\frac{1}{nt} \sum_{i=1}^{nt} z_i \right) \xrightarrow{\text{LLN}} nt \mathbb{E}(z_i) = t.$$

The integral version is the quadratic variation

$$\int_0^t (dW)^2 = t, \tag{21.1}$$

or written in differential form

$$(dW)^2 = dt. \quad (21.2)$$

It should be stressed, W is not differentiable. We use the notation dW , but it is not the same as conventional differentials. The calculus for Brownian motions, the Itô calculus, which will be introduced below, is a different class of calculus specifically designed for stochastic functions. Before we get to that, let's first have a look at some additional properties of Brownian motions.

By Central Limit Theorem, it holds that

$$\frac{1}{\sqrt{nt}}W(t) = \sqrt{nt}\frac{1}{nt}\sum_{i=1}^{nt}\epsilon_i \rightarrow N\left(0, \frac{1}{n}\right)$$

Therefore,

$$W(t) \sim \sqrt{nt}N\left(0, \frac{1}{n}\right) \sim N(0, t); \quad (21.3)$$

It follows that, for any $r < s$,

$$W(s) - W(t) \sim N(0, s - t); \quad (21.4)$$

As s and t become arbitrarily close, we have

$$dW \sim N(0, dt). \quad (21.5)$$

In essence, Brownian motion is the accumulation of tiny independent Gaussian innovations. We give the formal definition of Brownian motions below.

Brownian motion

A Brownian motion (Wiener process) is a stochastic function $W(t)$ such that

1. $W(0) = 0$;
2. For $0 \leq t \leq s \leq 1$, $W(s) - W(t) \sim N(0, s - t)$;
3. For any realization, $W(t)$ is continuous in t .

Brownian motions are frequently used to model stock returns. For a fixed horizon T , the returns are normally distributed, with volatility scaled by \sqrt{T} . And the returns over different periods are independent, which means no predictability from past returns to future returns.

21.3 Itô calculus

Lemma 21.1. *Let $F(W)$ be a “smooth” function of a Brownian motion $W(t)$. Then*

$$dF = F'dW + \frac{1}{2}F''dt.$$

Proof. For an informal proof, it immediately follows from the Taylor expansion

$$F(W(t+h)) - F(W(t)) = F'(W(t+h) - W(t)) + \frac{1}{2}F''(W(t+h) - W(t))^2 + \dots$$

As $h \rightarrow 0$,

$$dF = F'dW + \frac{1}{2}F''(dW)^2.$$

By Equation 21.2, $(dW)^2 = dt$. Therefore,

$$dF = F'dW + \frac{1}{2}F''dt.$$

□

This formula is known as the **Itô's lemma**, which is the key equation of Itô calculus. Note that how this differs from the differential formula for normal functions: $dF = F'dW$. The *second-order* term does not disappear precisely because the quadratic variation does not go to zero.

Example 21.1. $F(W) = W^2 \implies dF = 2WdW + dt$

Example 21.2. Let's do a more involved example to familiar ourselves with the Itô's lemma, especially how the second-order differentiation of the Brownian motions plays out in computation.

We like to model the continuous-time stock price with Brownian motions. Let S_t be the stock price at time t . Assume the behavior of S_t follows a stochastic differential equation:

$$\frac{dS_t}{S_t} = \mu dt + \sigma dW$$

That is, the percentage of S_t is a continuous time random walk with drift μ and volatility σ . Let's compute the log-return of the stock over the horizon T :

$$R_T = \ln(S_T) - \ln(S_0) = f(S_T)$$

Apply second-order Taylor expansion:

$$\begin{aligned}
dR_T &= f'(S_T) dS_T + \frac{1}{2} f''(S_T) (dS_T)^2 \\
&= \frac{1}{S_T} dS_T - \frac{1}{2} \frac{1}{S_T^2} (dS_T)^2 \\
&= \frac{dS_T}{S_T} - \frac{1}{2} \left(\frac{dS_T}{S_T} \right)^2 \\
&= (\mu dt + \sigma dW) - \frac{1}{2} (\mu dt + \sigma dW)^2 \\
&= (\mu dt + \sigma dW) - \frac{1}{2} (\mu^2 dt^2 + 2\mu\sigma dt dW + \sigma^2 dW^2) \\
&\rightarrow (\mu dt + \sigma dW) - \frac{1}{2} \sigma^2 dt \\
&= \left(\mu - \frac{1}{2} \sigma^2 \right) dt + \sigma dW
\end{aligned}$$

The second last step holds because, as $dt \rightarrow 0$, the terms dt^2 and $dt dW$ tend to zero faster than dW^2 . The only term left is $dW^2 = dt$. If we define integral as the inverse of differentiation, we have

$$\begin{aligned}
R_T &= \int_0^T dR_T = \left(\mu - \frac{1}{2} \sigma^2 \right) \int_0^T dt + \sigma \int_0^T dW \\
&= \left(\mu - \frac{1}{2} \sigma^2 \right) T + \sigma \sqrt{T} \epsilon_T
\end{aligned}$$

where ϵ_T is a standard normal variable. The model tells us that the log-return of a stock over a fixed horizon of T is normally distributed with mean $(\mu - \sigma^2/2)T$ and standard deviation of $\sigma\sqrt{T}$. Everything looks familiar, except the Ito's term, $\sigma^2/2$, which comes from the non-zero second-order differential dW^2 . The famous Black-Scholes formula for option pricing is derived from this model.

Definition 21.1. Stochastic integrals as the reverse operation of the stochastic differentiation. Note that we change $W(t)$ to $W(s)$ when it enters as the integrand.

1. $\int_0^t dW = W(t)$;
2. $F(W(t)) = \int_0^t f(W(s)) dW$, if $dF = f dW$;

3. $F(t, W(t)) = \int_0^t f(s, W(s))dW + \int_0^t g(s, W(s))ds$, if $dF = f dW + g dt$.

Example 21.3. Given $dW^2 = 2WdW + dt$, take integral on both sides:

$$\begin{aligned} W^2(t) &= 2 \int_0^t W(s)dW + \int_0^t ds \\ \implies \int_0^t WdW &= \frac{1}{2}[W^2(t) - t]. \end{aligned}$$

Setting $t = 1$, it follows that

$$\int_0^1 WdW = \frac{1}{2}[W^2(1) - 1].$$

Note that $W(1) \sim N(0, 1)$. So $W^2(1) \sim \chi^2(1)$ with expectation 1. Thus $\int_0^1 WdW$ is centered at 0 but skewed.

Example 21.4. Given $d(tW) = Wdt + t dW$ (verify this with Ito's lemma), we have

$$\begin{aligned} tW(t) &= \int_0^t W(s)ds + \int_0^t s dW \\ \implies \int_0^t W(s)ds &= tW(t) - \int_0^t s dW \\ &= t \int_0^t dW - \int_0^t s dW \\ &= \int_0^t (t - s)dW. \end{aligned}$$

Making Sense of Itô Calculus

Let $F(t)$ be a continuous-time trading strategy that holds an amount of $F(t)$ of a stock at time t . The stock price is a Brownian motion $W(t)$. dW represents the movement of stock price. Consider the Itô's integral

$$Y_t = \int_0^t F dW$$

The stochastic integral represents the payoff of the trading strategy up to time t . Note that the integral always evaluates *at the left*. So the strategy can only make decision based on available information.

Proposition 21.1. *Let $W(t)$ be a Brownian motion. Let $f(t)$ be a nonrandom function of time. Then*

1. $\mathbb{E} \left[\int_0^t f(s) dW \right] = 0;$
2. $\mathbb{E} \left[\left(\int_0^t f(s) dW \right)^2 \right] = \mathbb{E} \left[\int_0^t f^2 ds \right]$ (Itô isometry);
3. $\int_0^t f(s) dW \sim N \left(0, \int_0^t f^2 ds \right).$

If $f(t)$ represents a trading strategy and the stock price follows a Brownian motion, the theorem tells us the expected payoff of this strategy is zero; more precisely, the payoff follows a Gaussian distribution.

Example 21.5. Following the last example,

$$\int_0^t W ds = \int_0^t (t - s) dW$$

$f(s) = t - s$ is a nonrandom function, apply the theorem above

$$\text{var} \left[\int_0^t W ds \right] = \int_0^t (t - s)^2 ds = \frac{1}{3} t^3$$

Therefore,

$$\int_0^t W(s) ds \sim N \left(0, \frac{1}{3} t^3 \right).$$

Thus, the integral, the area under the curve of a Brownian motion, follows a Gaussian distribution.

21.4 Unit root process

Consider a unit root process,

$$y_t = y_{t-1} + \epsilon_t = \sum_{j=1}^t \epsilon_j,$$

where $\epsilon_t \sim \text{WN}(0, 1)$, $t = 1, 2, \dots, T$. It is not surprising to see the unit root process converges to Brownian motion if stabilizing it by $T^{-1/2}$:

$$\begin{aligned}
\frac{1}{\sqrt{T}}y_t &= \frac{1}{\sqrt{T}} \sum_{j=1}^t \epsilon_j \\
&= \frac{1}{\sqrt{T}} \sum_{j=1}^{Tr} \epsilon_j \quad (r = t/T) \\
&= \sqrt{r} \left(\frac{1}{\sqrt{Tr}} \sum_{j=1}^{Tr} \epsilon_j \right) \\
&\rightarrow \sqrt{r}N(0, 1) \sim N(0, r) \quad (\text{by CLT}) \\
&\rightarrow W(r) \quad (\text{by definition}).
\end{aligned}$$

If ϵ_t has variance σ^2 , we would have $T^{-1/2}y_t \rightarrow \sigma W(t/T)$. Note if y_t is stationary, y_t will not deviate too far from $\mathbb{E}(y_t)$, we would have $T^{-1/2}y_t \rightarrow 0$.

Now let's consider the behaviour of the mean: $\bar{y} = \frac{1}{T} \sum_{t=1}^T y_t$. Define $\xi_T(r) = T^{-1/2}y_t$, where $r = t/T$. We have

$$\begin{aligned}
\frac{1}{\sqrt{T}}\bar{y} &= \frac{1}{\sqrt{T}} \left(\frac{1}{T} \sum_{t=1}^T y_t \right) \\
&= \frac{1}{T} \sum_{t=1}^T \left(\frac{1}{\sqrt{T}} y_t \right) = \frac{1}{T} \sum_{t=1}^T \xi \left(\frac{t}{T} \right) \\
&= \Delta r \sum_{r=0}^1 \xi(r) \quad (r = t/T, \Delta r = 1/T) \\
&\rightarrow \int_0^1 \sigma W(r) dr \quad (\Delta r \rightarrow 0)
\end{aligned}$$

Remember for stationary process, we would have $\bar{y} \rightarrow \mathbb{E}(y_t)$. With unit root process, the mean no longer converges to a constant, but to a distribution $\int_0^1 W(r) dr \sim N(0, \frac{1}{3})$.

For higher orders of y_t , we have

$$\begin{aligned}
\frac{1}{T^2} \sum_{t=1}^T y_t^2 &= \frac{1}{T} \sum_{t=1}^T \left(\frac{1}{\sqrt{T}} y_t \right)^2 \\
&= \frac{1}{T} \sum_{t=1}^T \left[\xi \left(\frac{t}{T} \right) \right]^2 \\
&= \sum_{r=0}^1 [\xi(r)]^2 \Delta r \\
&\rightarrow \int_0^1 \sigma^2 W^2(r) dr
\end{aligned}$$

By continuous mapping theorem, it can be shown, in general

$$\frac{1}{T^{1+k/2}} \sum_{i=1}^T y_t^k \rightarrow \sigma^k \int_0^1 W^k(r) dr.$$

Consider the numerator of the OLS estimator of the unit root process,

$$\begin{aligned}
\sum_{t=1}^T y_{t-1} \epsilon_t &= \sum_{i=1}^T (\epsilon_1 + \cdots + \epsilon_{t-1}) \epsilon_t = \sum_{s < t}^T \epsilon_s \epsilon_t \\
&= \frac{1}{2} \left[(\epsilon_1 + \cdots + \epsilon_T)^2 - \sum_{t=1}^T \epsilon_t^2 \right] \\
&= \frac{1}{2} y_T^2 - \frac{1}{2} \sum_{t=1}^T \epsilon_t^2
\end{aligned}$$

Divide it by T , we have

$$\begin{aligned}
\frac{1}{T} \sum_{t=1}^T y_{t-1} \epsilon_t &= \frac{1}{2} \left(\frac{1}{\sqrt{T}} y_T \right)^2 - \frac{1}{2T} \sum_{t=1}^T \epsilon_t^2 \\
&= \frac{1}{2} [\xi^2(1) - \hat{\sigma}^2] \\
&\rightarrow \frac{1}{2} [\sigma^2 W^2(1) - \sigma^2] \\
&= \frac{1}{2} \sigma^2 [W^2(1) - 1] \\
&= \sigma^2 \int_0^1 W dW.
\end{aligned}$$

We summarize the important results below.

Key Rules Summary

1. $\int_0^T W dt = N(0, \frac{T^3}{3})$
2. $\int_0^T W dW = \frac{1}{2}[W^2(T) - T]$
3. $\frac{1}{T^{3/2}} \sum_{t=1}^T y_t \rightarrow \sigma \int_0^1 W(r) dr$
4. $\frac{1}{T^{1+k/2}} \sum_{i=1}^T y_t^k \rightarrow \sigma^k \int_0^1 W^k(r) dr$
5. $\frac{1}{T} \sum_{t=1}^T y_{t-1} \epsilon_t \rightarrow \sigma^2 \int_0^1 W dW$
6. $\frac{1}{T^2} \sum_{t=1}^T y_{1t} y_{2t} \rightarrow \sigma_1 \sigma_2 \int_0^1 W_1(r) W_2(r) dr$

The last rule was given without proof, as we will need it in the following chapters.

22 Unit Root Process (contd)

22.1 Univariate case

We now have all the ingredient to further analyse the unit root process

$$y_t = \phi y_{t-1} + \epsilon_t,$$

where $\phi = 1$, and its OLS estimator

$$T(\hat{\phi} - 1) = \frac{T^{-1} \sum_t y_{t-1} \epsilon_t}{T^{-2} \sum_t y_{t-1}^2}.$$

We have shown that

$$\begin{aligned} T^{-1} \sum_t y_{t-1} \epsilon_t &\rightarrow \sigma^2 \int_0^1 W dW = \frac{\sigma^2}{2} (W^2(1) - 1) \\ T^{-2} \sum_t y_{t-1}^2 &\rightarrow \sigma^2 \int_0^1 W^2 ds \end{aligned}$$

Therefore,

$$T(\hat{\phi} - 1) \rightarrow \frac{\int_0^1 W dW}{\int_0^1 W^2 ds}.$$

$\int W dW$ is centered around 0, meaning $\hat{\phi}$ is consistent for large samples. But it is biased in small samples. Moreover, the distribution is not Gaussian, rendering all conventional t -test or F -test meaningless. We contrast the properties of stationary processes and unit root processes below.

Table 22.1: Stationary AR(1) process vs unit root process

	Stationary	Unit Root
Model	$y_t = \phi y_{t-1} + \epsilon_t$	$y_t = y_{t-1} + \epsilon_t$

	Stationary	Unit Root
Asymptotic distribution of $\hat{\phi}$	$\sqrt{T}(\hat{\phi} - \phi) \rightarrow N(0, 1 - \phi^2)$	$\sqrt{T}(\hat{\phi} - 1) \rightarrow \frac{\int W dW}{\int W^2 dt}$
Asymptotic distribution of t -statistics	$t \rightarrow N(0, 1)$	$t \rightarrow \frac{\int W dW}{\sqrt{\int W^2 dt}}$

22.2 Spurious regression

We now dive deeper into the nature of spurious regression problem presented at the beginning of the chapter. We formulate the problem as below. Suppose

$$y_t = \alpha + \beta x_t + u_t,$$

where y_t and x_t are unit root processes and there does not exist (α, β) such that the residual u_t is stationary. In this case, OLS is likely to produce spurious result: even if y_t is completely unrelated to x_t , the estimated value of $\hat{\beta}$ is likely to appear to be statistically significantly different from zero.

Spurious regression happens when

1. Dependent/independent variables are non-stationary;
2. The residual is non-stationary for all possible values of the coefficient vector.

To understand why this happens, consider the OLS estimator:

$$\hat{b} = \begin{bmatrix} \hat{\alpha} \\ \hat{\beta} \end{bmatrix} = \begin{bmatrix} T & \sum x_t \\ \sum x_t & \sum x_t^2 \end{bmatrix}^{-1} \begin{bmatrix} \sum y_t \\ \sum x_t y_t \end{bmatrix}$$

To account for different convergent speed, similar to the trend-stationary case, we multiply the estimators by a matrix,

$$\begin{aligned} \begin{bmatrix} \sqrt{T}^{-1} & \\ & 1 \end{bmatrix} \begin{bmatrix} \hat{\alpha} \\ \hat{\beta} \end{bmatrix} &= \begin{bmatrix} \sqrt{T}^{-1} & \\ & 1 \end{bmatrix} \begin{bmatrix} T & \sum x_t \\ \sum x_t & \sum x_t^2 \end{bmatrix}^{-1} \begin{bmatrix} \sqrt{T}^{-1} & \\ & 1 \end{bmatrix}^{-1} \begin{bmatrix} \sqrt{T}^{-1} & \\ & 1 \end{bmatrix} \begin{bmatrix} \sum y_t \\ \sum x_t y_t \end{bmatrix} \\ &= \begin{bmatrix} 1 & T^{-3/2} \sum x_t \\ T^{-3/2} \sum x_t & T^{-2} \sum x_t^2 \end{bmatrix}^{-1} \begin{bmatrix} \sum T^{-3/2} y_t \\ T^{-2} \sum x_t y_t \end{bmatrix} \\ &\rightarrow \begin{bmatrix} 1 & \int W_X dt \\ \int W_X dt & \int W_X^2 dt \end{bmatrix}^{-1} \begin{bmatrix} \int W_Y dt \\ \int W_X W_Y dt \end{bmatrix} \end{aligned}$$

This means, $\hat{\alpha}$ actually diverges. Because it needs to be divided by \sqrt{T} to be able to converge to a stable distribution, rather than being multiplied by a stabilizing factor. $\hat{\beta}$ converges, but it is not consistent. If there is no α , we would have

$$\hat{\beta} = \frac{\sum x_t y_t}{\sum x_t^2} \rightarrow \frac{\int W_X W_Y dt}{\int W_X^2 dt},$$

which is inconsistent. So we won't get zero even in very large samples.

The OLS estimate of the variance of u_t also diverges. It needs to be divided by T to converge:

$$\begin{aligned} \frac{1}{T} \hat{\sigma}^2 &= \frac{1}{T^2} \sum (y_t - \hat{\beta} x_t)^2 \\ &= \frac{1}{T^2} \sum y_t^2 - 2\hat{\beta} \frac{1}{T^2} \sum x_t y_t + \hat{\beta}^2 \frac{1}{T^2} \sum x_t^2 \\ &\rightarrow \int W_Y^2 dt - 2\hat{\beta} \int W_X W_Y dt + \hat{\beta}^2 \int W_X^2 dt \\ &\rightarrow \int W_Y^2 dt - \frac{(\int W_X W_Y dt)^2}{\int W_X^2 dt}. \end{aligned}$$

The t or F statistics also diverge. t -stat has to be divided by \sqrt{T} to converge; F -stat needs to be divided by T to converge.

$$t = \frac{\hat{\beta}}{\hat{\sigma}} \sqrt{\sum x_t^2} = \sqrt{T} \frac{\hat{\beta}}{\sqrt{T^{-1} \hat{\sigma}^2}} \sqrt{T^{-2} \sum x_t^2} \rightarrow \sqrt{T} \cdot C$$

Thus, as sample size T grows, t -test will appear very large and significant, despite y_t and x_t are completely independent.

22.3 Cures for spurious regression

1. Include lagged values of both dependent and independent variables in the regression:

$$y_t = \alpha + \phi y_{t-1} + \beta x_t + \gamma x_{t-1} + u_t$$

Now there exists a coefficient vector $[\phi, \beta, \gamma] = [1, 0, 0]$ such that u_t is $I(0)$ stationary. In this case, OLS yields consistent estimates for all the coefficients. t -test converges to Gaussian, though F -test of joint hypotheses has non-standard asymptotic distribution. We will come back to this point later.

2. Difference the data to stationary:

$$\Delta y_t = \alpha + \beta \Delta x_t + u_t$$

Because Δy_t and Δx_t are all $I(0)$. Standard OLS is valid.

3. Estimate with Cochrane-Orcutt adjustment for first-order correlations in the residuals. This method is asymptotically equivalent to the second method.

22.4 Summary

! Key Point Summary

1. The OLS estimator for unit root coefficient converges to non-standard distributions involving Brownian motions. Thus, standard statistical inferences are meaningless.
2. Regressing unit root processes lead to spurious results, because the diverging behavior of t -stats makes artificially significant values.
3. Include lagged values or difference the data to stationary when working with non-stationary time series.

23 Unit Root Test

A presence of unit root necessitates special treatment in empirical applications. Therefore it is of vital importance to pre-test the existence of unit root.

However, there is no clear cut between stationary and unit root processes for finite samples. Consider an AR(1) process with $\phi = 0.999$, which is a stationary process that behaves very close to a unit root process. In other words, unit root and stationary processes differ in their implications at *infinite* time horizons, but for any *finite* number of observations, there is always a representation from either class of models that could account for all the observed features of the data. So it is *not* possible to tell whether the DGP is stationary or not. We can formulate testable hypothesis only if we were willing to restrict the class of models being considered. Suppose we were committed to an AR(1) model: $y_t = \phi y_{t-1} + \epsilon_t$. The hypothesis $\phi = 1$ is definitely testable.

23.1 Dickey-Fuller Test

Consider an AR(1) process

$$y_t = \phi y_{t-1} + u_t,$$

assuming no series correlation in the innovations $u_t \sim IID(0, \sigma^2)$. We have shown that under the hypothesis $\phi = 1$,

$$T(\hat{\phi} - 1) \rightarrow \frac{\int W dW}{\int W^2 dt}.$$

We can the hypothesis $H_0 : \phi = 1$ utilizing this distribution. The critical values can be obtained by Monte Carlo simulations. The test was proposed by Fuller (1976).

The Dickey-Fuller tests involve three sets of equations depending whether a drift or trend is included, assuming *iid* innovations.

$$\begin{aligned}\Delta y_t &= \gamma y_{t-1} + u_t \\ \Delta y_t &= \alpha_0 + \gamma y_{t-1} + u_t \\ \Delta y_t &= \alpha_0 + \gamma y_{t-1} + \alpha_2 t + u_t\end{aligned}$$

Testing $\phi = 1$ is equivalent to testing $\gamma = 0$. The critical values depends on the form of the regression and the sample size (including a drift or trend results in different limiting distributions for γ).

Table 23.1: Critical values of Dickey-Fuller tests

Model	Hypothesis	95%	99%
Default	$\gamma = 0$	-1.95	-2.60
With drift	$\gamma = 0$	-2.89	-3.51
	$\alpha_0 = \gamma = 0$	4.71	6.70
With drift and trend	$\gamma = 0$	-3.45	-4.04
	$\gamma = \alpha_2 = 0$	6.49	8.73
	$\alpha_0 = \gamma = \alpha_2 = 0$	4.88	6.50

23.2 Augmented Dickey-Fuller Test

The assumption that ϵ_t being uncorrelated is too strong for empirical applications. Suppose the data is generated by an $AR(p)$ process with an unit root,

$$\begin{aligned} a(L)y_t &= \epsilon_t \\ (1 - L)y_t &= \underbrace{b^{-1}(L)\epsilon_t}_{u_t} \end{aligned}$$

where $a(L) = (1 - L)b(L)$ in which $b(L)$ is stationary. In this case, u_t will be autocorrelated. In empirical works, it is more reasonable to assume u_t being serially correlated.

If we difference y_t once, we have

$$\begin{aligned} b(L)\Delta y_t &= \epsilon_t \\ y_t &= y_{t-1} + \sum_{j=1}^p \beta_j \Delta y_{t-j} + \epsilon_t \end{aligned}$$

This motivates Dickey-Fuller tests with lags $\{\Delta y_{t-j}\}$. This is called augmented Dickey-Fuller test. The set of equations change to

$$\begin{aligned}\Delta y_t &= \gamma y_{t-1} + \sum_{j=1}^p \beta_j \Delta y_{t-j} + u_t \\ \Delta y_t &= \alpha_0 + \gamma y_{t-1} + \sum_{j=1}^p \beta_j \Delta y_{t-j} + u_t \\ \Delta y_t &= \alpha_0 + \gamma y_{t-1} + \alpha_2 t + \sum_{j=1}^p \beta_j \Delta y_{t-j} + u_t\end{aligned}$$

The coefficients on Δy_{t-j} converge to Gaussian. The coefficient on y_{t-1} converges to non-standard distribution. The critical values are unchanged with lags are included.

23.3 Phillips-Perron Test

Another approach to test unit root is proposed by Phillips and Perron (1988), which also assumes autocorrelated errors. Our Brownian motion theories derived from *iid* innovations can be extended to autocorrelated innovations:

$$\xi_T(r) = \frac{1}{\sqrt{T}} \sum_{t=1}^{Tr} u_t \rightarrow \omega W(r)$$

where $\omega^2 = \sum_{-\infty}^{\infty} \gamma_j$ is the long-run variance for the autocorrelated process $\{u_t\}$.

But, with autocorrelated errors, the limiting distribution of $\hat{\phi}$ is slightly different, since

$$\begin{aligned}\frac{1}{T} \sum_t y_{t-1} u_t &= \frac{1}{2T} y_T^2 - \frac{1}{2T} \sum_t u_t^2 \\ &\rightarrow \frac{1}{2} [\omega^2 W^2(1) - \sigma_u^2] \\ &\rightarrow \omega^2 \int W dW + \frac{\omega^2 - \sigma_u^2}{2}\end{aligned}$$

where $\frac{\omega^2 - \sigma_u^2}{2}$ is a “nuisance” parameter. The Phillips and Perron proposed a test statistics correcting the nuisance parameter:

$$T(\hat{\phi} - 1) + \frac{\frac{1}{2}\hat{\omega}^2 - \hat{\sigma}_u^2}{\frac{1}{T^2} \sum y_t^2} \rightarrow \frac{\int W dW}{\int W^2 dt}$$

where $\hat{\omega}^2$ can be estimated by Newey-West, $\hat{\sigma}_u^2$ is the estimated variance of the residuals. After the correction, the unit root test can be applied to processes with autocorrelated errors.

24 Cointegration

We have shown that regressing on non-stationary time series might lead to spurious regressions that produce nonsensible large t -values. The conventional asymptotic theorems no longer hold for series with very high persistence such as unit root processes. But unit roots are not always an enemy, sometimes it can be a friend, as in the case of cointegration. In this case, the OLS estimates are super-consistent. They are consistent even when there is an endogeneity issue (which is amazing!)

24.1 Cointegration and super-consistency

Consider a regression with two random walks x_t and y_t :

$$y_t = \beta x_t + e_t$$

where e_t is stationary. In this case, y_t and x_t are **cointegrated**, because the linear combination of the two $I(1)$ processes becomes $I(0)$. If this is the case, we no longer have a spurious regression. In fact, $\hat{\beta}$ is not only consistent, but **super-consistent**. Consider the OLS estimator

$$T(\hat{\beta} - \beta) = \frac{T^{-1} \sum x_t e_t}{T^{-2} \sum x_t^2} \rightarrow \frac{\int W_1 dW_2}{\int W_1^2 dt}.$$

Note that $\int W_1 dW_2$ is centered at zero. Hence $\hat{\beta}$ is consistent, despite the distribution is non-Gaussian. It converges at rate T , faster than the standard case \sqrt{T} . So it is called super-consistency.

The argument extends to general regressions with multiple regressors. As long as there exists a vector of coefficients that makes the non-stationary variables cointegrated, the OLS estimator is super-consistent. For example,

$$y_t = \phi y_{t-1} + \beta x_t + e_t$$

If y_t is $I(1)$, then $y_t - y_{t-1}$ is $I(0)$ (cointegrated with itself). Therefore, $[0, 1]$ is the cointegration vector that makes e_t stationary even if x_t is not cointegrated with y_t . In this case, OLS estimates for $\hat{\phi}$ and $\hat{\beta}$ will be super-consistent.

The intuition of super-consistency is that OLS minimizes squared residuals. If the coefficients deviate from the cointegration vector, \hat{e}_t^2 would diverge. \hat{e}_t^2 is minimized only when the coefficients coincide with the cointegration vector. That makes OLS converges even faster.

The super-consistency is so strong, that the OLS estimators for cointegrated variables are consistent even when there is endogeneity problem. We demonstrate this with an example. Suppose x_t follows a random walk, and y_t cointegrates with x_t :

$$\begin{aligned}x_t &= x_{t-1} + u_t \\y_t &= \beta x_t + e_t\end{aligned}$$

Assume u_t and e_t are correlated with $\mathbb{E}(e_t^2) = \mathbb{E}(u_t^2) = 1$ and $\text{cov}(e_t, u_t) = \phi$. For simplicity, also assume $e_t = \phi u_t + \sqrt{1 - \phi^2} \eta_t$ where η_t is *iid* standard normal. As e_t is correlated with x_t through u_t , there is clearly an endogeneity problem. The OLS estimator is given by

$$T(\hat{\beta} - \beta) = \frac{\frac{1}{T} \sum x_t e_t}{\frac{1}{T^2} \sum x_t^2} = \frac{\frac{1}{T} \sum x_{t-1} e_t + \frac{1}{T} \sum u_t e_t}{\frac{1}{T^2} \sum x_t^2}$$

where

$$\begin{aligned}\frac{1}{T} \sum x_{t-1} e_t &= \frac{\phi}{T} \sum x_{t-1} u_t + \frac{\sqrt{1 - \phi^2}}{T} \sum x_{t-1} \eta_t \\&\rightarrow \phi \int W_1 dW_1 + \sqrt{1 - \phi^2} \int W_1 dW_2\end{aligned}$$

Therefore,

$$T(\hat{\beta} - \beta) \rightarrow \frac{\phi \int W_1 dW_1 + \sqrt{1 - \phi^2} \int W_1 dW_2 + \phi}{\int W_1^2 dt}.$$

When $\phi \neq 0$, that is endogeneity exists, the limiting distribution is shifted. As a result $\hat{\beta}$ has a finite sample bias of order $\frac{1}{T}$. But as $T \rightarrow \infty$, the estimator is consistent. Figure 24.1 demonstrates the convergence of $\hat{\beta}$ as sample size increases (the true $\beta = 0.5$).




Figure 24.1: Superconsistency with cointegrated variables (even with endogeneity)

24.2 Inference under cointegration

We have seen the limiting distributions of persistent regressors could be non-standard. But luckily, in many cases, we can still get Gaussian distributions. This is hard to believe. But it can be shown. Consider a regression of the **canonical form** — that is a regression with four types of regressors: stationary $I(0)$, non-stationary $I(1)$, constant and trend. It can be shown any regression can be rewritten in the canonical form.

$$y_t = \gamma z_t + e_t$$

where

$$z_t = \begin{bmatrix} F_1(L) & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ F_2(L) & G & H & 0 \\ F_3(L) & T & K & 1 \end{bmatrix} \begin{bmatrix} \epsilon_t \\ 1 \\ \eta_t \\ t \end{bmatrix}$$

and e_t is stationary. Consider the OLS estimator: $\hat{\gamma} = \gamma + (Z'Z)^{-1}Z'e$ with a scaling matrix

$$Q = \begin{bmatrix} \sqrt{T}\mathbf{I}_{k_1} & & & \\ & \sqrt{T} & & \\ & & T\mathbf{I}_{k_3} & \\ & & & T^{3/2} \end{bmatrix}$$

Multiply them together,

$$\begin{aligned} Q(\hat{\gamma} - \gamma) &= (Q^{-1}Z'ZQ^{-1})^{-1}Q^{-1}Z'e \\ &= \begin{bmatrix} \text{const}_{k_1 \times k_1} & 0 & 0 & 0 \\ 0 & \text{Functions of} & & \\ 0 & \text{Brownian} & & \\ 0 & \text{motions} & & \end{bmatrix}^{-1} \begin{bmatrix} N(0, ?) \\ N(0, ?) \\ ? \int W dW? \\ N(0, ?) \end{bmatrix} \end{aligned}$$

We don't care the specific functional forms of the converged distribution. What matters is the speed of convergence. Note that

$$\begin{aligned} \sqrt{T}(\hat{\gamma}_1 - \gamma_1) &\rightarrow N(0, ?) \\ \sqrt{T}(\hat{\gamma}_2 - \gamma_2) &\rightarrow \text{Something 2} \\ T(\hat{\gamma}_3 - \gamma_3) &\rightarrow \text{Something 3} \\ T^{3/2}(\hat{\gamma}_4 - \gamma_4) &\rightarrow \text{Something 4} \end{aligned}$$

Constant and stationary regressors have the slowest converging speed \sqrt{T} . Only stationary regressors converge to Gaussian distribution.

Now consider a general regression,

$$y_t = \beta \mathbf{x}_t + e_t$$

where e_t is stationary. Sims (1990) shows that we can always find a linear combination $\mathbf{D}\mathbf{x}_t = \mathbf{z}_t$ that transforms the regression into a canonical form

$$y_t = \gamma \mathbf{z}_t + e_t$$

where $\gamma = \beta \mathbf{D}^{-1}$. This means, if a component of $\hat{\beta}$ can be written as a linear combination of $\hat{\gamma}_1, \hat{\gamma}_3, \hat{\gamma}_4$, its distribution will be dominated by the behavior of $\hat{\gamma}_1$ due to its slower convergence speed. As such, it will behave like asymptotic normal and converge at speed \sqrt{T} .

In essence, the coefficients that can be represented as a linear combination involving stationary regressors will be asymptotically normal and converge at rate \sqrt{T} . Consider an example,


$$y_t = \alpha + \rho_1 y_{t-1} + \rho_2 y_{t-2} + \beta_1 x_t + \beta_2 x_{t-1} + e_t$$

where y_t and x_t are $I(1)$. The regression can be rewritten as

$$\begin{aligned} y_t &= \alpha + \rho_1 \Delta y_{t-1} + (\rho_1 + \rho_2) y_{t-2} + \beta_1 \Delta x_t + (\beta_1 + \beta_2) x_{t-1} + e_t \\ &= \alpha + \rho_1 \Delta y_{t-1} + \lambda y_{t-2} + \beta_1 \Delta x_t + \delta x_{t-1} + e_t \end{aligned}$$

in which, ρ_1 and β_1 are coefficients on stationary regressors, therefore converging to Gaussian; $\rho_2 = \lambda - \rho_1$ and $\beta_2 = \delta - \beta_1$ are linear combinations involving coefficients on stationary regressors, whose distributions will be dominated by that of ρ_1 and β_1 . Hence, all coefficients will have asymptotically normal distributions. Standard inference applies. We verify the claim with a Monte Carlo simulation.

```
library(dynlm)
beta = sapply(1:1000, function(i) {
  x = arima.sim(list(order=c(1,1,0),ar=.5), 200)
  y = arima.sim(list(order=c(0,1,1),ma=.7), 200)
  coef(dynlm(y ~ L(y,1:2) + L(x,0:1)))
}) |> t()
{
  par(mfrow=c(2,2), mar=c(2,2,2,2))
  hist(beta[, 'L(y, 1:2)1'], freq=F, main="y(t-1)")
  hist(beta[, 'L(y, 1:2)2'], freq=F, main="y(t-2)")
  hist(beta[, 'L(x, 0:1)0'], freq=F, main="x(t)")
  hist(beta[, 'L(x, 0:1)1'], freq=F, main="x(t-1)")
}
```



24.3 Conclusion

Key Takeaways

1. Regressions with non-stationary regressors are likely to be spurious regressions. Do not regress two non-stationary series unless they are cointegrated.
2. Including lags of dependent and independent variables protects you from spurious regressions.
3. Persistence in time series makes statistical inference complicated. Sometimes standard inference works, but not always.

Part V

Vector Autoregression

25 System of Equations

We have discussed how to estimate the effect of one economic variable on another, and the assumptions on which the estimate would have a (dynamic) causal interpretation. But one equation is often inadequate to characterize the economy, as it does not take into account the feedback between economic variables. For example, an oil price shock would have impact on the price levels, which would trigger adjustments in the monetary policy, which would further exert impact on price levels, real output and so on. To capture the intertwined relationships, it would require a **system of equations**.

Consider an example of a backward-looking Keynesian system:

$$\begin{aligned}y_t &= \phi y_{t-1} - \psi(r_{t-1} - \phi_{t-1}) + \epsilon_t^{IS} \\ \pi_t &= \delta \pi_{t-1} + \kappa(y_{t-1} - y_t^n) + \epsilon_t^S \\ r_t &= \beta \pi_t + \gamma(y_t - y_t^n) + \epsilon_t^{MP}\end{aligned}$$

The first equation is the IS curve, which states the negative relationship between output and real interest rate. ϵ_t^{IS} is a structural shock of investment-saving decisions that moves the IS curve. We call it structural shock, because it is associated with a structural meaning, not a mere residual from a regression. The second equation describes the Phillips curve, which postulates a positive correlation between inflation and output gap (where y_t^n is the potential output level). ϵ_t^S is the supply shock, which originates from exogenous supply conditions (such as weather), that could also affect inflation. The third equation is the Taylor's rule for monetary policy, which sets the interest rate in response to inflation and output gap. ϵ_t^{MP} is the monetary policy shock, which is the unpredictable part of the monetary policy decision making.

The set of equations are called structural equations, in a sense that they describe the structure of the economy according to some economic theories (particularly the Keynesian theory). These equations were very popular in 70s and 80s until Sims (1980) questioned their validity. The fact is, these equations impose a lot of restrictions on the relationships between the variables. For example, why output responds to real interest rate but not inflation? Why interest rate does not enter the equation of inflation? Yes, the equations are justified by the theory. But who knows the theory is correct? In reality, economic variables influence each other, often in a way unknown to theorists. So why not model the economy unrestrictively and let the data tell us the relationships between the variables?

$$\begin{aligned}
y_t &= \phi_{11}y_{t-1} + \phi_{12}\pi_{t-1} + \phi_{13}r_{t-1} + \cdots \\
\pi_t &= \phi_{21}y_{t-1} + \phi_{22}\pi_{t-1} + \phi_{23}r_{t-1} + \cdots \\
r_t &= \phi_{31}y_{t-1} + \phi_{32}\pi_{t-1} + \phi_{33}r_{t-1} + \cdots
\end{aligned}$$

This gives rise to a vector autoregressive system:

$$\begin{bmatrix} y_t \\ \pi_t \\ r_t \end{bmatrix} = \sum_{j=1}^p \begin{bmatrix} \phi_{j,11} & \phi_{j,12} & \phi_{j,13} \\ \phi_{j,21} & \phi_{j,22} & \phi_{j,23} \\ \phi_{j,31} & \phi_{j,32} & \phi_{j,33} \end{bmatrix} \begin{bmatrix} y_{t-j} \\ \pi_{t-j} \\ r_{t-j} \end{bmatrix} + \begin{bmatrix} u_t^y \\ u_t^\pi \\ u_t^r \end{bmatrix}$$

This is called a **vector autoregression (VAR)**. Ever since being proposed by Sims (1980), VARs have been the Swiss knife for empirical macroeconomists. This chapter offers a thorough introduction of this Nobel prize winning technique. We start by introduce the general general vector processes and the estimation methods. We then explain how VARs map to the structural framework (SVAR). We finish the chapter by a discussion on dimension reduction techniques and the cases when a VAR system is not stationary.

26 Vector Processes

26.1 Definitions

Let \mathbf{y}_t be an $n \times 1$ vector. An vector autoregressive process is defined as

$$\mathbf{y}_t = \alpha + \Phi_1 \mathbf{y}_{t-1} + \Phi_2 \mathbf{y}_{t-2} + \cdots + \Phi_p \mathbf{y}_{t-p} + \epsilon_t$$

where ϵ_t is the vector white noise with $\mathbb{E}(\epsilon_t) = \mathbf{0}$ and $\mathbb{E}(\epsilon_t \epsilon_t') = \Omega$. In a vector form,

$$\begin{bmatrix} y_{1,t} \\ y_{2,t} \\ \vdots \\ y_{n,t} \end{bmatrix} = \begin{bmatrix} \alpha_1 \\ \alpha_2 \\ \vdots \\ \alpha_n \end{bmatrix} + \sum_{j=1}^p \begin{bmatrix} \phi_{j,11} & \phi_{j,12} & \cdots & \phi_{j,1n} \\ \phi_{j,21} & \phi_{j,22} & \cdots & \phi_{j,2n} \\ \vdots & \vdots & \ddots & \vdots \\ \phi_{j,n1} & \phi_{j,n2} & \cdots & \phi_{j,nn} \end{bmatrix} \begin{bmatrix} y_{1,t-j} \\ y_{2,t-j} \\ \vdots \\ y_{n,t-j} \end{bmatrix} + \begin{bmatrix} \epsilon_{1,t} \\ \epsilon_{2,t} \\ \vdots \\ \epsilon_{n,t} \end{bmatrix}$$

Each component $y_{j,t}$ corresponds to T observations in the data. So from the perspective of data, each VAR is represented by a dataset with T rows and n columns. To unpack the matrix notation, the first row of the vector system is

$$\begin{aligned} y_{1,t} = & \alpha_1 + \phi_{1,11}y_{1,t-1} + \phi_{1,12}y_{2,t-1} + \cdots + \phi_{1,1n}y_{n,t-1} \\ & + \phi_{2,11}y_{1,t-2} + \phi_{2,12}y_{2,t-2} + \cdots + \phi_{2,1n}y_{n,t-2} \\ & \vdots \\ & + \phi_{p,11}y_{1,t-p} + \phi_{p,12}y_{2,t-p} + \cdots + \phi_{p,1n}y_{n,t-p} \\ & + \epsilon_{1,t} \end{aligned}$$

So each variable in a VAR system is a function of the lags of itself and all other variables. In the spirit of Sims, the VAR system is intended to impose minimal restrictions. All variables are treated as endogenous and influencing each other (though we can also include exogenous variables).

26.2 VAR and VMA

We can rewrite it more compactly with the lag operator:

$$(I_n - \Phi_1 L - \Phi_2 L^2 - \cdots - \Phi_p L^p)y_t = \Phi(L)y_t = \alpha + \epsilon_t.$$

Similarly, we can generalize an MA process to the vector form:

$$y_t = \mu + \epsilon_t + \Theta_1 \epsilon_{t-1} + \Theta_2 \epsilon_{t-2} + \cdots = \mu + \Theta(L)\epsilon_t.$$

Similar to scalar processes, with stationary y_t , VAR and VMA processes can be converted to each other by inverting the lag polynomial

$$\Psi(L) = \Phi^{-1}(L)$$

where the inverse is defined as

$$[I_n - \Phi_1 L - \Phi_2 L^2 - \cdots][I_n + \Psi_1 L + \Psi_2 L^2 + \cdots] = I_n.$$

Computationally, we can expand the product of the lag polynomials

$$I_n + (\Psi_1 - \Phi_1)L + (\Psi_2 - \Phi_2 - \Phi_1\Psi_1)L^2 + \cdots = I_n$$

The coefficients of the inverse lag polynomial can be computed recursively

$$\begin{aligned}\Psi_1 &= \Phi_1 \\ \Psi_2 &= \Phi_2 + \Phi_1\Psi_1 \\ &\vdots \\ \Psi_s &= \Phi_1\Psi_{s-1} + \Phi_2\Psi_{s-2} + \cdots + \Phi_p\Psi_{s-p}\end{aligned}$$

with $\Psi_0 = I_n$, $\Psi_s = 0$ for $s < 0$.

26.3 Stationary Conditions

Proposition 26.1. *A VAR(p) process is covariance-stationary if all roots of*

$$\det|I_n - \Phi_1 z - \Phi_2 z^2 - \dots - \Phi_p z^p| = 0$$

lie outside the unit circle (z is a complex scalar).

The VAR is said to contain at least one unit root if

$$\det|I_n - \Phi_1 - \Phi_2 - \dots - \Phi_p| = 0.$$

Any VMA(q) process is covariance-stationary.

Example

Consider a two-variable VAR(1) process

$$\begin{bmatrix} y_{1t} \\ y_{2t} \end{bmatrix} = \begin{bmatrix} 0.7 & 0.1 \\ 0.3 & 0.9 \end{bmatrix} \begin{bmatrix} y_{1,t-1} \\ y_{2,t-1} \end{bmatrix} + \begin{bmatrix} u_{1t} \\ u_{2t} \end{bmatrix}$$

We compute the determinant of the matrix

$$\det \left| \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} - \begin{bmatrix} 0.7z & 0.1z \\ 0.3z & 0.9z \end{bmatrix} \right| = 0$$

$$(1 - 0.7z)(1 - 0.9z) - 0.1z \cdot 0.3z = 0$$

which solves to $z_1 = 1$, $z_2 = \frac{5}{3}$. So the VAR process is not stationary.

If the whole VAR system is stationary it follows that every single component is stationary, but not vice versa (requires proof). Similar to the univariate case, with stationarity, standard OLS and statistical inference applies. In all the sections in the chapter, we assume stationary VARs; we address the unit roots in a VAR at the end of the chapter.

26.4 Autocovariance Matrix

The autocovariance matrix for a vector process is defined as

$$\Gamma_j = \mathbb{E}[(y_t - \mu)(y_{t-j} - \mu)'].$$

For demeaned y_t , we have

$$\begin{aligned}
\Gamma_j &= \mathbb{E}(y_t y'_{t-j}) = \mathbb{E} \begin{bmatrix} y_{1,t} \\ y_{2,t} \\ \vdots \\ y_{n,t} \end{bmatrix} \begin{bmatrix} y_{1,t-j} & y_{2,t-j} & \cdots & y_{n,t-j} \end{bmatrix} \\
&= \begin{bmatrix} \mathbb{E}(y_{1,t} y_{1,t-j}) & \mathbb{E}(y_{1,t} y_{2,t-j}) & \cdots & \mathbb{E}(y_{1,t} y_{n,t-j}) \\ \mathbb{E}(y_{2,t} y_{1,t-j}) & \mathbb{E}(y_{2,t} y_{2,t-j}) & \cdots & \mathbb{E}(y_{2,t} y_{n,t-j}) \\ \vdots & \vdots & \ddots & \vdots \\ \mathbb{E}(y_{n,t} y_{1,t-j}) & \mathbb{E}(y_{n,t} y_{2,t-j}) & \cdots & \mathbb{E}(y_{n,t} y_{n,t-j}) \end{bmatrix} \\
&= \begin{bmatrix} \gamma_1(j) & \gamma_{12}(j) & \cdots & \gamma_{1n}(j) \\ \gamma_{21}(j) & \gamma_2(j) & \cdots & \gamma_{2n}(j) \\ \vdots & \vdots & \ddots & \vdots \\ \gamma_{n1}(j) & \gamma_{n2}(j) & \cdots & \gamma_n(j) \end{bmatrix}
\end{aligned}$$

Note that $\Gamma_j \neq \Gamma_{-j}$, but $\Gamma'_j = \Gamma_{-j}$.

27 Estimating VAR

Consider a more general VAR specification with n endogenous variables and m exogenous variables:

$$y_t = A_1 y_{t-1} + A_2 y_{t-2} + \cdots + A_p y_{t-p} + C x_t + u_t \quad (27.1)$$

where $t = 1, 2, \dots, T$. y_t is a $n \times 1$ vector of endogenous data, A_1, A_2, \dots, A_p are p matrices of dimension $n \times n$; C is an $n \times m$ matrix, and x_t is an $m \times 1$ vector of exogenous regressors which can be e.g. constant terms, time trends, or exogenous data series. u_t is the vector of residuals.

We assume: (1) all variables are stationary; (2) the p lags are sufficient to summarize all the dynamic correlations among elements of y_t ; and (3) u_t is uncorrelated with y_{t-1}, \dots, y_{t-p} and is free of serial correlation. Thus, $\mathbb{E}(u_t u_t') = \Omega$, while $\mathbb{E}(u_t u_s') = 0$ for $t \neq s$. Ω is an $n \times n$ symmetric positive definite variance-covariance matrix, with variance terms on the diagonal and covariance terms off diagonal.

T is the size of the sample. There are $k = np + m$ coefficients to estimate for each equation, and a total of $q = nk = n(np + m)$ coefficients to estimate for the full model.

27.1 Stacked Form

For computation, it is more convenient to rewrite the VAR system in transpose:

$$y'_t = y'_{t-1} A'_1 + y'_{t-2} A'_2 + \cdots + y'_{t-p} A'_p + x'_t C' + u'_t$$

We can stack observations to represent the whole data set:

$$\begin{bmatrix} y'_1 \\ y'_2 \\ \vdots \\ y'_T \end{bmatrix}_{T \times n} = \sum_{j=1}^p \begin{bmatrix} y'_{1-j} \\ y'_{2-j} \\ \vdots \\ y'_{T-j} \end{bmatrix}_{T \times n} A'_j + \begin{bmatrix} x'_1 \\ x'_2 \\ \vdots \\ x'_T \end{bmatrix}_{T \times m} C' + \begin{bmatrix} u'_1 \\ u'_2 \\ \vdots \\ u'_T \end{bmatrix}_{T \times n}$$

Gathering the regressors into a single matrix, one obtains:

$$\begin{bmatrix} y'_1 \\ y'_2 \\ \vdots \\ y'_T \end{bmatrix}_{T \times n} = \underbrace{\begin{bmatrix} y'_0 & y'_{-1} & \cdots & y'_{1-p} & x'_1 \\ y'_1 & y'_0 & \cdots & y'_{2-p} & x'_2 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ y'_{T-1} & y'_{T-2} & \cdots & y'_{T-p} & x'_T \end{bmatrix}}_{T \times (np+m)} + \underbrace{\begin{bmatrix} A'_1 \\ A'_2 \\ \vdots \\ A'_p \\ C' \end{bmatrix}}_{(np+m) \times n} + \begin{bmatrix} u'_1 \\ u'_2 \\ \vdots \\ u'_T \end{bmatrix}_{T \times n}$$

Or, more compactly:

$$Y_{T \times n} = X_{T \times k} B_{k \times n} + U_{T \times n}. \quad (27.2)$$

Once the model is stacked this way, obtaining OLS estimates of the VAR is straightforward. An estimate \hat{B} is obtained from:

$$\hat{B} = (X'X)^{-1}X'Y$$

This is equivalent to applying OLS on each column variable:

$$\begin{bmatrix} \hat{B}^{(1)} & \hat{B}^{(2)} & \cdots & \hat{B}^{(n)} \end{bmatrix} = (X'X)^{-1}X' \begin{bmatrix} Y^{(1)} & Y^{(2)} & \cdots & Y^{(n)} \end{bmatrix}$$

where $Y^{(i)}$ denotes i -th column of matrix Y . An estimate of the covariance matrix Ω can be obtained from:

$$\hat{\Omega} = \frac{1}{T - k - 1} \hat{U}'\hat{U}$$

Under the assumptions (1)-(3), the parameters are consistently estimated by OLS regressions. Standard asymptotic results apply.

27.2 Vectorized Form

Alternatively, one can vectorize the VAR system as:

$$\begin{bmatrix} Y^{(1)} \\ Y^{(2)} \\ \vdots \\ Y^{(n)} \end{bmatrix}_{nT \times 1} = \begin{bmatrix} X & & & \\ & X & & \\ & & \ddots & \\ & & & X \end{bmatrix}_{nT \times nk} \begin{bmatrix} B^{(1)} \\ B^{(2)} \\ \vdots \\ B^{(n)} \end{bmatrix}_{nk \times 1} + \begin{bmatrix} U^{(1)} \\ U^{(2)} \\ \vdots \\ U^{(n)} \end{bmatrix}_{nT \times 1}$$

where $Y^{(i)}$ denotes i -th column of matrix Y . The vectorized system can be compactly written as

$$y = \bar{X}\beta + u \quad (27.3)$$

where $y = \text{vec}(Y)$, $\bar{X} = I_n \otimes X$, $\beta = \text{vec}(B)$, $u = \text{vec}(U)$. Also, one has $\mathbb{E}(uu') = \bar{\Omega}$, where $\bar{\Omega} = \Omega \otimes I_T$. An OLS estimate of the vectorised β can be obtained as:

$$\hat{\beta} = (\bar{X}'\bar{X})^{-1}\bar{X}'y$$

It should be noted that Equation 27.2 and Equation 27.3 are just alternative but equivalent representations of the same VAR model Equation 27.1. We opt to use one representation or the other according to which one is most convenient for our purposes. Equation 27.2 is typically faster to compute (because smaller matrices produce more accurate estimates), while statistical inference works more naturally with the vectorized form.

28 Granger Causality

28.1 Definition

Granger causality conceptualizes the usefulness of some variables in forecasting other variables.

Definition 28.1 (Granger Causality). x fails to Granger-cause y if for all $s > 0$, the MSE of a forecast \hat{y}_{t+s} based on $(y_t, y_{t-1}, \dots, x_t, x_{t-1}, \dots)$ is the same as the MSE of the forecast based on (y_t, y_{t-1}, \dots) :

$$\text{MSE}[\hat{y}_{t+s}|y_t, y_{t-1}, \dots] = \text{MSE}[\hat{y}_{t+s}|y_t, y_{t-1}, \dots, x_t, x_{t-1}, \dots].$$

Equivalently, we say x is exogenous to y , or x is not linearly informative about y .

It must be noted that Granger causality has nothing to do with the causality defined by counterfactuals. The word “causality” is unfortunately misleading. It would be better named as “Granger predictability”.

28.2 Granger Causality Test

Consider the Granger causality test in a single equation setup:

$$y_t = \alpha + \phi_1 y_{t-1} + \dots + \phi_p y_{t-p} + \beta_1 x_{t-1} + \dots + \beta_q x_{t-q} + u_t$$

Testing whether x Granger-causes y is equivalent to test

$$H_0 : \beta_1 = \beta_2 = \dots = \beta_q = 0$$

The test can be done by comparing the residual sum of squares (RSS) with and without x as the regressors. Assuming H_0 holds, we would have the restricted model:

$$y_t = \alpha + \phi_1 y_{t-1} + \dots + \phi_p y_{t-p} + u_t^R$$

Compute RSS for the restricted model:

$$\text{RSS}_0 = \sum_{t=1}^T (\hat{u}_t^R)^2$$

Also compute the RSS for the unrestricted model, i.e. including all x as the regressors:

$$\text{RSS}_1 = \sum_{t=1}^T \hat{u}_t^2$$

The joint significance can be tested by the F ratio:

$$S = \frac{(\text{RSS}_0 - \text{RSS}_1)/q}{\text{RSS}_1/(T - 2q - 1)} \sim F(q, T - 2q - 1).$$

28.3 Granger Causality in VAR

In a VAR setting, x does not Granger-cause y if the coefficient matrix are lower triangular for all j :

$$\begin{aligned} \begin{bmatrix} y_t \\ x_t \end{bmatrix} &= \begin{bmatrix} \alpha_1 \\ \alpha_2 \end{bmatrix} + \begin{bmatrix} \phi_{1,11} & 0 \\ \phi_{1,21} & \phi_{1,22} \end{bmatrix} \begin{bmatrix} y_{t-1} \\ x_{t-1} \end{bmatrix} + \begin{bmatrix} \phi_{2,11} & 0 \\ \phi_{2,21} & \phi_{2,22} \end{bmatrix} \begin{bmatrix} y_{t-2} \\ x_{t-2} \end{bmatrix} + \dots \\ &= \begin{bmatrix} \alpha_1 \\ \alpha_2 \end{bmatrix} + \begin{bmatrix} \phi_{11}(L) & 0 \\ \phi_{21}(L) & \phi_{22}(L) \end{bmatrix} \begin{bmatrix} y_t \\ x_t \end{bmatrix} + \begin{bmatrix} u_t \\ v_t \end{bmatrix} \end{aligned}$$

It is equivalent to the MA representation:

$$\begin{bmatrix} y_t \\ x_t \end{bmatrix} = \begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix} + \begin{bmatrix} \theta_{11}(L) & 0 \\ \theta_{21}(L) & \theta_{22}(L) \end{bmatrix} \begin{bmatrix} u_t \\ v_t \end{bmatrix}$$

To test Granger causality in a VAR setting, we need to introduce the likelihood ratio (LR) test.

28.4 Likelihood Ratio Test

Let's quickly derive the maximum likelihood estimator (MLE) for the VAR. The joint probability density function is

$$\begin{aligned} f(y_T, y_{T-1}, \dots, y_1 | y_0, \dots, y_{1-p}; \Theta) &= f(y_T | y_{T-1} \dots) f(y_{T-1} | y_{T-2} \dots) \cdots f(y_1 | y_0 \dots) \\ &= \prod_{t=1}^T f(y_t | y_{t-1}, y_{t-2}, \dots, y_{1-p}; \Theta) \end{aligned}$$

Define $x'_t = \begin{bmatrix} 1 & y_{t-1} & \dots & y_{t-p} \end{bmatrix}$ the collection of all regressors, and $B' = \begin{bmatrix} \alpha & \Phi_1 & \dots & \Phi_p \end{bmatrix}$ the collection of all parameters. The log-likelihood function can be written as

$$\begin{aligned} \ell(\theta) &= \sum_{t=1}^T \log f(y_t | y_{t-1}, y_{t-2}, \dots, y_{1-p}; \theta) \\ &= -\frac{T}{2} \log |2\pi\Omega^{-1}| - \frac{1}{2} \sum_{t=1}^T [(y_t - B'x_t)' \Omega^{-1} (y_t - B'x_t)] \end{aligned}$$

where Ω is the variance-covariance matrix of the residuals. Maximizing the log-likelihood function gives the ML estimator:

$$\begin{aligned} \hat{B}'_{n \times (np+1)} &= \left[\sum_{t=1}^T y_t x'_t \right] \left[\sum_{t=1}^T x_t x'_t \right]^{-1} \\ \hat{\Omega}_{n \times n} &= \frac{1}{T} \sum_{t=1}^T \hat{e}_t \hat{e}'_t \end{aligned}$$

where $e_t = [u_t \ v_t]'$.

The likelihood ratio (LR) test is motivated by the fact that different specifications give different likelihood evaluates. By comparing the likelihood difference, we can test the significance of one specification versus the alternative.

The null hypothesis (H_0) could be a specification with a particular lag length, or a particular-ization with certain exogeneity restrictions. We compute the covariance matrix under the null hypothesis (H_0) and the alternative (H_1) respectively:

$$\begin{aligned} \hat{\Omega}_0 &= \frac{1}{T} \sum_t \hat{e}_t(H_0) \hat{e}_t(H_0)' \\ \hat{\Omega}_1 &= \frac{1}{T} \sum_t \hat{e}_t(H_1) \hat{e}_t(H_1)' \end{aligned}$$

The corresponding log-likelihoods are

$$\begin{aligned}\ell_0^* &= -\frac{T}{2} \log |2\pi\hat{\Omega}_0^{-1}| - \frac{Tn}{2} \\ \ell_1^* &= -\frac{T}{2} \log |2\pi\hat{\Omega}_1^{-1}| - \frac{Tn}{2}\end{aligned}$$

The difference between the log-likelihoods

$$2(\ell_1^* - \ell_0^*) = T(\log |\hat{\Omega}_0| - \log |\hat{\Omega}_1|)$$

has a χ^2 distribution with degree of freedom $n^2(p_1 - p_0)$.

Now consider Granger causality test in a VAR setting

$$\begin{bmatrix} y_t \\ x_t \end{bmatrix} = \begin{bmatrix} \alpha_1 \\ \alpha_2 \end{bmatrix} + \begin{bmatrix} \phi_{11}(L) & \phi_{12}(L) \\ \phi_{21}(L) & \phi_{22}(L) \end{bmatrix} \begin{bmatrix} y_t \\ x_t \end{bmatrix} + \begin{bmatrix} u_t \\ v_t \end{bmatrix}$$

Testing x fails to Granger-cause y is equivalent to test $\phi_{12} = 0$. Therefore, the restricted regression under H_0 is

$$y_t = \alpha_1 + \phi_{11}(L)y_t + u_t^R$$

The unrestricted regression is

$$y_t = \alpha_1 + \phi_{11}(L)y_t + \phi_{12}(L)x_t + u_t^U$$

Estimate the variance-covariance matrices

$$\begin{aligned}\hat{\Omega}^U &= \frac{1}{T} \sum_t u_t^U u_t^{U'} \\ \hat{\Omega}^R &= \frac{1}{T} \sum_t u_t^R u_t^{R'}\end{aligned}$$

Form the LR test statistics

$$\text{LR} = T(\log |\hat{\Omega}^U| - \log |\hat{\Omega}^R|) \sim \chi^2$$

If the LR statistics is significant, it would mean x is informative about y (x Granger-causes y). Otherwise, if it makes no difference including x as regressors, it would mean x fails to Granger-cause y .

29 Structural VAR

29.1 The Structural Framework

Now we reconsider the problem of estimating structural shocks. We have introduced the structural shock framework in Chapter 15 the underlying econometric framework of understanding our economy. To briefly recap, we envision our economy as an MA process in which multiple structural shocks are the fundamental driving forces:

$$y_t = \Theta(L)\epsilon_t \quad (29.1)$$

Structural shocks must be distinguished from residuals in a regression. Residuals are prediction errors based on past observations. Residuals can be cross-sectionally or serially correlated. Structural shocks are attached with specific economic meaning. They are also assumed to be unforeseeable and uncorrelated. We do not usually observe structural shocks but they are the conceptualized driving forces in the background.

If $\Theta(L)$ is invertible, we would have

$$\Theta^{-1}(L)y_t = \epsilon_t$$

which is an infinite order AR process. This motivates us to estimate structural shocks via vector autoregressive processes. Suppose we have a VAR process:

$$A(L)y_t = u_t \quad (29.2)$$

where $u_t = y_t - \text{Proj}(y_t|y_{t-1}, y_{t-2}, \dots)$ are the projection residuals. The question is, to what extend, or under what conditions, can we identify the structural shocks from this VAR specification?

The answer is easier than you might have thought. We only need to identify

$$u_t = \Theta_0\epsilon_t \quad (29.3)$$

where Θ_0 is the first coefficient matrix in the lag polynomial. That is, the condition for identification is that we can find a linear transformation to decompose u_t into ϵ_t .

29.2 Invertibility

The structural MA process is said to be **invertible** if ϵ_t can be linearly determined from current and lagged values of y_t :

$$\epsilon_t = \text{Proj}(\epsilon_t | y_t, y_{t-1}, \dots).$$

This means there is no “omitted variable” in the observable space, in the sense that the space spanned by $\{\epsilon_t, \epsilon_{t-1}, \dots\}$ is fully covered by $\{y_t, y_{t-1}, \dots\}$.

$$\text{span}\{\epsilon_t, \epsilon_{t-1}, \dots\} = \text{span}\{y_t, y_{t-1}, \dots\} = \text{span}\{u_t, u_{t-1}, \dots\}$$

This is a strong assumption. Under invertibility, the knowledge of the past true shocks would not even improve the the VAR forecast. But it does not require our VAR system being exhaustive, including everything observable variables in our economy. In a particular application, we would only be interested in a few structural shocks. The invertibility condition requires the observables fully cover the space spanned by the structural shocks of particular interests.

With the above condition satisfied, we can show that the identification problem is reduced to identify Θ_0 . Given Equation 29.1 and Equation 29.2, we have

$$u_t = A(L)y_t = A(L)\Theta(L)\epsilon_t \stackrel{?}{=} \Theta_0\epsilon_t$$

By definition,

$$\begin{aligned} u_t &= y_t - \text{Proj}[y_t | y_{t-1}, y_{t-2}, \dots] \\ &= \Theta(L)\epsilon_t - \text{Proj}[\Theta(L)\epsilon_t | y_{t-1}, y_{t-2}, \dots] \\ &= \Theta_0\epsilon_t + \Theta_1\epsilon_{t-1} + \dots + \text{Proj}[\Theta_0\epsilon_t + \Theta_1\epsilon_{t-1} + \dots | y_{t-1}, y_{t-2}, \dots] \\ &= \Theta_0\epsilon_t - \underbrace{\Theta_0 \text{Proj}[\epsilon_t | y_{t-1}, \dots]}_{=0 \text{ by definition}} + \sum_{j=1}^{\infty} \Theta_j \{ \epsilon_{t-j} - \underbrace{\text{Proj}[\epsilon_{t-j} | y_{t-1}, \dots]}_{=\epsilon_{t-j} \text{ by invertibility}} \} \\ &= \Theta_0\epsilon_t. \end{aligned}$$

Proposition 29.1 (Assumptions of Structural VAR).

1. All variables are stationary;
2. The space spanned by the innovations and the structural shocks coincide such that $u_t = \Theta_0\epsilon_t$;
3. The structural process $y_t = \Theta(L)\epsilon_t$ is invertible.

Under the assumptions, identifying Θ_0 is equivalent to identify the structural shocks $\epsilon_t = \Theta_0^{-1}u_t$.

In essence, structural identification is equivalent to sorting out the contemporaneously correlated residuals into uncorrelated shocks that can be attached to certain economic meanings. As we will see, the decomposition is largely subjective, according to researchers' understanding of how structural shocks are correlated contemporaneously.

If the invertibility assumption fails, that means there exists no mapping from VAR residuals to the structural shocks. Non-invertibility arises when the observed variables fail to span the space of the state variables (structural shocks). If this is the case, we can include more variables to expand the space; or we may choose to simply ignore it if we believe the wedge between the spaces spanned by VAR residuals and structural shocks are small.

29.3 Identification

With invertibility, the essential task of SVAR is to decompose Equation 29.3 to recover the structural shocks. The key is to estimate Θ_0 . Consider the second-order moments of Equation 29.3:

$$\Theta_0 \mathbb{E}(\epsilon_t \epsilon_t') \Theta_0' = \mathbb{E}(u_t u_t')$$

Estimating the VAR system Equation 29.2 by OLS gives $\hat{\Omega} = \hat{\mathbb{E}}(u_t u_t')$. By definition, elements of ϵ_t are orthogonal to each other, so $D = \mathbb{E}(\epsilon_t \epsilon_t')$ is diagonal. Estimated $\hat{\Omega}$ gives $n(n+1)/2$ distinct values. Identification of D requires n values. So no more than $n(n-1)/2$ parameters in Θ_0 can be identified. That means, we cannot identify the full $n \times n$ matrix Θ_0 without restrictions.

29.3.1 Recursive restriction

One common way to impose restrictions on Θ_0 is to require it being lower triangular. Thus eliminating $n(n-1)/2$ entries. We also assume the structural shocks have the same magnitude as the residuals, so the diagonal entries are 1s. For example, in the three variable Keynesian system, we may assume

$$\begin{bmatrix} u_t^\pi \\ u_t^y \\ u_t^m \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ * & 1 & 0 \\ * & * & 1 \end{bmatrix} \begin{bmatrix} \epsilon_t^S \\ \epsilon_t^{IS} \\ \epsilon_t^{MP} \end{bmatrix}$$

The recursive structure is equivalent to imposing restrictions on the contemporaneous relationships between variable, or imposing different reaction speed to the structural shocks. In

the above example, we assume the observed monetary policy (interest rate) responds to IS shock, supply shock and monetary policy shock contemporaneously; but inflation and output respond to monetary policy shock with a lag (sluggish response). Output responds to IS shock and supply shock contemporaneously, but inflation responds to IS shock with a lag. Of course, one can question the validity of these assumptions, or even the validity of the conceptualization of the three structural shocks. But this is a structural question, not an econometric one. Economists have always been debating what are the proper structures to describe the economy.

In the recursive identification scheme, the ordering of the variables is the vital decision to make. Typically, the slow-moving variables are ordered first, and the fast-moving variables last, provided the Θ_0 is upper triangular. In the literature involving monetary policy, a “slow-r-fast” scheme is widely adopted. That is, low-moving variables such as real output and price levels are ordered before interest rates; and fast-moving variables such as financial market indexes are ordered after interest rates. Because financial market absorbs information in real time, even ahead of the monetary policy decision. But it take time for real variables to materialize the impact of monetary policy changes.

29.3.2 Non-recursive restriction

We may also impose non-recursive structure based on theories or intuitions. Consider a model constituted of the demand and supply of an agriculture product, and the weather condition that affects the supply of the product. We assume weather does not depend on market behaviors. In addition, the supply but not the demand is influenced by the weather. This results in an identification matrix as follows

$$\begin{bmatrix} u_t^d \\ u_t^s \\ u_t^w \end{bmatrix} = \begin{bmatrix} 1 & -\beta & 0 \\ 1 & -\gamma & -\delta \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \epsilon_t^d \\ \epsilon_t^s \\ \epsilon_t^w \end{bmatrix}$$

Note that there are only three parameters to be estimated in the matrix. So Θ_0 in this case can also be identified.

Structural VAR literature has invented lots of identification schemes, such as long-run restrictions, sign restrictions, and so on. These are left for the readers to explore themselves.

Takeaways

1. Reduced-form VARs only require errors be free of serial correlation, but allow cross-sectional correlations. The errors do not have structural interpretation.
2. Structural identification means to decompose reduced-form residuals into uncorrelated structural shocks, so that we can attach structural meaning to the identified shocks.

3. Under the assumption of invertibility, structural identification boils down to restricting the contemporaneous correlations between the endogenous variables. There are various identification schemes including Cholesky decomposition, sign restrictions, long-run restrictions, and so on.

30 IRF and FEVD

30.1 Estimating SVAR

Given the Structural VAR,

$$\Theta_0 y_t = \Phi_1 y_{t-1} + \Phi_2 y_{t-2} + \cdots + \Phi_p y_{t-p} + \epsilon_t$$

We estimate the reduced-form VAR using the methods in Chapter 27:

$$y_t = A_1 y_{t-1} + A_2 y_{t-2} + \cdots + A_p y_{t-p} + u_t$$

where $A_j = \Theta_0^{-1} \Phi_j$ and $u_t = \Theta_0^{-1} \epsilon_t$. Θ_0 can be estimated using the identification scheme discussed in Chapter 29, so that the structural parameters can be recovered accordingly.

However, the coefficients of a VAR are difficult to interpret. For empirical analysis, we are particularly interested in estimating the impact of a structural shock on other economic variables (e.g. the impact of monetary policy shock on inflation). We cannot read off this information from the estimated VAR coefficients.

In this section, we introduce two reporting techniques: **impulse-response functions (IRF)** and **forecast error variance decomposition (FEVD)**.

30.2 Impulse-Response Functions

With VAR parameters estimated, we can convert it to the MA form:

$$y_t = A^{-1}(L)u_t = u_t + \Psi_1 u_{t-1} + \Psi_2 u_{t-2} + \cdots$$

or written in structural shocks:

$$y_t = \Theta_0^{-1} \epsilon_t + \Psi_1 \Theta_0^{-1} \epsilon_{t-1} + \Psi_2 \Theta_0^{-1} \epsilon_{t-2} + \cdots$$

With this structural MA form, we can directly read off the impact of structural shock j on observable variable k :

$$\frac{\partial y_{t+h}^k}{\partial \epsilon_t^j} = (\Psi_h \Theta_0^{-1})_{kj}$$

A sequence of the marginal impacts over time $\left\{ \frac{\partial y_t^k}{\partial \epsilon_t^j}, \frac{\partial y_{t+1}^k}{\partial \epsilon_t^j}, \frac{\partial y_{t+2}^k}{\partial \epsilon_t^j}, \dots \right\}$ constitute the dynamic response of variable k in response to structural shock j , also known as the impulse-response function.

Note that $\frac{\partial y_{t+h}^k}{\partial u_t^j}$ cannot be interpreted as the response of y_{t+h}^k after a shock on u_t^j . Because u_t^j could be a linear combination of multiple structural shocks, e.g.

$$u_t^1 = a_1 \epsilon_t^1 + a_2 \epsilon_t^2 + a_3 \epsilon_t^3$$

Therefore, no structural meaning can be attached to $\frac{\partial y_{t+h}^k}{\partial u_t^j}$. But $\frac{\partial y_{t+h}^k}{\partial \epsilon_t^j}$ is interpretable, as structural shocks ϵ_t are uncorrelated.

If Θ_0 has a recursive structure, Θ_0 can be found simply by applying decomposition to the OLS-estimated $\hat{\Omega}$:

$$\hat{\Omega} = LDL'$$

If $\hat{\Omega}$ is positive definite, such decomposition always exist, with D diagonal and L lower triangular with 1s on the diagonal.

If we restrict D to be an identity matrix, the decomposition becomes

$$\hat{\Omega} = PP'$$

where $P = LD^{1/2}$. This is known as the *Cholesky decomposition*. By applying the Cholesky decomposition, we have

$$\zeta_t = P^{-1}u_t = D^{-1/2}L^{-1}u_t = D^{-1/2}\epsilon_t$$

in which 1 unit shock to ζ_t is equivalent to 1 standard deviation shock to ϵ_t . The IRFs estimated with Cholesky decomposition thus have the interpretation of response to standard-deviation shocks.

30.3 IRF Standard Errors

The IRF confidence intervals are usually constructed by bootstrapping.

1. Estimate the VAR by OLS. Save residuals $\hat{u}_1, \hat{u}_2, \dots, \hat{u}_T$.
2. Randomly pick u_1 from $\{\hat{u}_1, \hat{u}_2, \dots, \hat{u}_T\}$, which will be used to construct an artificial sample. Generate

$$y_1^{(1)} = \hat{c} + \hat{\Phi}_1 y_0 + \hat{\Phi}_2 y_{-1} + \dots + \hat{\Phi}_p y_{1-p} + u_1^{(1)}$$

where $y_0, y_{-1}, \dots, y_{1-p}$ are pre-sampled values that were actually observed. Take a second draw u_2 . Generate

$$y_2^{(1)} = \hat{c} + \hat{\Phi}_1 y_1 + \hat{\Phi}_2 y_0 + \dots + \hat{\Phi}_p y_{2-p} + u_2^{(1)}$$

Proceed until $\{y_1^{(1)}, y_2^{(1)}, \dots, y_T^{(1)}\}$ are generated. Run OLS on the simulated data, calculate the impulse-response function $\{\Psi_h^{(1)}\}$.

3. Repeat the above step, generate $\{y_1^{(2)}, y_2^{(2)}, \dots, y_T^{(2)}\}$. Estimate the IRF again $\{\Psi_h^{(2)}\}$. Repeat the process N times and get N IRFs. Sort them by $\Psi^{(1)} \leq \Psi^{(2)} \dots \leq \Psi^{(N)}$.
4. The α confidence interval is constructed as $[\Psi_{[N\alpha/2]}, \Psi_{[N(1-\alpha/2)]}]$.

30.4 FEVD

We would also like to know the relative importance of each structural force, which can be gauged by computing the **forecast error variance decomposition (FEVD)**. Consider the forecast error

$$y_{t+h} - \hat{y}_{t+h|t} = \Theta_0^{-1} \epsilon_{t+h} + \Psi_1 \Theta_0^{-1} \epsilon_{t+h-1} + \dots + \Psi_{h-1} \Theta_0^{-1} \epsilon_{t+1}$$

The mean squared error (MSE) is

$$\begin{aligned} \text{MSE}(\hat{y}_{t+h|t}) &= \mathbb{E}[(y_{t+h} - \hat{y}_{t+h|t})(y_{t+h} - \hat{y}_{t+h|t})'] \\ &= \Omega + \Psi_1 \Omega \Psi_1' + \Psi_2 \Omega \Psi_2' + \dots + \Psi_{h-1} \Omega \Psi_{h-1}' \end{aligned}$$

where

$$\begin{aligned}
\Omega &= \mathbb{E}(\Theta_0^{-1} \epsilon_t \epsilon_t' \Theta_0^{-1'}) \\
&= \begin{bmatrix} a_1 & a_2 & \dots & a_n \end{bmatrix} \begin{bmatrix} \text{var}(\epsilon_{1,t}) & & & \\ & \text{var}(\epsilon_{2,t}) & & \\ & & \ddots & \\ & & & \text{var}(\epsilon_{n,t}) \end{bmatrix} \begin{bmatrix} a_1' \\ a_2' \\ \vdots \\ a_n' \end{bmatrix} \\
&= a_1 a_1' \text{var}(\epsilon_{1,t}) + a_2 a_2' \text{var}(\epsilon_{2,t}) + \dots + a_n a_n' \text{var}(\epsilon_{n,t})
\end{aligned}$$

Therefore,

$$\text{MSE}(\hat{y}_{t+h|t}) = \sum_{j=1}^n \{ \text{var}(\epsilon_{j,t}) [a_j a_j' + \Psi_1 a_j a_j' \Psi_1' + \dots + \Psi_{h-1} a_j a_j' \Psi_{h-1}'] \}$$


The contribution of the j -th structural force to the MSE of the h -period-ahead forecast is given by

$$\text{var}(\epsilon_{j,t}) [a_j a_j' + \Psi_1 a_j a_j' \Psi_1' + \dots + \Psi_{h-1} a_j a_j' \Psi_{h-1}'].$$

30.5 Example




```
library(vars)
data = readRDS("data/md.Rds")
y = cbind(p = data$CPI, y = data$NGDP, r = data$Repo7D)
var = VAR(as.ts(y), p = 4)
irf(var, impulse = "r", response = "p", ortho = TRUE) |> plot()
```

Orthogonal Impulse Response from r



95 % Bootstrap CI, 100 runs

```
fevd(var) |> plot()
```



31 Factor Models

VARs are not parsimonious models. A VAR with n endogenous variables and p lags would have n^2p parameters to be estimated. With more variables included in the model, the degree of freedom is quickly exhausted. On the other hand, in many applications, we would want to include more variables into the model. The economy is complex and the endogenous variables are huge. Also, including more variables would mitigate the omitted variable problem that would lead to more credible identification. To tackle the “big data” challenge, we would need some dimension reduction technology.

31.1 Principle Component Analysis

Question: how to summarize the movements of a large number of time series with fewer time series?




Figure 31.1: 50 economic time series in monthly percentage changes

The native approach is to take the average after some standardization. A better approach is to find a linear combination of them such that:

- (1) The 1st linear combination captures most of the variances among the series;
- (2) The 2nd linear combination, while being orthogonal to the 1st one, captures most of the remaining variations;
- (3) ...

To formulate the question mathematically, let $X = [x_1 \ x_2 \ \dots \ x_p]'$ represent p time series. In the form of data matrix

$$X = \begin{bmatrix} x_{11} & x_{21} & \dots & x_{p1} \\ x_{12} & x_{22} & \dots & x_{p2} \\ \vdots & \vdots & \ddots & \vdots \\ x_{1T} & x_{2T} & \dots & x_{pT} \end{bmatrix}_{T \times p}$$

We want to find a linear combination of X :

$$a'X = a_1x_1 + a_2x_2 + \dots + a_px_p$$

such that $a'X$ captures the greatest variance. If X is a 2-dimensional vector $X = [x_1 \ x_2]'$. We want to find a vector $[a_1 \ a_2]$ such that projecting the data onto this direction gives the largest variance.




Figure 31.2: PCA of a multivariate Gaussian distribution

Expressed as an optimization problem, we want to solve

$$\max_a \mathbb{E}(a' X X' a)$$

where we normalize a to a unit vector $a'a = 1$. Define the Lagrangian

$$\mathcal{L} = a' \Sigma a - \lambda_1 (a'a - 1)$$

where $\Sigma = \mathbb{E}(X X')$. The first-order condition gives

$$\frac{\partial \mathcal{L}}{\partial a} = 2\Sigma a - 2\lambda_1 a = 0 \implies \Sigma a = \lambda_1 a$$

This means, a is an eigenvector of Σ . The linear combination has variance

$$a' \Sigma a = a' \lambda_1 a = \lambda_1 a' a = \lambda_1$$

which is the eigenvalue. So, if we choose a to be the eigenvector associated with the largest eigenvalue of the covariance matrix Σ and project the data onto it, the projected variance will be maximized. The eigenvector points to the direction that the components of X co-vary the most. It is called the **first principle component**.

If we want a second principle component that points to the direction that the data co-vary second to the most, by similar reasoning, we pick the second largest eigenvalue λ_2 :

$$\Sigma b = \lambda_2 b$$

Then b is the **second principle component**. Since Σ is symmetric, we know that the eigenvectors are orthogonal, $b \perp a$.

$$\lambda_1 a' b = (\lambda_1 a)' b = (\Sigma a)' b = a' \Sigma' b = a' \Sigma b = a' \lambda_2 b = \lambda_2 a' b$$

$$\implies (\lambda_1 - \lambda_2) a' b = 0$$

Since $\lambda_1 \neq \lambda_2$, we have $a' b = 0$.

PCA Procedure

1. Standardize the dataset $X_{T \times n}$ by mean and variance;
2. Calculate the covariance matrix $\Sigma_{n \times n}$ for the dataset;
3. Calculate the eigenvalues and eigenvectors for the covariance matrix;

4. Sort eigenvalues and their corresponding eigenvectors;
5. Pick the first k eigenvalues and form a matrix of eigenvectors $A_{n \times k} = [a_1 \ a_2 \ \dots \ a_k]$;
6. Transform the original matrix $Y_{T \times k} = X_{T \times n} \times A_{n \times k}$.

The transformed matrix Y is the reduced-dimension dataset $k \ll n$ that captures as much variance as possible of the original dataset. One of the drawbacks of PCA is that the principle components do not have an economic meaning. We do not know how to interpret the principle components except that they represents the co-movements among the original variables. If we want a clear meaning of the principle components, it is suggested we group the original dataset by categories and extract principle components for each category. For example, if we put all price indexes in a group and the principle components of that group would likely be interpreted as the most representative price movement.

31.2 Factor-augmented VAR

Standard VAR assumes the “shocks” are identified by the VAR residuals after imposing some restriction. However, due to the “curse of dimensionality”, standard VAR can only include a limited number of variables. Three problems would arise because of that:

- (1) Small number of variables is unlikely to span the space of structural shocks. For example, in the application of identifying monetary policy shocks, if the information set used by the central bank is not fully captured by the VAR, the residuals would not span the space of structural shocks.
- (2) It is questionable that specific observable measures correspond precisely to the theoretical constructs. For example, the concept of “economic activity” may not be precisely represented by GDP.
- (3) Standard VAR can only generate a limited number of impulse responses that we care about. In many applications, we would care about a wide range of impulse responses from various aspects of the economy.

One solution is to augment a standard VAR with a few principle components (factors) estimated from big dataset. A factor-augmented VAR (FAVAR) is specified as follows:

$$X_t = \Lambda^f f_t + \Lambda^y y_t + u_t$$

$$\begin{bmatrix} f_t \\ y_t \end{bmatrix} = \Phi(L) \begin{bmatrix} f_{t-1} \\ y_{t-1} \end{bmatrix} + v_t$$

where X_t is an $n \times 1$ vector representing the information set, which is assumed to be spanned by $k \times 1$ factors and $m \times 1$ observable measures. Λ^f is called the factor loading matrix. We assume the factors and observables follow a vector autoregressive process.

We follow a two-step procedure to estimate the FAVAR:

1. Estimate the factors by principle components of X_t , denoted by \hat{C}_t . \hat{F}_t is the part of \hat{C}_t not spanned by y_t .
2. Estimate the FAVAR with the factors \hat{F}_t . Apply identification similar to standard VARs.

We use the example of Bernanke et al. (2005) as an illustration, in which the authors use an FAVAR to identify the impact of US monetary policy shocks. They treat only the Fed's policy instrument r_t as observed, all other variables including output and inflation, as unobserved (captured by f_t). They adopt a “slow-r-fast” identification scheme:

Slow-moving variables x_t^S	Policy instrument r_t	Fast-moving variables x_t^F
Output	Fed fund rate	Asset price
Employment		Financial shocks
Inflation		News shocks
.....	

The FAVAR is specified as

$$\begin{bmatrix} x_t^S \\ x_t^F \end{bmatrix} = \begin{bmatrix} \Lambda_{SS} & 0 & 0 \\ \Lambda_{FS} & \Lambda_{FR} & \Lambda_{FF} \end{bmatrix} \begin{bmatrix} f_t^S \\ r_t \\ f_t^F \end{bmatrix} + u_t$$

$$\Phi(L) \begin{bmatrix} f_t^S \\ r_t \\ f_t^F \end{bmatrix} = \begin{bmatrix} \eta_t^S \\ \eta_t^R \\ \eta_t^F \end{bmatrix}$$

with recursive identification

$$\begin{bmatrix} \eta_t^S \\ \eta_t^R \\ \eta_t^F \end{bmatrix} = \begin{bmatrix} H_{SS} & 0 & 0 \\ H_{RS} & 1 & 0 \\ H_{FS} & H_{FR} & H_{FF} \end{bmatrix} \begin{bmatrix} \epsilon_t^S \\ \epsilon_t^R \\ \epsilon_t^F \end{bmatrix}.$$

The IRFs for every variable in X_t can be constructed with the factor loading matrix. FEVDs follow immediately from the coefficients of the MA representation of the VAR system and the variance of the structural shocks.

Bernanke et al. (2005) finds adding the unobserved factors change the result dramatically. The IRF from a standard VAR exhibits the so-called “price puzzle”, that is inflation does not drop

immediately after a tightening monetary policy shock. The FAVAR reduces the price puzzle significantly, which indicates the FAVAR indeed incorporates useful information that is missing in standard VARs.

32 Unit Roots in VAR

So far we have only considered stationary VARs. Now we discuss what happens if a VAR contains unit roots. Recall that a VAR(p) process

$$y_t = \Phi_1 y_{t-1} + \Phi_2 y_{t-2} + \cdots + \Phi_p y_{t-p} + \epsilon_t$$

contains at least one unit root iff

$$|I_n - \Phi_1 - \Phi_2 - \cdots - \Phi_p| = 0.$$

Suppose some or all components of y_t are $I(1)$, others are $I(0)$. Rewrite the equation as follows:

$$\begin{aligned} y_t &= \Phi_1 y_{t-1} + \cdots + \Phi_{p-2} y_{t-p+2} + \Phi_{p-1} y_{t-p+1} + \Phi_p y_{t-p} \\ &= \Phi_1 y_{t-1} + \cdots + \Phi_{p-2} y_{t-p+2} - (\Phi_{p-1} + \Phi_p) y_{t-p+1} + \Phi_p \Delta y_{t-p+1} \\ &= \Phi_1 y_{t-1} + \cdots - (\Phi_{p-2} + \Phi_{p-1} + \Phi_p) y_{t-p+2} - (\Phi_{p-1} + \Phi_p) \Delta y_{t-p+2} + \Phi_p \Delta y_{t-p+1} \\ &\vdots \end{aligned}$$

Therefore, the VAR process can be rewritten as

$$y_t = \rho y_{t-1} + \zeta_1 \Delta y_{t-1} + \cdots + \zeta_{p-2} \Delta y_{t-p+2} + \zeta_{p-1} \Delta y_{t-p+1} + \epsilon_t$$

where

$$\begin{aligned} \rho &= \Phi_1 + \Phi_2 + \cdots + \Phi_{p-1} + \Phi_p \\ -\zeta_1 &= \Phi_2 + \cdots + \Phi_{p-1} + \Phi_p \\ &\vdots \\ -\zeta_{p-2} &= \Phi_{p-1} + \Phi_p \\ -\zeta_{p-1} &= \Phi_p \end{aligned}$$

Thus, $\Phi_1 = \rho + \zeta_1$, $\Phi_s = \zeta_s - \zeta_{s-1}$, $\Phi_p = -\zeta_{p-1}$. So the coefficients of the original VAR $\{\Phi_s\}$ can be written as linearly combinations of coefficients on stationary regressors $\{\zeta_s\}$. According

to the theorem in Chapter 24, the asymptotic distribution of Φ_s would be dominated by slower converging ζ_s . It follows that $\sqrt{T}(\hat{\Phi}_s - \Phi_s)$ is asymptotically Gaussian for $s = 1, 2, \dots, p$. The usual OLS t -test and F -test are asymptotically valid. However, tests for Granger-causality based on VAR with unit roots do not have the usual χ^2 distribution, hence would not be valid.

32.1 Monte Carlo

Below is a Monte Carlo simulation of a 2-dimensional VAR process with unit root, which verifies the Gaussian distribution of its coefficients.

```
library(tsDyn)
library(vars)
set.seed(0)
bhat = sapply(1:1000, function(i) {
  # this is a VAR with unit root
  B = matrix(c(0.7, 0.1, 0.3, 0.9), 2)
  # simulate the VAR process
  sim <- VAR.sim(B, n = 300, include = "none")
  mod = VAR(sim); b = coef(mod)
  # extract the coefficients
  c(B11 = b$y1['y1.l1', 'Estimate'],
    B12 = b$y1['y2.l1', 'Estimate'],
    B21 = b$y2['y1.l1', 'Estimate'],
    B22 = b$y2['y2.l1', 'Estimate'])
}) |> t()
# plot the distribution of the coefficients
{
  par(mfrow=c(2,2), mar=c(2,2,2,2))
  hist(bhat[, 'B11'], freq=F, main="B11")
  hist(bhat[, 'B12'], freq=F, main="B12")
  hist(bhat[, 'B21'], freq=F, main="B21")
  hist(bhat[, 'B22'], freq=F, main="B22")
}
```




Figure 32.1: Distributions of the VAR coefficients by Monte Carlo simulation

32.2 Conclusions

Economic time series usually comes in seasonally-adjusted (log) levels, which often involve unit roots. Researchers have to make the choice whether to difference the data to stationary or leave it as it is when modelling. There is no single principle to rule them all. It depends on the purpose of the research. It might feel safe to work with stationary time series only. Though stationarity is not necessary for VARs to work properly. Here are the tips from Walter Enders:

To difference or not to difference

- If the coefficient of interest can be written as a coefficient on a stationary variable, then a t -test is appropriate.
- You can use t -tests or F -tests on the stationary variables.
- You can perform a lag length test on any variable or any set of variables.
- Generally, you cannot use Granger causality tests concerning the effects of a non-stationary variable.
- The issue of differencing is important. If the VAR can be written entirely in first differences, hypothesis tests can be performed on any equation or any set of equations

using t -tests or F -tests.

- It is possible to write the VAR in first differences if the variables are $I(1)$ and are not cointegrated. If the variables in question are cointegrated, the VAR cannot be written in first differences.

33 VECM*

33.1 Cointegrated Systems

Definition 33.1. An $n \times 1$ vector y_t is said to be **cointegrated** if each of its elements individually is $I(1)$ and there exists a non-zero vector $a \in \mathbb{R}^n$ such that $a'y_t$ is stationary. a is called a **cointegrating vector**.

If there are $h < n$ linearly independent cointegrating vectors (a_1, a_2, \dots, a_h) , then any linear combination $k_1 a_1 + k_2 a_2 + \dots + k_h a_h$ is also a cointegrating vector. Thus, we say (a_1, a_2, \dots, a_h) form a basis for **the space of cointegrating vectors**.

Cointegrated systems can be represented by a **Vector Error Correlation Model (VECM)**:

$$\Delta y_t = \Pi y_{t-1} + \Gamma_1 \Delta y_{t-1} + \Gamma_2 \Delta y_{t-2} + \dots + \Gamma_{p-1} \Delta y_{t-p+1} + \mu + \epsilon_t \quad (33.1)$$

Theorem 33.1 (Engle-Granger Representation Theorem). *Any set of $I(1)$ variables are cointegrated if and only if there exists an error correlation (ECM) representation for them.*

Therefore, it is inappropriate to model a cointegrated system with differenced VAR. Because the term Πy_{t-1} is missing out, which means a misspecification.

The cointegration term can be further factored as $\Pi = \underset{n \times h}{A} \times \underset{h \times n}{B'}$, in which B contains the cointegrating vectors and A hosts the adjustment coefficients. In a bivariate example, it looks like

$$\begin{bmatrix} \Delta y_{1t} \\ \Delta y_{2t} \end{bmatrix} = \begin{bmatrix} \alpha_1 \\ \alpha_2 \end{bmatrix} \begin{bmatrix} \beta_1 & \beta_2 \end{bmatrix} \begin{bmatrix} y_{1t-1} \\ y_{2t-1} \end{bmatrix} + \sum_{j=1}^{p-1} \begin{bmatrix} \gamma_{j,11} & \gamma_{j,12} \\ \gamma_{j,21} & \gamma_{j,22} \end{bmatrix} \begin{bmatrix} \Delta y_{1t-j} \\ \Delta y_{2t-j} \end{bmatrix} + \begin{bmatrix} \epsilon_{1t} \\ \epsilon_{2t} \end{bmatrix}$$

The economic interpretation is that $\beta_1 y_{1t} + \beta_2 y_{2t}$ represents some *long-run equilibrium* relationship of the two variables. Parameters α_1, α_2 describe the speed of adjustment, that is how each variable reacts to the deviations from the equilibrium path. Small values of α_i would imply a relatively unresponsive reaction, which means it takes a long time to return to the equilibrium.

Note that $\Pi = AB'$ cannot be a full rank matrix. If there are n independent cointegrating vectors, it follows that any linear combination of the components of y_t is stationary, which effectively means y_t is stationary. Π cannot be zero either. If this is the case, the system is fully characterized by differenced VAR, there is no cointegration. Therefore, for a cointegrated system, it necessitates $0 < h < n$.

A three variable example would be like:

$$\begin{bmatrix} \Delta y_{1t} \\ \Delta y_{2t} \\ \Delta y_{3t} \end{bmatrix} = \begin{bmatrix} \alpha_{11} & \alpha_{12} \\ \alpha_{21} & \alpha_{22} \\ \alpha_{31} & \alpha_{32} \end{bmatrix} \begin{bmatrix} \beta_{11} & \beta_{12} & \beta_{13} \\ \beta_{21} & \beta_{22} & \beta_{23} \end{bmatrix} \begin{bmatrix} y_{1t-1} \\ y_{2t-1} \\ y_{3t-1} \end{bmatrix} + \dots$$

In general, if y_t has n non-stationary components, there could be at most $n - 1$ cointegrating vectors. The number of cointegrating vectors is also called the **cointegrating rank**.

Note that a single equation ECM is equivalent to an ARDL model:

$$\begin{aligned} \Delta y_t &= \alpha(y_{t-1} - \delta - \beta x_{t-1}) + \gamma \Delta x_t + u_t \\ \Leftrightarrow y_t &= (\alpha + 1)y_{t-1} + \gamma x_t - (\alpha\beta + \gamma)x_{t-1} - \alpha\delta + u_t \\ \Leftrightarrow y_t &= b_1 y_{t-1} + b_2 x_t + b_3 x_{t-1} + c + u_t \end{aligned}$$

Given a (possibly) cointegrated system, we would like to know if any cointegrating relations exist and how many cointegrating vectors there are. Johansen (1991) provides a likelihood-based method to test and estimate a cointegrated system. But before we introduce the Johansen's method, we need the prerequisite knowledge of canonical correlations.

33.2 Canonical Correlation

Principle component analysis (PCA) finds a linear combination of $[x_1 \ x_2 \ \dots x_n]$ that produces the largest variance. What if we want to extend the analysis to the correlations between two datasets: $\underset{T \times n}{X} = [x_1 \ x_2 \ \dots x_n]$ and $\underset{T \times m}{Y} = [y_1 \ y_2 \ \dots y_m]$? The cross-dataset covariance matrix is

$$\Sigma_{XY} = \underset{n \times m}{\begin{bmatrix} \text{cov}(x_1, y_1) & \text{cov}(x_1, y_2) & \dots & \text{cov}(x_1, y_m) \\ \text{cov}(x_2, y_1) & \text{cov}(x_2, y_2) & \dots & \text{cov}(x_2, y_m) \\ \vdots & \vdots & \ddots & \vdots \\ \text{cov}(x_n, y_1) & \text{cov}(x_n, y_2) & \dots & \text{cov}(x_n, y_m) \end{bmatrix}}$$

Canonical correlation analysis (CCA) seeks two vectors $a \in \mathbb{R}^n$ and $b \in \mathbb{R}^m$ such that $a'X$ and $b'Y$ maximize the correlation $\rho = \text{corr}(a'X, b'Y)$. The random transformed random variable

$U = a'X$ and $V = b'Y$ are the first pair of canonical variables. The second pair of canonical variables are orthogonal to the first pair and maximize the same correlation, and so on.

Suppose we want to choose a and b to maximize

$$\rho = \frac{a'\Sigma_{XY}b}{\sqrt{a'\Sigma_{XX}a}\sqrt{b'\Sigma_{YY}b}}$$

We may impose the constraint such that $a'\Sigma_{XX}a$ and $b'\Sigma_{YY}b$ normalize to 1. Form the Lagrangian

$$\mathcal{L} = a'\Sigma_{XY}b - \frac{\mu}{2}(a'\Sigma_{XX}a - 1) - \frac{\nu}{2}(b'\Sigma_{YY}b - 1)$$

The first-order conditions are

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial a} &= \Sigma_{XY}b - \mu \Sigma_{XX}a = 0 \\ \frac{\partial \mathcal{L}}{\partial b} &= \Sigma_{YX}a - \nu \Sigma_{YY}b = 0\end{aligned}$$

which implies

$$\Sigma_{XX}^{-1}\Sigma_{XY}\Sigma_{YY}^{-1}\Sigma_{YX}a = \mu\nu a = \lambda a$$

$$\Sigma_{YY}^{-1}\Sigma_{YX}\Sigma_{XX}^{-1}\Sigma_{XY}b = \mu\nu b = \lambda b$$

Therefore, a is the eigenvector of $\Sigma_{XX}^{-1}\Sigma_{XY}\Sigma_{YY}^{-1}\Sigma_{YX}$. The associated eigenvalue λ is equivalent to the maximized correlation squared ρ^2 . To see this, just multiply the first-order condition by a' :

$$a'\Sigma_{XY}b = \mu \cdot a'\Sigma_{XX}a = \mu = \rho^*$$

Symmetrically, b is the eigenvector of $\Sigma_{YY}^{-1}\Sigma_{YX}\Sigma_{XX}^{-1}\Sigma_{XY}$.

So the canonical correlation can be computed as follows:

1. Compute the eigenvalues and eigenvectors of

$$\Sigma_{XX}^{-1}\Sigma_{XY}\Sigma_{YY}^{-1}\Sigma_{YX}$$

2. Sort the eigenvalues as $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_n$ and a_1, a_2, \dots, a_n are the corresponding eigenvectors such that $a'\Sigma_{XX}a = 1$.

3. Compute the eigenvalues and eigenvectors of

$$\Sigma_{YY}^{-1} \Sigma_{YX} \Sigma_{XX}^{-1} \Sigma_{XY}$$

4. Sort the eigenvalues as $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_m$ and b_1, b_2, \dots, b_m are the corresponding eigenvectors such that $b' \Sigma_{XX} b = 1$.
5. (a_k, b_k) is the k -th pair of canonical variables, where $k \leq \min\{m, n\}$; and γ_k is the k -th largest canonical correlation.

Note that, if $A = [a_1 \ a_2 \ \dots \ a_k]_{n \times k}$, $B = [b_1 \ b_2 \ \dots \ b_k]_{m \times k}$, we would have $A' \Sigma_{XY} B = \Lambda$ where

$$\Lambda = \begin{bmatrix} \rho_1 & & & \\ & \rho_2 & & \\ & & \ddots & \\ & & & \rho_k \end{bmatrix}$$

is a diagonal matrix. Therefore, canonical variables transform the covariance matrix between X and Y into a diagonal matrix, where entries on the diagonal best summarize the correlations between the two datasets.

33.3 Johansen's Procedure

1. Estimate by OLS two regressions

$$\begin{aligned} \Delta y_t &= \hat{\Psi}_0 + \hat{\Psi}_1 \Delta y_{t-1} + \dots + \hat{\Psi}_{p-1} \Delta y_{t-p+1} + \hat{u}_t \\ y_{t-1} &= \hat{\Theta}_0 + \hat{\Theta}_1 \Delta y_{t-1} + \dots + \hat{\Theta}_{p-1} \Delta y_{t-p+1} + \hat{v}_t \end{aligned}$$

Save \hat{u}_t and \hat{v}_t .

2. Compute the canonical correlations between \hat{u}_t and \hat{v}_t . Find the eigenvalues of $\Sigma_{vv}^{-1} \Sigma_{vu} \Sigma_{uu}^{-1} \Sigma_{uv}$. Sort them from the largest to the smallest: $\hat{\lambda}_1 \geq \hat{\lambda}_2 \geq \dots \geq \hat{\lambda}_n$. Then the maximum log-likelihood function subject to the constraint that there are h cointegrating relations is given by

$$\ell^* = -\frac{Tn}{2} \log(2\pi) - \frac{Tn}{2} - \frac{T}{2} \log |\hat{\Sigma}_{uu}| - \frac{T}{2} \sum_{i=1}^h \log(1 - \hat{\lambda}_i)$$

The test for there being h cointegrating relations is equivalent to testing

$$-\frac{T}{2} \sum_{i=h+1}^n \log(1 - \hat{\lambda}_i) = 0.$$

3. Calculate the MLE of the parameters. The cointegrating matrix is given by

$$\hat{B} = \begin{bmatrix} \hat{b}_1 & \hat{b}_2 & \dots & \hat{b}_h \end{bmatrix}$$

where \hat{b}_i are the eigenvectors used to normalize v_t . The adjustment matrix and other parameters are given by

$$\begin{aligned} \hat{A} &= \hat{\Sigma}_{uv} \hat{B} \\ \hat{\Pi} &= \hat{A} \hat{B}' \\ \hat{\Gamma}_i &= \hat{\Psi}_i - \hat{\Pi} \hat{\Theta}_i \\ \mu &= \hat{\Psi}_0 - \hat{\Pi} \hat{\Theta}_0. \end{aligned}$$

To understand this procedure, note that if we treat Π as given, the MLE for Equation 33.1 is equivalent to estimating the coefficients by OLS:

$$\Delta y_t - \Pi y_{t-1} = \hat{\Gamma}_1 \Delta y_{t-1} + \hat{\Gamma}_2 \Delta y_{t-2} + \dots + \hat{\Gamma}_{p-1} \Delta y_{t-p+1} + \hat{\mu} + \hat{\epsilon}_t$$

The log-likelihood function becomes

$$\ell(\Pi, \Omega) = -\frac{T}{2} \log |2\pi\Omega| - \frac{1}{2} \sum_{t=1}^T (\hat{\epsilon}_t' \Omega^{-1} \hat{\epsilon}_t).$$

Step 1 does this in two separate regressions, in which

$$\begin{aligned} \hat{\Gamma}_i &= \hat{\Psi}_i - \Pi \hat{\Theta}_i \\ \hat{\epsilon}_t &= \hat{u}_t - \Pi \hat{v}_t \end{aligned}$$

Thus, the log-likelihood can be rewritten as

$$\ell(\Pi, \Omega) = -\frac{T}{2} \log |2\pi\Omega| - \frac{1}{2} \sum_{t=1}^T [(\hat{u}_t - \Pi \hat{v}_t)' \Omega^{-1} (\hat{u}_t - \Pi \hat{v}_t)]$$

Further concentrating Ω :

$$\hat{\Omega} = \frac{1}{T} \sum_{t=1}^T [(\hat{u}_t - \Pi \hat{v}_t)(\hat{u}_t - \Pi \hat{v}_t)']$$

Substituting this into the ℓ function

$$\ell(\Pi) = -\frac{Tn}{2} \log(2\pi) - \frac{Tn}{2} - \frac{T}{2} \log \left| \frac{1}{T} \sum_{t=1}^T [(\hat{u}_t - \Pi \hat{v}_t)(\hat{u}_t - \Pi \hat{v}_t)'] \right|$$

Thus, maximizing ℓ is equivalent to minimizing

$$\left| \frac{1}{T} \sum_{t=1}^T [(\hat{u}_t - \Pi \hat{v}_t)(\hat{u}_t - \Pi \hat{v}_t)'] \right|$$

by choosing Π . If u_t is a single variable, the optimal $\hat{\Pi}$ that minimizes $|T^{-1} \sum_t (u_t - \Pi v_t)^2|$ would be simply the OLS estimator. Similarly, for the vector case, we have

$$\hat{\Pi} = \left(\frac{1}{T} \sum_t u_t v_t' \right) \left(\frac{1}{T} \sum_t v_t v_t' \right)^{-1}$$

Suppose we have the canonical decomposition $A' \Sigma_{uv} B = \Lambda$ where Λ is a diagonal matrix of canonical correlations, and $A' \Sigma_{uu} A = I$, $B' \Sigma_{vv} B = I$. The estimator can be reduced to

$$\begin{aligned} \hat{\Pi} &= (A'^{-1} \Lambda B^{-1})(B'^{-1} B^{-1})^{-1} = A'^{-1} \Lambda B' \\ &= A'^{-1} \begin{bmatrix} r_1 & & & \\ & r_2 & & \\ & & \ddots & \\ & & & r_n \end{bmatrix} B' \end{aligned}$$

The minimized ‘squared residuals’ is

$$\begin{aligned} \left| \frac{1}{T} \sum_{t=1}^T [(\hat{u}_t - \Pi \hat{v}_t)(\hat{u}_t - \Pi \hat{v}_t)'] \right| &= \left| A'^{-1} \frac{1}{T} \sum_{t=1}^T (A' u_t - \Lambda B' v_t)(A' u_t - \Lambda B' v_t)' A^{-1} \right| \\ &= |A|^{-2} |I - \Lambda \Lambda'| \\ &= |A|^{-2} \begin{vmatrix} 1 - r_1^2 & & & \\ & 1 - r_2^2 & & \\ & & \ddots & \\ & & & 1 - r_n^2 \end{vmatrix} \\ &= |A|^{-2} \prod_{i=1}^n (1 - \lambda_i). \end{aligned}$$

This explains the likelihood function in Step 2. The MLE for $\hat{\Pi}$ above is subject to no constraint. However, if there is any cointegrating relations, Π cannot be full rank. If we restrict the rank of Π to h , the minimized squared residuals is achieved by picking the h largest λ s.

33.4 Hypothesis Testing

Test 1: At most h cointegrations

Hypothesis:

H_0 : there are no more than h cointegrating relations

H_1 : there are more than h cointegrating relations

Test statistics:

$$\lambda_{\text{trace}} = 2(\ell_1^* - \ell_0^*) = -T \sum_{i=h+1}^n \log(1 - \hat{\lambda}_i)$$

If H_0 is true, λ_{trace} should be close to zero. The critical values are provided by the table below. *Case 1* means there is no constant or deterministic trend; *Case 2* contains constants in cointegrating vectors but no deterministic trend; *Case 3* contains deterministic trend.

Table 33.1: Critical values of Johansen's likelihood ratio test of the null hypothesis of h integrating relations against the alternative of no restrictions

$n - h$	T	0.1	0.05	0.025	0.001
Case 1					
1	400	2.86	3.84	4.93	6.51
2	400	10.47	12.53	14.43	16.31
...					
Case 2					
1	400	6.69	8.08	9.66	11.58
2	400	15.58	17.84	19.61	21.96
...					
Case 3					
1	400	2.82	3.96	5.33	6.94
2	400	13.34	15.20	17.30	19.31
...					

Test 2: h cointegrations vs h+1

Hypothesis:

H_0 : there are h cointegrating relations

H_1 : there are $h + 1$ cointegrating relations

Test statistics:

$$\lambda_{\max} = 2(\ell_1^* - \ell_0^*) = -T \log(1 - \hat{\lambda}_{h+1})$$

The critical values are given as below.

Table 33.2: Critical values of Johansen's likelihood ratio test of the null hypothesis of h integrating relations against the alternative of $h + 1$ relations

$n - h$	T	0.1	0.05	0.025	0.001
Case 1					
1	400	2.86	3.84	4.96	6.51
2	400	9.52	11.44	13.27	15.69
...					
Case 2					
1	400	6.69	8.08	9.66	11.58
2	400	12.78	14.60	16.40	18.78
...					
Case 3					
1	400	2.82	3.96	5.33	6.94
2	400	12.10	14.04	15.81	17.94
...					

Part VI

Bayesian Analysis

34 Intro to Bayes

We have completed our venture on classical time series topics covering univariate and multivariate, stationary and non-stationary models. This chapter taps into the Bayesian approach to time series analysis, which is not an essential component for classical treatment of the subject. But given its rising popularity and importance, it is an exciting topic that cannot be missed. Bayesian statistics is a whole new world for frequentist statisticians. We introduce the Bayesian approach with an example.

34.1 The Sunrise Problem

Question: What is the probability that the sun will rise tomorrow?

We do not consider the physics here. Suppose we want to answer this question by purely statistics. What we need to do is to observe how many days the sun had risen in the past, and make some inference about the future. If we had collect the data on the past n days, we would have observed the sun had risen everyday for sure (the sun rises even in cloudy or rainy days). If we want to calculate the probability of a sunrise event, denoted by A , the frequentist approach would give $P(A) = \frac{n}{n} = 1$. The probability is always 1 no matter how many observations we have. That's a bit quirky, even though we haven't looked at the confidence interval. The 100 percent probability does not sound correct, as nothing can be so certain.

Let's have a look at how Laplace in the 18th century solves this problem. Let x_t be a random variable such that

$$x_t = \begin{cases} 1, & \text{the sun rise in day } t \text{ with probability } \theta \\ 0, & \text{otherwise} \end{cases}$$

In other words, x_t follows a Bernoulli distribution $x_t \sim \text{Bern}(\theta)$. There is an unknown parameter θ , which is our goal to estimate. Before we have observed any data, we have no knowledge about this θ . We assume it is distributed uniformly, $\theta \sim \text{Unif}(0, 1)$. That is, it can be any value between 0 and 1.

Suppose we have observed the data for n days: x_1, x_2, \dots, x_n . Assume these events are *i.i.d.* Define S_n as the total number of sunrises that had happened:

$$S_n = x_1 + x_2 + \cdots + x_n$$

We know S_n follows a Binomial distribution $S_n \sim \text{Bin}(n, \theta)$, with the probability mass function

$$P(S_n = k|\theta) = \binom{n}{k} \theta^k (1 - \theta)^{n-k}$$

Our goal is to find: $\theta|S_n = ?$ An estimation of the probability of a sunrise after observing the data.

Recall that the Bayesian rule allows us to invert the conditional probability:

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

Using this formula, we have

$$\begin{aligned} P(\theta|S_n = k) &= \frac{P(S_n = k|\theta)P(\theta)}{P(S_n = k)} = \frac{P(S_n = k|\theta)P(\theta)}{\int_0^1 P(S_n = k|\theta)P(\theta)d\theta} \\ &= \frac{\binom{n}{k} \theta^k (1 - \theta)^{n-k} \cdot \mathbb{I}(0 \leq \theta \leq 1)}{\int_0^1 \binom{n}{k} \theta^k (1 - \theta)^{n-k} \cdot \mathbb{I}(0 \leq \theta \leq 1) d\theta} \\ &= \begin{cases} \frac{\binom{n}{k} \theta^k (1 - \theta)^{n-k}}{\int_0^1 \binom{n}{k} \theta^k (1 - \theta)^{n-k} d\theta} & \text{if } 0 \leq \theta \leq 1 \\ 0 & \text{otherwise} \end{cases} \end{aligned}$$

If $k = n$, we have

$$P(\theta|S_n = n) = \frac{\theta^n}{\int_0^1 \theta^n d\theta} = (n + 1)\theta^n$$

for $0 \leq \theta \leq 1$. Now we are ready to calculate the probability of the sun rising tomorrow after observing n sunrises:

$$\begin{aligned} P(x_{n+1} = 1|S_n = n) &= \int_0^1 P(x_{n+1} = 1|\theta)P(\theta|S_n = n)d\theta \\ &= \int_0^1 \theta \cdot (n + 1)\theta^n d\theta \\ &= \frac{n + 1}{n + 2} \end{aligned}$$

As $n \rightarrow \infty$, the probability approaches 1. Personally, I think this is much more reasonable answer than the frequentist approach, in which you always get probability 1.

34.2 The Bayesian Approach

This illustration literally shows every tenet of the Bayesian approach. We start with an **prior distribution** about an unknown parameter θ . In the previous example, we model it as a uniform distribution because of our ignorance. But the prior can be any distribution reflecting our subjective belief about the parameter before we see the data. Note how this contrasts with the frequentist approach, in which the parameter is a fixed unknown number.

The principle of Bayesian analysis is then to combine the prior information with the information contained in the data to obtain an updated distribution accounting for both sources of information, known as the **posterior distribution**. This is done by using the Bayes rule:

$$\underbrace{p(\theta|X)}_{\text{posterior}} = \frac{\overbrace{p(X|\theta)}^{\text{likelihood}} \times \overbrace{p(\theta)}^{\text{prior}}}{\underbrace{p(X)}_{\text{normalizing scalar}}}$$

The posterior distribution $p(\theta|X)$ is the central object for Bayesian inference as it combines all the information we have about θ . Note that $p(\theta|X)$ is a function of θ . Since the denominator $p(X)$ is independent of θ , it only plays the role of a normalizing constant to ensure the posterior is a valid probability density function that integrates to 1. It is therefore convenient to ignore it and rewrite the posterior as

$$p(\theta|X) \propto p(X|\theta)p(\theta)$$

One difficulty of Bayesian inference is that the denominator $p(X)$ is often impossible to compute, especially for high dimensional parameters:

$$p(X) = \int_{\theta_1} \int_{\theta_2} \cdots \int_{\theta_n} p(X, \theta_1, \theta_2, \dots, \theta_n) d\theta_1 d\theta_2 \dots d\theta_n$$

But the relative frequencies of parameter values are easy to compute

$$\frac{p(\theta_A|X)}{p(\theta_B|X)} = \frac{p(X|\theta_A)p(\theta_A)}{p(X|\theta_B)p(\theta_B)}$$

This allows us to sample from the posterior distribution even the *pdf* of the distribution is unknown. We will return to this point when we discuss computational Bayesian methods.

The relative weight of the prior versus the data in determining the posterior depends on (i) how strong the prior is, and (ii) how many data we have. If the prior is so strong (very small variance / uncertainty) that seeing the data will not change our beliefs, the posterior would be mostly determined by the prior. On the contrary, if the data is so abundant that the evidence overwhelms any prior belief, the impact of prior would be negligible.

34.3 Frequentist vs Bayesian

Frequentists and Bayesians hold different philosophy about statistics. Frequentists view our sample as the result of one of an infinite number of exactly repeated experiments. The data are **randomly sampled** from a fixed population distribution. The unknown parameters are properties of the population, and therefore are fixed. The purpose of statistics is to make inference about the population parameters (the ultimate truth) with limited samples. The uncertainty associated with this process arises from sampling. Because we do not have the entire population, each sample only tells partial truth about the population. Therefore our inference about the parameters can never be perfect due to sampling errors. Frequentists conduct hypothesis tests assuming a hypothesis (about the population parameter) is true and calculating the probability of obtaining the observed sample data.

In Bayesians' world view, probability is an expression of **subjective beliefs** (a measure of certainty in a belief), which can be updated in light of new data. Parameters are probabilistic rather than fixed, which reflects the uncertainties about the parameters. The essence of Bayesian inference is to update the probability of a 'hypothesis' given the data we have obtained. The Bayes' rule is all we need. All information is summarized in the posterior probability and there is not need for explicit hypothesis testing.

Frequentist	Bayesian
Probability is the limit of frequency	Probability is uncertainty
Parameters are fixed unknown numbers	Parameters are random variables
Data is a random sample from the population	Data is fixed/given
LLN/CLT	Bayes' rule

In time series analysis, there are good reasons to be Bayesian. Perhaps the frequentist perspective makes sense in a cross section, where it is intuitive to image taking different samples from the population. However, in time series we have only one realization. It is difficult to imagine where we would obtain another sample. It is more natural to take a Bayesian perspective. For example, we have some prior belief on how inflation and unemployment might be related (the Phillips curve), then we update our belief with data.

Frequentists often criticize Bayesians' priors as entirely subjective. Bayesians would respond that frequentists also have prior assumptions that they are not even aware of. Frequentist

inference utilizes the LLN and CLT, which inevitably assumes the speed of convergence. In settings like VAR models, where there are a large number of parameters to estimate but only a limited amount of observations. Are the asymptotically properties really plausible? Bayesians believe it would be better to make our assumptions explicit.

Apart from the philosophical difference, in practice Frequentists and Bayesians might well give similar results (though the results should be interpreted differently). After all, if the data is plenty, the influence of priors would diminish to zero.

35 Linear Model

35.1 Linear Regression with Known σ^2

Let's use the Bayesian principle to estimate a simple linear regression:

$$y_t = x_t\beta + \epsilon_t$$

where $\epsilon_t \sim N(0, \sigma^2)$. For simplicity, we assume σ^2 is known. So the only unknown parameter is β . Assume it has a Gaussian prior

$$\beta \sim N(\beta_0, V_\beta)$$

Gaussian prior is a handy prior to express our belief about the mean and the degree of certainty of that belief (expressed by the variance).

Traditional OLS works without specifying the distribution of ϵ_t . However, for Bayesian inference to work, we always need to specify the full distribution of the model. With Gaussian errors, we have

$$(\mathbf{Y}|\beta) \sim N(\mathbf{X}\beta, \sigma^2 \mathbf{I}_T)$$

Under the *i.i.d* assumption, the joint likelihood function is

$$\begin{aligned} p(\mathbf{Y}|\beta) &= \prod_{t=1}^T \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(y_t - x_t\beta)^2} \\ &= (2\pi\sigma^2)^{-\frac{T}{2}} e^{-\frac{1}{2\sigma^2} \sum_t (y_t - x_t\beta)^2} \end{aligned}$$

By the Bayes rule, the posterior distribution is

$$\begin{aligned}
p(\beta|\mathbf{Y}) &\propto p(\mathbf{Y}|\beta)p(\beta) \\
&\propto e^{-\frac{1}{2\sigma^2} \sum_t (y_t - x_t\beta)^2} \cdot e^{-\frac{1}{2V_\beta} (\beta - \beta_0)^2} \\
&\propto e^{-\frac{1}{2} \left(\frac{\sum x_t}{\sigma^2} + \frac{1}{V_\beta} \right) \beta^2 + \left(\frac{\sum x_t y_t}{\sigma^2} + \frac{\beta_0}{V_\beta} \right) \beta}
\end{aligned}$$

which is the kernel of a Gaussian distribution. Therefore,

$$p(\beta|\mathbf{Y}) \sim N(\hat{\beta}, D_\beta) \propto e^{-\frac{1}{2D_\beta} \beta^2 + \frac{\hat{\beta}}{D_\beta} \beta}$$

where

$$D_\beta = \left(\frac{\sum x_t}{\sigma^2} + \frac{1}{V_\beta} \right)^{-1}, \quad \hat{\beta} = \left(\frac{\sum x_t y_t}{\sigma^2} + \frac{\beta_0}{V_\beta} \right) D_\beta.$$

Note that if we have a very loose prior $V_\beta \rightarrow \infty$, or abundant data $N \rightarrow \infty$, we would have

$$D_\beta = \frac{\sigma^2}{\sum x_t}, \quad \hat{\beta} = \frac{\sum x_t y_t}{\sum x_t}$$

which is exactly the same as the OLS estimator.

So with a Gaussian prior and a Gaussian likelihood, the posterior distribution is also Gaussian. It is this particular choice of the prior and the likelihood function that the posterior has a closed-form solution. Not many prior choice has this property. This is what we called a **conjugate prior**.

35.2 Linear Regression with Unknown σ^2

For simplicity, we have assumed the variance σ^2 is known. In reality, if σ^2 is unknown, we also need to assign it a prior distribution. A common choice is an inverse-Gamma distribution:

$$\sigma^2 \sim IG(\nu_0, S_0)$$

whose density function is given by

$$p(\sigma^2) = \frac{S_0^{\nu_0}}{\Gamma(\nu_0)} (\sigma^2)^{-(\nu_0+1)} e^{-\frac{S_0}{\sigma^2}}$$

One reason of this choice is that an inverse-Gamma can never be negative. Another reason is that it is also a conjugate prior. It can be shown, with an inverse-Gamma as the prior for the variance and an Gaussian likelihood, the posterior for σ^2 is also an inverse-Gamma:

$$(\sigma^2|\mathbf{Y}, \beta) \sim IG\left(\nu_0 + \frac{T}{2}, S_0 + \frac{1}{2} \sum_{t=0}^T (y_t - x_t\beta)^2\right).$$

35.3 Credible Interval

Once the posterior distribution is obtained, the question becomes how to report and interpret the results. Similar to conventional OLS results, we would like to report the mean or median of the parameter, and the associated “credible interval”. The credible interval is directly obtained from the distribution:

$$P(\beta_L \leq \beta \leq \beta_U) = \alpha$$

which indicate that β falls between the range $[\beta_L, \beta_U]$ with a probability α . In a frequentist approach, a p -value is not the probability of the parameter, nor does confidence interval represent the distribution of the parameter. However, the credible interval obtained from a Bayesian posterior is the probability for particular values of the parameter. It is more straightforward to interpret. After all, parameters are themselves probabilistic in a Bayesian world.

36 Bayesian VAR

Once we understand the Bayesian approach to estimate simple linear regression models, it is easy to extend it to more complicated linear models. There are many reasons why a Bayesian approach to VAR models is preferred over a frequentist approach. VAR models have a proliferation of parameters to estimate. But economic time series typically do not come with many observations. The “curse of dimensionality” is a serious problem for frequentists. The benefit of using a Bayesian approach is that the priors can provide “shrinkage” over the parameters, that is strong priors for some parameters to reduce the burden of the data.

36.1 Vectorized Form

For Bayesian inference, it is easier to work with the vectorized form of a VAR in Chapter 27:

$$y = \bar{X}\beta + u$$

where $y = \text{vec}(Y)$, $\bar{X} = I_n \otimes X$, $\beta = \text{vec}(B)$, $u = \text{vec}(U)$, and $\bar{\Sigma} = \Sigma \otimes I_T$. For a VAR with n variables and p lags, the vectorized form looks like

$$\begin{bmatrix} y_{1,1} \\ \vdots \\ y_{1,T} \\ \vdots \\ y_{n,1} \\ \vdots \\ y_{n,T} \end{bmatrix} = \begin{bmatrix} y'_0 & \cdots & y'_{1-p} & 0 & \cdots & 0 \\ \vdots & \ddots & \vdots & \vdots & & \vdots \\ y'_{T-1} & \cdots & y'_{T-p} & 0 & \cdots & 0 \\ 0 & \cdots & 0 & \ddots & y'_0 & \cdots & y'_{1-p} \\ \vdots & & \vdots & \vdots & \ddots & \vdots \\ 0 & \cdots & 0 & y'_{T-1} & \cdots & y'_{T-p} \end{bmatrix} \begin{bmatrix} A_1^{(1)} \\ \vdots \\ A_p^{(1)} \\ \vdots \\ A_1^{(n)} \\ \vdots \\ A_p^{(n)} \end{bmatrix} + \begin{bmatrix} u_{1,1} \\ \vdots \\ u_{1,T} \\ \vdots \\ u_{n,1} \\ \vdots \\ u_{n,T} \end{bmatrix}$$

Assume multivariate Gaussian distribution for the residuals $u \sim N(0, \bar{\Sigma})$, the likelihood for y is given by

$$p(y|\beta) = |2\pi\bar{\Sigma}|^{-1/2} \exp \left[-\frac{1}{2}(y - \bar{X}\beta)' \bar{\Sigma}^{-1} (y - \bar{X}\beta) \right]$$

Also assume that β has a multivariate Gaussian prior

$$\beta \sim N(\beta_0, \Omega_0)$$

The key is to specify β_0 and Ω_0 . We now introduce one of the simplest and yet the most popular prior setting for VAR models.

36.2 Minnesota Prior

The Minnesota prior is proposed by Litterman (1986). It is assumed that the VAR residual covariance matrix Σ is known. The only prior required is for parameters β . The essence of the Minnesota prior is to shrink the parameters of longer lags to zero.

The prior setting for β_0 is as following: as most observed macroeconomic variables seem to be characterized by a unit root, our prior belief should be that each endogenous variable included in the model presents a unit root in its first own lags, and the coefficients equal to zero for further lags and cross-variable coefficients. Therefore,

$$\mathbb{E}[A_k^{(ij)}] = \begin{cases} 1, & \text{if } i = j, k = 1 \\ 0, & \text{otherwise} \end{cases}$$

For stationary variables, the coefficient 1 can be replaced by, say, 0.8. Regarding the uncertainty of our belief, expressed in Ω_0 , it is assumed that no covariance exists between terms in β so that Ω_0 is diagonal. Furthermore, our prior shall become stronger (variance becomes smaller) for longer lags that they are closer to zero (i.e. shrinking longer lags to zero). Besides, correlations with lags on other variables are likely weaker than the correlations on their own lags (stronger shrinkage on coefficients relating to other variables).

$$\text{Var}[A_k^{(ij)}] = \lambda_1 \lambda_2^{1 - \mathbb{I}(i=j)} \frac{1}{k^{\lambda_3}} \frac{\sigma_i^2}{\sigma_j^2}$$

where λ_1 controls the overall tightness, λ_2 controls the tightness on cross-variable coefficients, λ_3 controls the speed at which coefficients on longer lags shrink to zero. σ_i^2 and σ_j^2 denote the OLS residual variance of the auto-regressive models estimated for variables i and j .

Finally, if any exogenous variables are included in the model, they should have priors centered at zero with large variance, as little is known about exogenous variables. A typical set of values for these hyper-parameters found in the literature are: $\lambda_1 = 0.1$, $\lambda_2 = 0.5$, $\lambda_3 = 1$ or 2 , λ_4 for exogenous variables should be greater than 100 .

Since the Minnesota prior assumes Σ is known, one has to obtain it beforehand. One method is to set the diagonal of Σ equal to the residual variance of individual AR models run on each

variable in the VAR. Alternatively, one can use the variance-covariance matrix of the VAR estimated by OLS.

36.3 The Posterior

Once the prior is determined, we can derive the posterior as follows

$$\begin{aligned} p(\beta|y) &\propto p(y|\beta)p(\beta) \\ &\propto \exp \left[-\frac{1}{2}(y - \bar{X}\beta)' \bar{\Sigma}^{-1}(y - \bar{X}\beta) \right] \times \exp \left[-\frac{1}{2}(\beta - \beta_0)' \Omega_0^{-1}(\beta - \beta_0) \right] \end{aligned}$$

After some manipulation, it can be shown that

$$p(\beta|y) \propto \exp \left[-\frac{1}{2}(\beta - \bar{\beta})' \bar{\Omega}^{-1}(\beta - \bar{\beta}) \right]$$

where

$$\bar{\Omega} = [\Omega_0^{-1} + \Sigma^{-1} \otimes X'X]^{-1}$$

$$\bar{\beta} = \bar{\Omega}[\Omega_0^{-1}\beta_0 + (\Sigma^{-1} \otimes X')y]$$

This is again the kernel of a multivariate Gaussian distribution. Therefore, the posterior distribution of β is characterized by

$$p(\beta|y) \sim N(\bar{\beta}, \bar{\Omega}).$$

Once an estimate for β is obtained, one can compute the IRF or FEVD accordingly. Typically in a Bayesian procedure, with one draw of $\beta^{(1)}$ from the posterior distribution, we compute one round of IRF⁽¹⁾; with a second draw of $\beta^{(2)}$, we compute another round of IRF⁽²⁾, and so on. With a collection of IRFs, we can get the median and the credible bands for the IRF.

37 Conjugate Priors

We have mentioned, many of the times, we do not have closed-form solution for the posterior. However, there is a class of models — pairs of likelihoods and priors — that an analytic posterior exists. These pairs of likelihoods and priors are referred as **conjugate**.

Table 37.1: Some common conjugate pairs

Likelihood	Prior	Posterior
Bernoulli	$\text{Beta}(\alpha, \beta)$	$\text{Beta}(\alpha + \sum x_i, \beta + n - \sum x_i)$
Binomial	$\text{Beta}(\alpha, \beta)$	$\text{Beta}(\alpha + \sum x_i, \beta + \sum n_i - \sum x_i)$
Multinomial	$\text{Dirichlet}(\alpha)$	$\text{Dirichlet}(\alpha + \sum x_i)$
Normal (known σ^2)	$N(\mu_0, \sigma_0^2)$	$N\left(\frac{1}{\frac{1}{\sigma_0^2} + \frac{n}{\sigma^2}} \left(\frac{\mu_0}{\sigma_0^2} + \frac{\sum x_i}{\sigma^2}\right), \left(\frac{1}{\sigma_0^2} + \frac{n}{\sigma^2}\right)^{-1}\right)$
Normal (known μ)	$IG(\alpha, \beta)$	$IG\left(\frac{\alpha+n}{2}, \frac{\beta + \sum (x_i - \mu)^2}{2}\right)$
Poisson	$\Gamma(\alpha, \beta)$	$\Gamma(\alpha + \sum x_i, \beta + n)$

Using conjugate priors, we can plug-in the data into the formula and get the exact posterior distribution. But the limitation is obvious. We are confined to use a given set of distributions, whereas other distributions do not have conjugate properties.

38 Gibbs Sampling

38.1 Computational Bayes

If we move away from conjugate priors, we are likely to have no closed-form solution for the posterior distribution. But we can approximate the posterior distribution by some computational algorithms. The class of these algorithms are famously known as **Markov chain Monte Carlo (MCMC)** methods. That is, one can construct a Markov chain that has the desired distribution as its equilibrium distribution.

One strategy to approximate the shape of a distribution is through **random sampling**. If we can draw random samples from a distribution, as the sample size grows, sampling frequency will approach the probability density. Properties of a distribution can also be estimated from finite samples, e.g. $\frac{1}{n} \sum_{i=1}^n x_i \rightarrow \mathbb{E}(x)$.

The problem is how to draw samples from a distribution without knowing its PDF. It seems to be an impossible task. But what if the conditional PDF is easier to compute?

38.2 Gibbs Sampler for Linear Regression Models

Consider the example in Chapter 35 with unknown variance. The joint distribution of $p(\beta, \sigma^2 | Y)$ is hard to come by. But the condition distribution is easier to compute:

$$\begin{aligned} p(\beta | Y, \sigma^2) &\sim N(\hat{\beta}, D_{\beta}) \\ p(\sigma^2 | Y, \beta) &\sim IG(\bar{\nu}, \bar{S}) \end{aligned}$$

Suppose we know how to draw random samples from the normal distribution and the inverse-Gamma distribution (in fact, if the CDF is known, a random sample can be generated by $y = \text{CDF}^{-1}(x)$, where $x \sim U(0, 1)$), we can generate samples $\{(\beta, \sigma^2)\}$ by the following algorithm.

Algorithm (Gibbs Sampling)

Pick some initial value $\beta^{(0)} = a_0$ and $\sigma^{2(0)} = b_0$. Repeat the following steps for $i = 1 \dots N$:

1. Draw $\sigma^{2(i)} \sim p(\sigma^2|Y, \beta^{(i-1)})$ (Inverse-Gamma);
2. Draw $\beta^{(i)} \sim p(\beta|Y, \sigma^{2(i)})$ (Normal).

The most efficient sampling is direct independent sampling (β, σ^2) from the posterior distribution. If this is not feasible, Gibbs sampler is a reasonable alternative. It can be imagined, in the stationary situation, sampling from the conditional distributions would be statistically identical to sampling from the joint probability distribution. However, Gibbs sampling can be highly inefficient if the posterior variables are highly correlated.

38.3 Gibbs Sampler for General Models

Here is the more general Gibbs algorithm. Suppose we have a model with k parameters $\boldsymbol{\theta} = (\theta_1, \theta_2, \dots, \theta_k)$. Assume we can derive the exact expressions for each of the conditional distributions: $p(\theta_i|\theta_1, \dots, \theta_{i-1}, \theta_{i+1}, \dots, \theta_k, Y)$. Further assume we can generate independent samples from each of them. Then the Gibbs sampler runs as follows.

Algorithm (Gibbs Sampling)

Initialize the algorithm with $(\theta_1^{(0)}, \theta_2^{(0)}, \dots, \theta_k^{(0)})$. Then repeat the following steps:

1. Choose a random parameter ordering, e.g. $(\theta_3, \theta_1, \theta_2, \dots)$. Denote the parameters with the new ordering $(\theta_1, \theta_2, \theta_3, \dots)$.
2. Sample from the conditional distribution for each parameter using the most up-to-date parameters. That is, draw samples from the following distributions subsequently:

$$\begin{aligned}
 & p(\theta_1^{(i)}|\theta_2^{(i-1)}, \theta_3^{(i-1)}, \dots, \theta_k^{(i-1)}, Y) \\
 & p(\theta_2^{(i)}|\theta_1^{(i)}, \theta_3^{(i-1)}, \dots, \theta_k^{(i-1)}, Y) \\
 & p(\theta_3^{(i)}|\theta_1^{(i)}, \theta_2^{(i)}, \dots, \theta_k^{(i-1)}, Y) \\
 & \vdots \\
 & p(\theta_k^{(i)}|\theta_1^{(i)}, \theta_2^{(i)}, \dots, \theta_{k-1}^{(i)}, Y)
 \end{aligned}$$

Repeat the process until the algorithm is reasonably converged.

39 Metropolis–Hastings

39.1 Dependent Sampling

We have mentioned, despite the posterior density is hard to compute, it is easy to determine the relative frequencies of the parameter values:

$$\frac{p(\theta_A|X)}{p(\theta_B|X)} = \frac{p(X|\theta_A)p(\theta_A)}{p(X|\theta_B)p(\theta_B)} \quad (39.1)$$

To generate a sample that approximate the posterior distribution, we require the values with higher probability density to be sampled proportionally more often.

Imagine we are exploring an unknown map. We start somewhere, each step forward depends on the current location. We require some areas (with high density) be explored more than others. This means the exploration moves cannot be purely “random walk”. We need some rules to dictate the “direction” of the next move. So that the exploration approximates the relative frequency suggested by Equation 39.1.

Note that this exploration (sampling) algorithm necessarily has dependencies. The next sample value depends on the current sample value. This is called dependent sampling as opposed to independent sampling. Dependent sampling naturally takes more samples to reach a reasonable approximation of the posterior than independent sampling. Because of the dependency, each move gives less information than independent sampling.

39.2 Random Walk Metropolis

We now introduce one of the most famous MCMC algorithm:

Algorithm (Random Walk Metropolis-Hastings)

Choose an arbitrary point θ_t to be the first sample value in the posterior space. Propose the next value according to a random walk:

$$\theta_{t+1} = \theta_t + \epsilon_{t+1}$$

If the proposed value has a higher density than the current value, we accept the proposal and move forward. Precisely, we accept the proposal with probability

$$r = \begin{cases} 1, & \text{if } p(\theta_{t+1}|X) \geq p(\theta_t|X) \\ \frac{p(\theta_t|X)}{p(\theta_{t+1}|X)}, & \text{if } p(\theta_{t+1}|X) < p(\theta_t|X) \end{cases}$$

r is the probability that we accept θ_{t+1} as our next sample value.

This Metropolis criteria is the only one that will work. It strikes the perfect balance between random exploring and paying attention to the posterior shape — If we pay too little attention, we get uniform sampling; if we pay too much attention, we will get stuck on a mode forever. It can be proved the algorithm converges to the posterior distribution. That is, if the algorithm reaches the posterior density, it stays there.

39.3 Remarks

Step size. The rate at which Metropolis converges to the posterior distribution is highly sensitive to the step size. If the step size is too small, we obtain a density that is highly dependent on the initial value. It takes a long time to find areas of high density. On the other hand, if the step size is too big, we would reject the majority of proposals, since most of the parameter space is low and flat, and we would get a highly autocorrelated chain with low number of effective samples. We would tune the step size so that rate of acceptance is optimized (0.44 for one-dimensional models, 0.23 for high-dimensional models).

Multiple chains. It is never a good idea to run a single chain (exploring from one starting point forever). Different chains (starting from different initial values) are likely exploring different density areas of the posterior space. It takes longer time for a single chain to explore the entire space. If different chains converge to the same distribution, we are confident the whole posterior space is explored.

Convergence. We can compare the within-chain variance and the between-chain variance to gauge the convergence.

$$\text{within-chain variance: } W = \frac{1}{m} \sum_{j=1}^m s_j^2$$

$$\text{between-chain variance: } B = \frac{n}{m-1} \sum_{j=1}^m (\bar{\theta}_j - \bar{\theta})^2$$

where s_j^2 is the sample variance of chain j , θ represents the sample mean, m is the number of chains. The convergence can be gauged by the R -ratio:

$$\hat{R} = \sqrt{\frac{W + \frac{1}{n}(B - W)}{W}}$$

Initially, $\hat{R} \gg 1$. The better the convergence, the closer \hat{R} is to 1.

Warm-up. The first part of the chain is selected in a haphazard fashion, and unlikely to be representative of the posterior. So we should not include the first few samples in our final posterior sample. Usually, it is recommended to discard the first half of the chains that appear to have converged as a default method.

Effective sample size. Dependent sampling naturally converges slower than independent sampling. For independent sampling, CLT predicts $\sqrt{T}(\hat{\theta} - \theta) \rightarrow N(0, \sigma^2)$. We say the convergence speed is $T^{-1/2}$. The effective sample size, n_{eff} , for a dependent sampler is defined so that its convergence speed is $n_{\text{eff}}^{-1/2}$.

Thinning. Once convergence is reached, we could make our samples look “more independent” if we keep only every tenth, or hundredth, of the samples. These will naturally be less correlated than the original samples. This process is known as “thinning”.

References

- Hamilton, James D. 1994. *Time Series Analysis*. Princeton University Press.
- Enders, Walter. 2008. *Applied Econometric Time Series*. John Wiley & Sons.
- Verbeek, Marno. 2008. *A Guide to Modern Econometrics*. John Wiley & Sons.
- Hayashi, Fumio. 2011. *Econometrics*. Princeton University Press.
- Hansen, Bruce. 2022. *Econometrics*. Princeton University Press.
- Mikusheva, Anna, and Paul Schrimpf. 2007. *Time Series Analysis*. MIT OpenCourseWare.
- Hyndman, Rob J, and George Athanasopoulos. 2018. *Forecasting: Principles and Practice* (2nd Edition). [OTexts.com/fpp2](https://otexts.com/fpp2).
- Stock, James H, and Mark W Watson. 2020. *Introduction to Econometrics*. Pearson.
- Stock, James H, and Mark W Watson. 2016. “Dynamic Factor Models, Factor-Augmented Vector Autoregressions, and Structural Vector Autoregressions in Macroeconomics.” In *Handbook of Macroeconomics*, 2:415–525.
- Ramey, V.A., 2016. *Macroeconomic Shocks and Their Propagation*. In *Handbook of Macroeconomics*, 2, pp.71-162.
- Chan, Joshua CC. 2017. “Notes on Bayesian Macroeconometrics.” Unpublished Manuscript.
- Dieppe, Alistair, Romain Legrand, and Björn Van Roye. 2016. *The Bayesian Estimation, Analysis and Regression (BEAR) Toolbox*. European Central Bank.